Inferential analysis

Jason G. Fleischer, Ph.D.
Asst. Teaching Professor
Department of Cognitive Science, UC San Diego
jfleischer@ucsd.edu


@jasongfleischer

https://jgfleischer.com

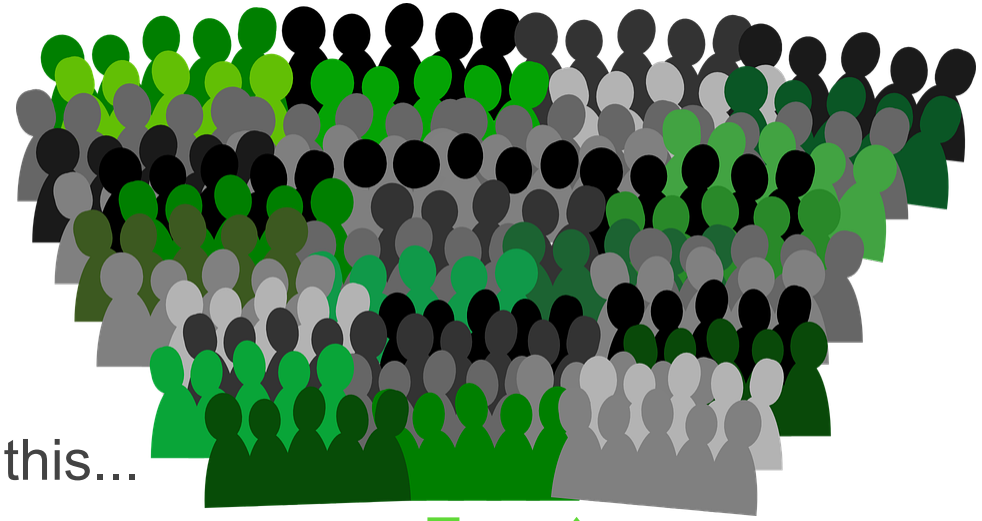

Population

We want to learn
something about this...

Sampling

Inference

Sample

....but we can only *actually*
collect data from this

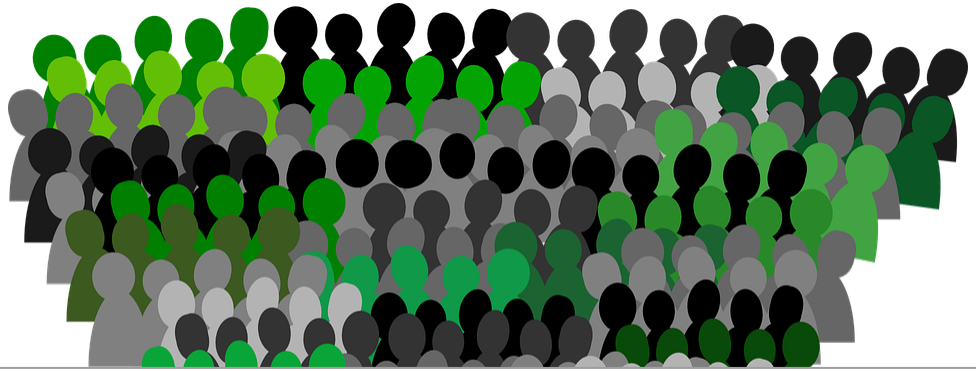

Random samples are random!
They differ from each other and the population!

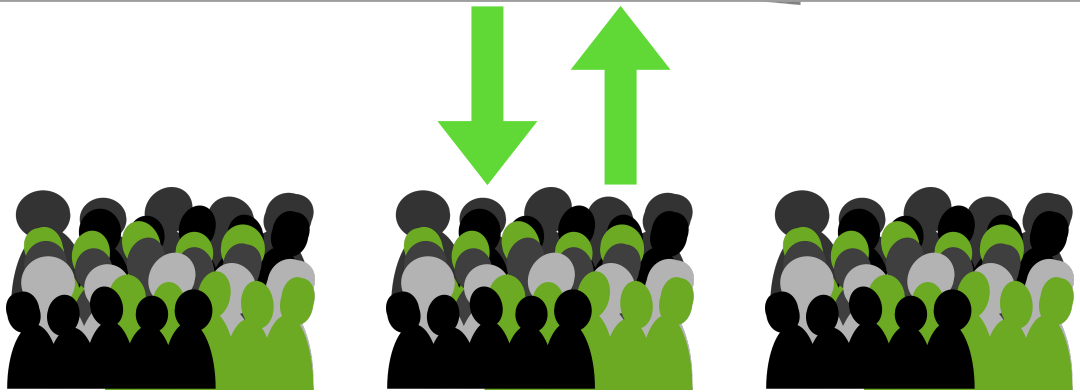

Published in final edited form as:
Epidemiology. 2013 January ; 24(1): 23–31. doi:10.1097/EDE.0b013e3182770237.

The Effect of Air Pollution Control on Life Expectancy in the United States: An Analysis of 545 US counties for the period 2000 to 2007

Andrew W. Correia,
Department of Biostatistics, Harvard School of Public Health, 655 Huntington Avenue, HSPH Building 2, 4th Floor, Boston, MA 02115

C. Arden Pope III,
Department of Economics, Brigham Young University, 142 Faculty Office Building, Provo, UT 84602

Douglas W. Dockery,
Departments of Environmental Health and Epidemiology, Harvard School of Public Health, 655 Huntington Avenue, HSPH Building 1, 1301B, Boston, MA 02115

Yun Wang,
Department of Biostatistics, Harvard School of Public Health, 655 Huntington Avenue, HSPH Building 2, 4th Floor, Boston, MA 02115

Majid Ezzati, and
MRC-HPA Centre for Environment and Health and Department of Epidemiology and Biostatistics, Imperial College London, Norfolk Place, St Mary's Campus, London W2 1PG

Francesca Dominici
Department of Biostatistics, Harvard School of Public Health, 655 Huntington Avenue, HSPH Building 2, 4th Floor, Boston, MA 02115, fdominic@hsph.harvard.edu, P: (617) 432-1056; F: (617)-739-1781

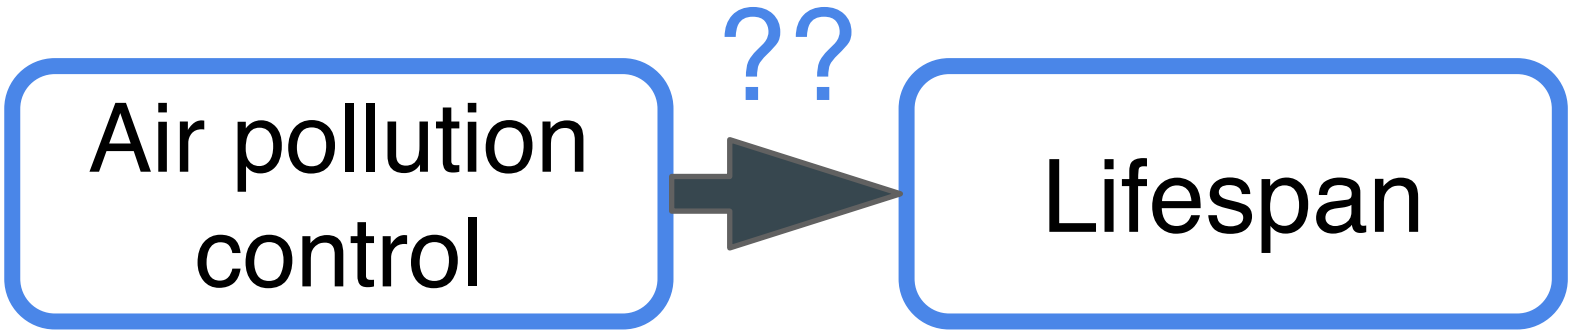

Is there a relationship between air pollution control and lifespan?

A decrease of 10 μg/m3 in the concentration of PM2.5 was associated with an increase in mean life expectancy of 0.35 years SD= 0.16 years, p = 0.033). This association was stronger in more urban and densely populated counties.

What if we want to know the effect of air pollution on everyone in the United States?

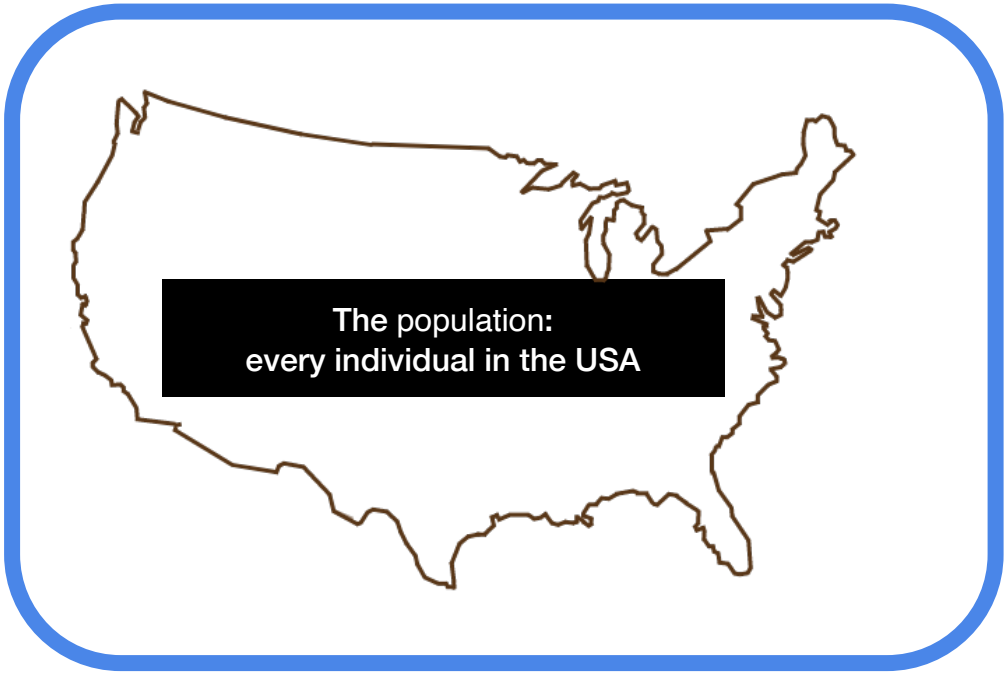

The population:
every individual in the USA

The sample:
545 US counties

#	State	Total Number of Counties	Total Area
1	Texas	254	268,596 mi²
2	Georgia	159	59,425 mi²
3	Virginia	133	42,775 mi²
4	Kentucky	120	40,408 mi²
5	Missouri	115	69,707 mi²
6	Kansas	105	82,278 mi²
7	Illinois	102	57,914 mi²
8	North Carolina	100	53,819 mi²
9	Iowa	99	56,273 mi²
10	Tennessee	95	42,144 mi²

3143 counties in the US

CA is #27

58 counties on 163k mi²

Random Sampler

How would you want to select those 545 counties?

What criteria are ideal?

What do you think they did?

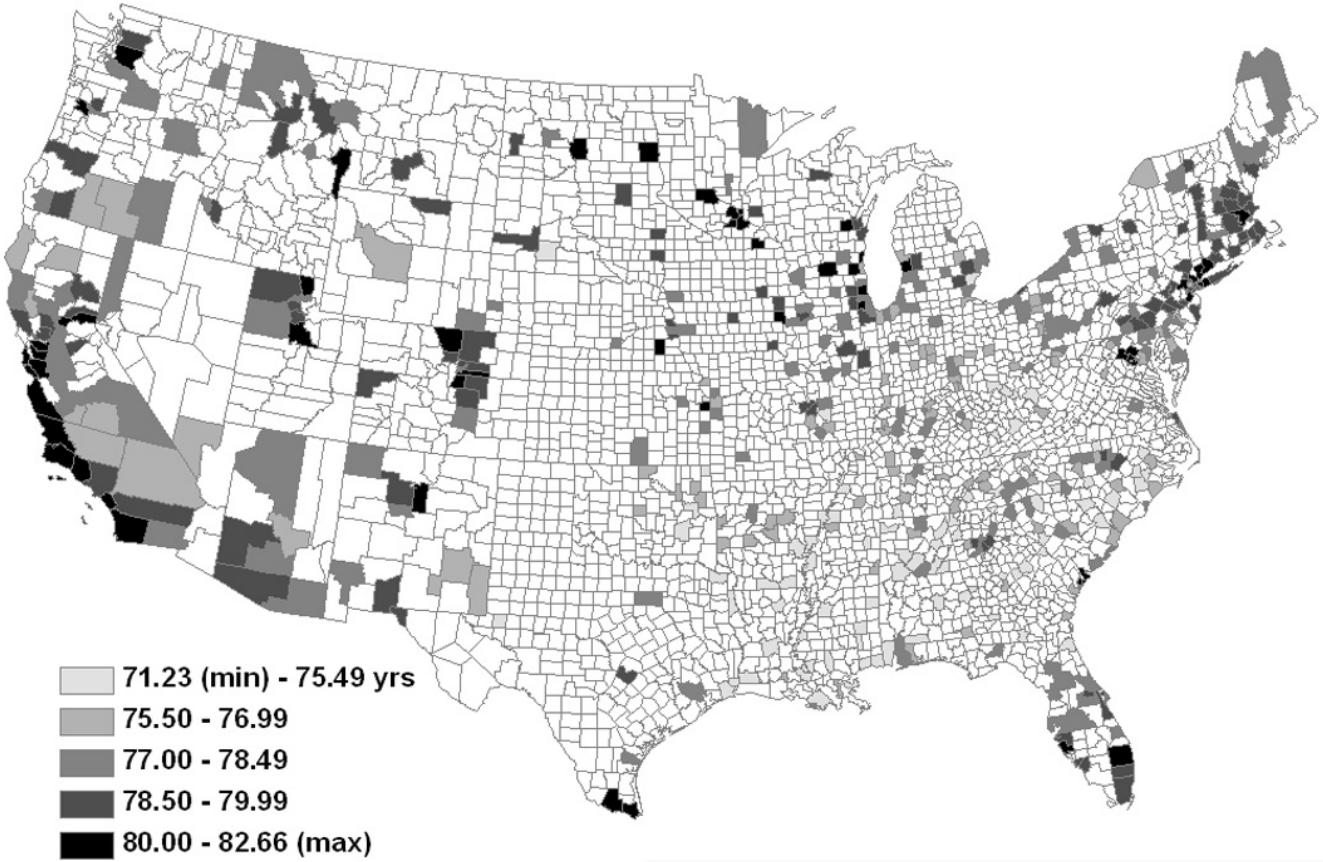

All counties with with available matching PM2.5 data for 2000 and 2007 from the EPA's Air Quality System. Includes both metropolitan and non-metro counties

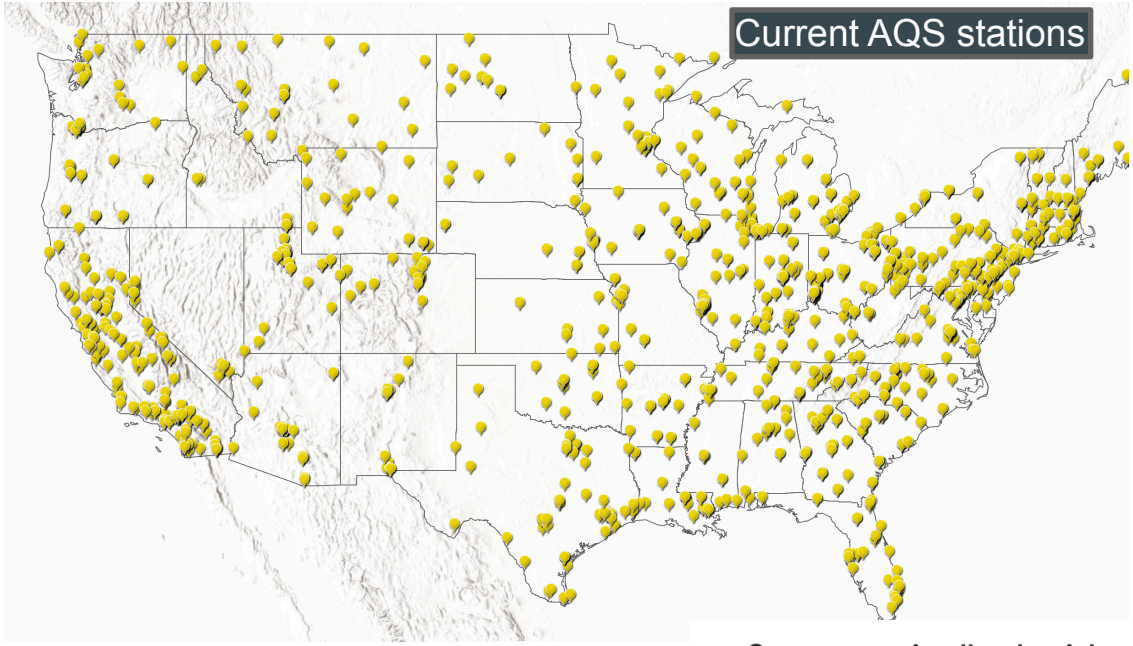

States typically decide where monitors are placed based on areas of relatively high population and/or areas believed to have relatively higher pollutant concentrations. Each state is responsible for developing its own monitoring plan, which is then reviewed and revised every five years. Aug 28, 2023

United States Environmental Protection Agency (.gov)
https://www.epa.gov › outdoor-air-quality-data › who-de...

Who decides where monitors get placed? | US EPA

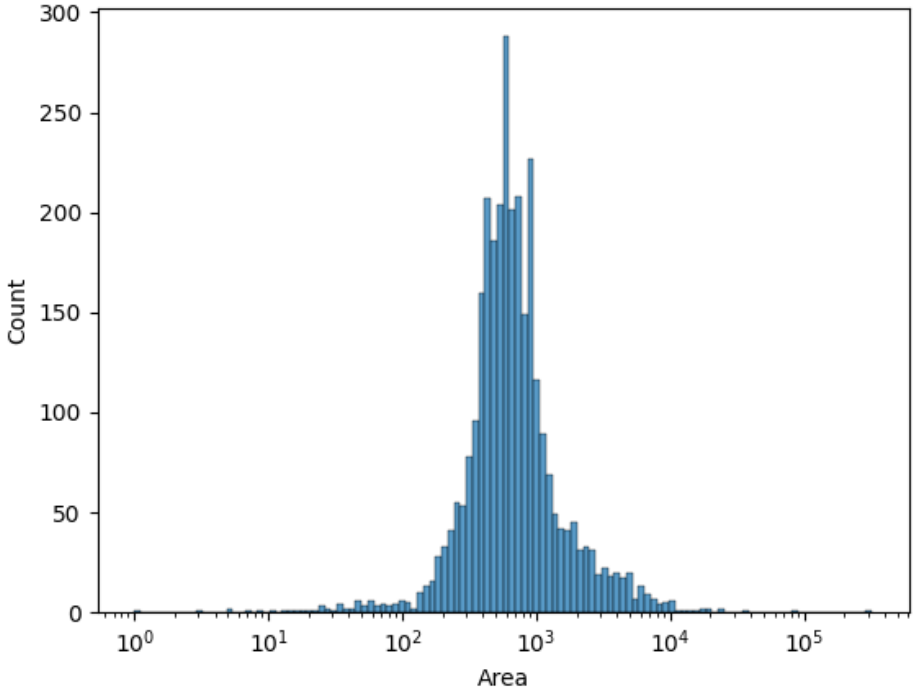
300
250
200
150
100
50
0
Count
10^0
10^1
10^2
10^3
10^4
10^5
Area

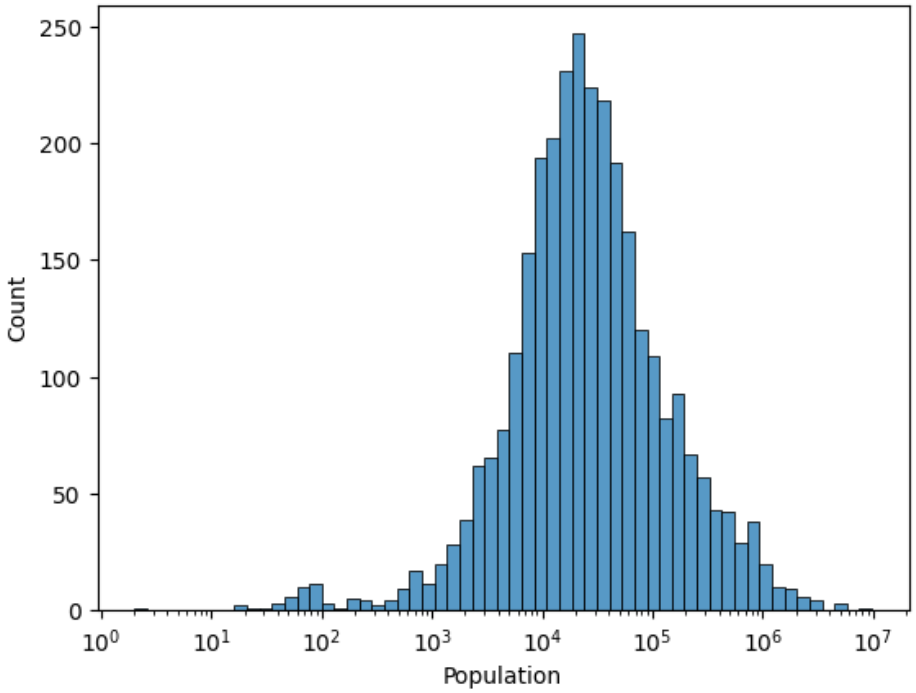
250
200
150
100
50
0
Count
10^0
10^1
10^2
10^3
10^4
10^5
10^6
10^7
Population

Random Sampler

How might conclusions differ if…

the sample locations were instead completely spatially at random across the USA?

The sample came from the 1980s or the 2020s instead of the 2000s?

Approaches to Inference

CORRELATION **COMPARISON OF MEANS** **REGRESSION** **NON-PARAMETRIC TESTS**

CORRELATION

ASSOCIATION BETWEEN VARIABLES

i.e. Pearson Correlation, Spearman Correlation, chi-square test

COMPARISON OF MEANS

DIFFERENCE IN MEANS BETWEEN VARIABLES

i.e. t-test, ANOVA

REGRESSION

DOES CHANGE IN ONE VARIABLE MEAN CHANGE IN ANOTHER?

I.e. simple regression, multiple regression

NON-PARAMETRIC TESTS

FOR WHEN ASSUMPTIONS IN THESE OTHER 3 CATEGORIES ARE NOT MET

i.e. Wilcoxon rank-sum test, Wilcoxon sign-rank test, sign test

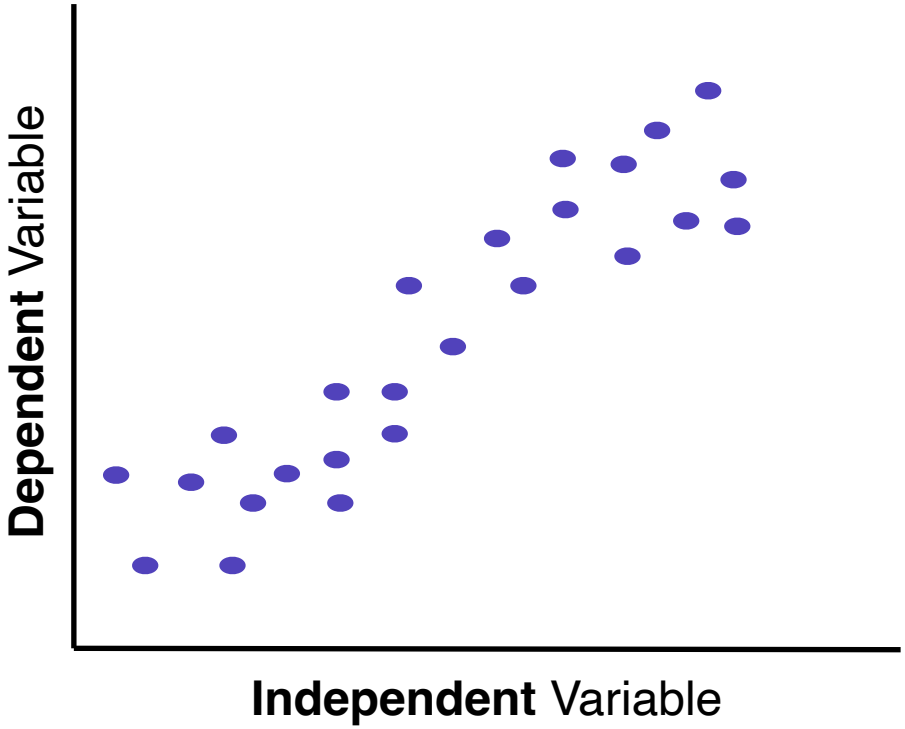
Dependent Variable
Independent Variable

weaker relationship

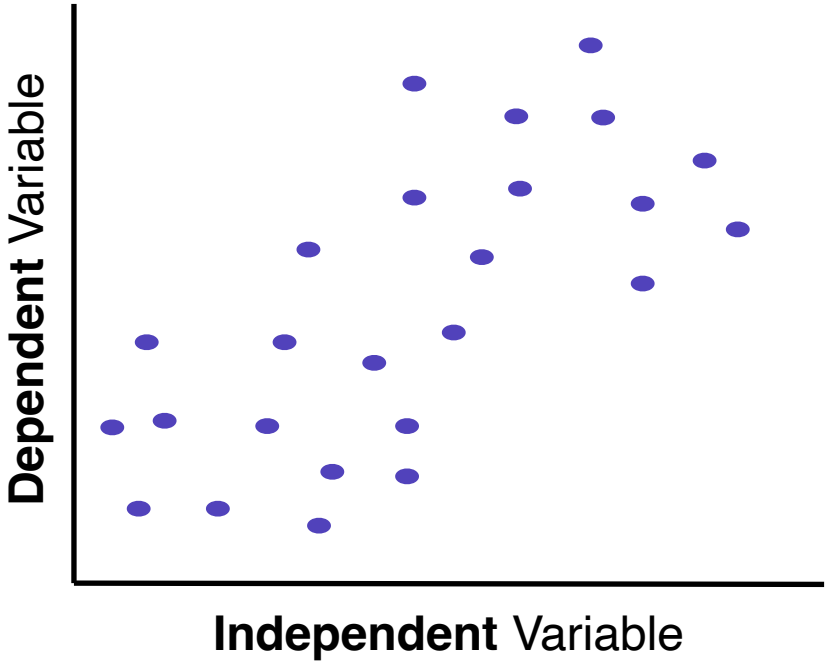

stronger relationship

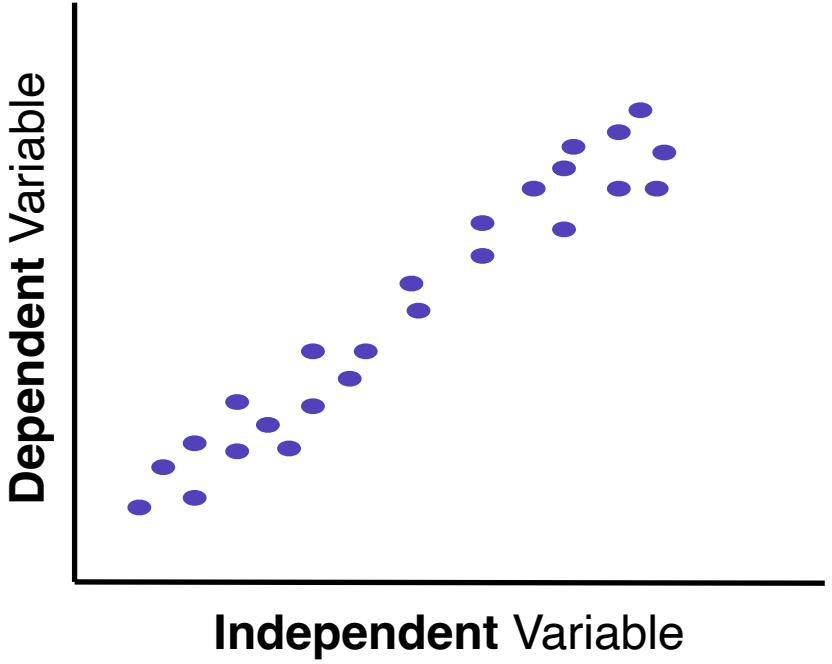

stronger relationship = higher correlation

The *larger* the independent variable value, the *larger* the dependent variable tends to be

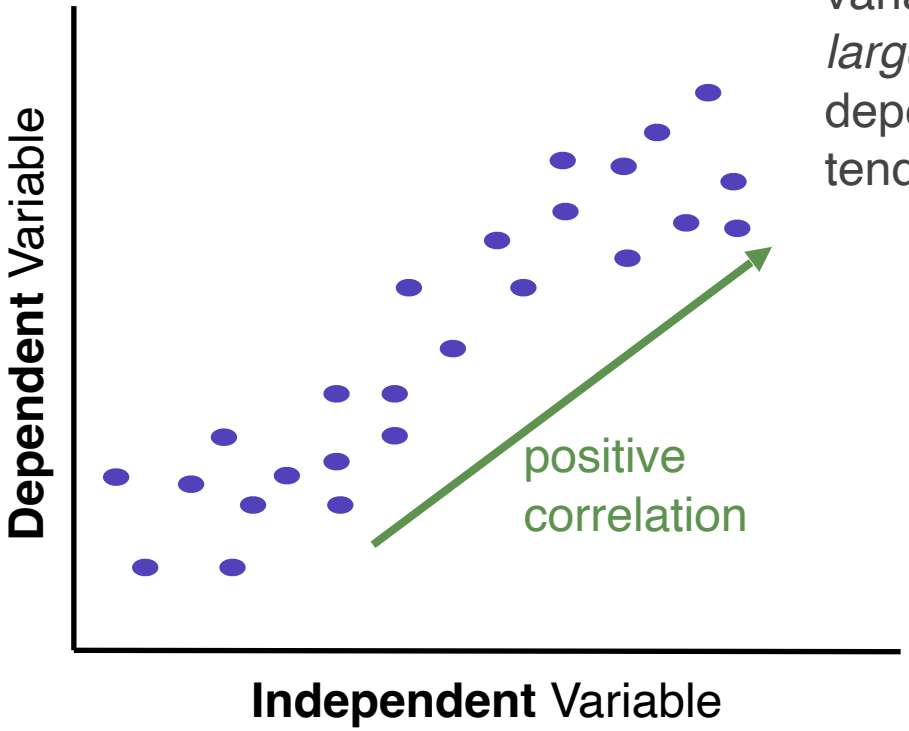

The *smaller* the independent variable value, the *smaller* the dependent variable tends to be

The *taller* the father, the taller his son tends to be

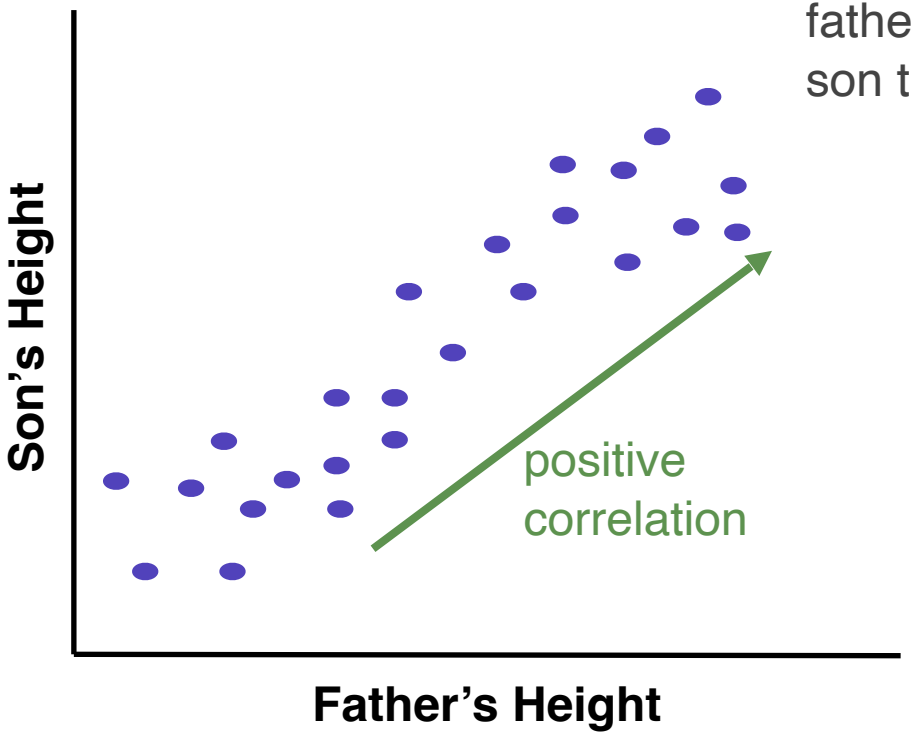

The *shorter* the father, the shorter his son tends to be

The *smaller* the independent variable value, the *larger* the dependent variable tends to be

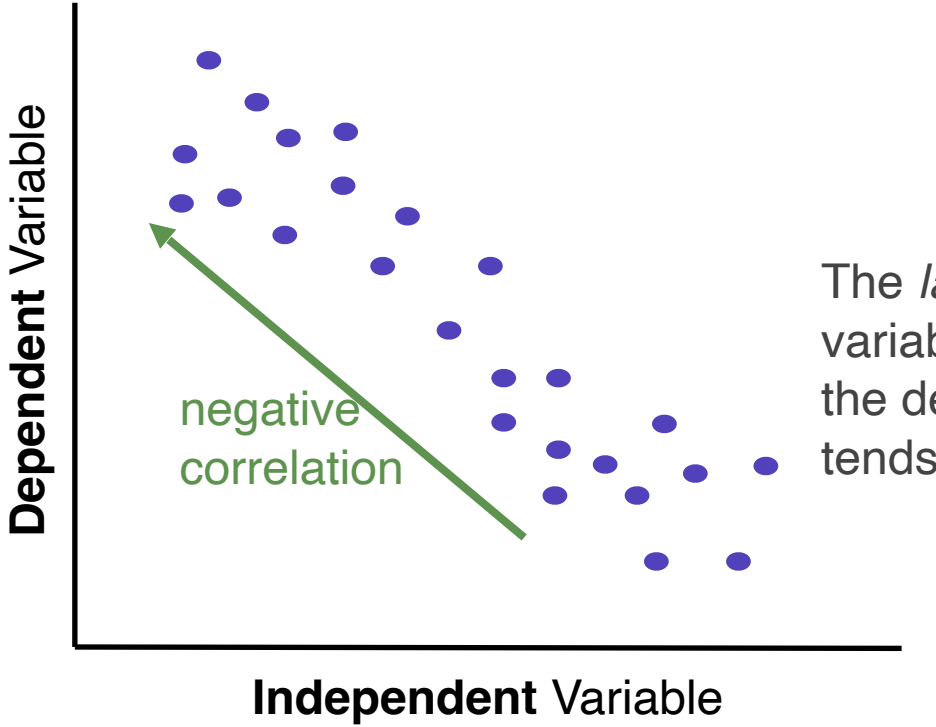

The *larger* the independent variable value, the *smaller* the dependent variable tends to be

The *lower* the number of absences, the *higher* the students' grades tend to be

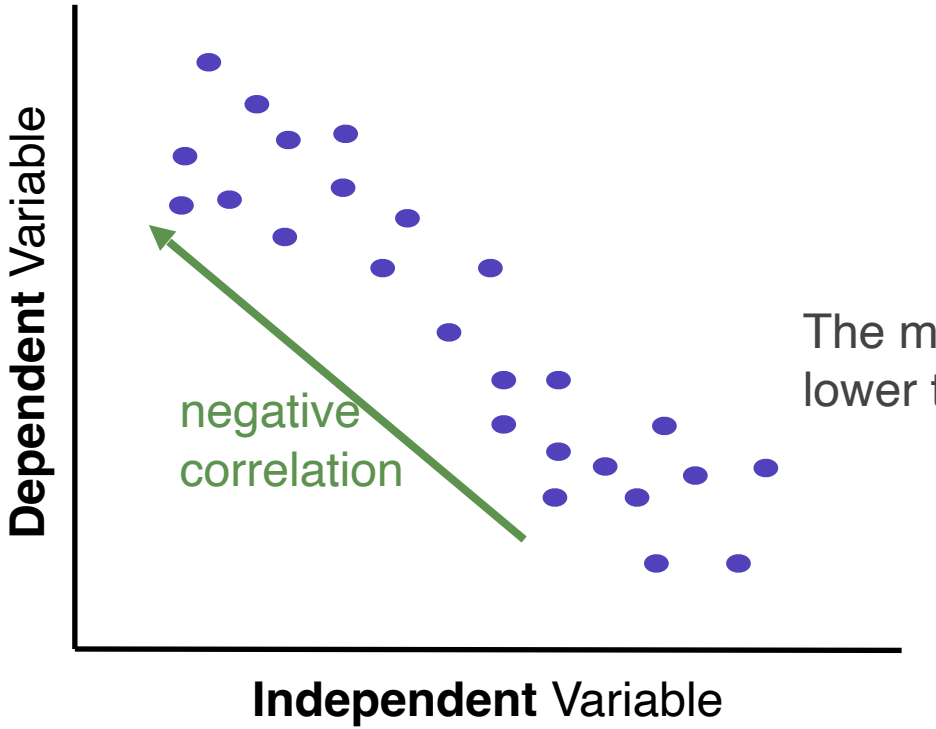

The more absences, the lower the students' grades

Pearson's *r* :

linear correlation between two variables

takes values [-1,1]

Correlation is how close the data are to being in a line…
BUT IT HAS NOTHING TO DO WITH THE SLOPE

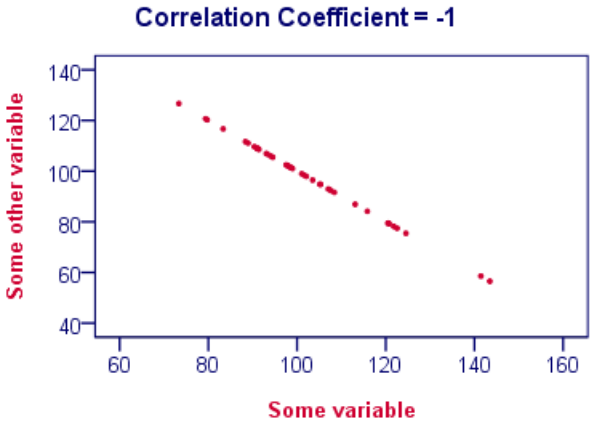

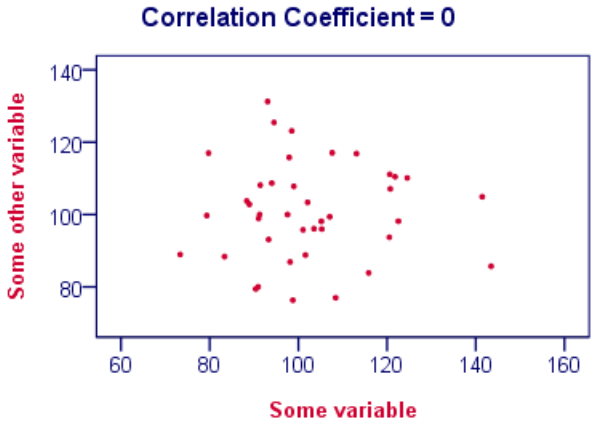

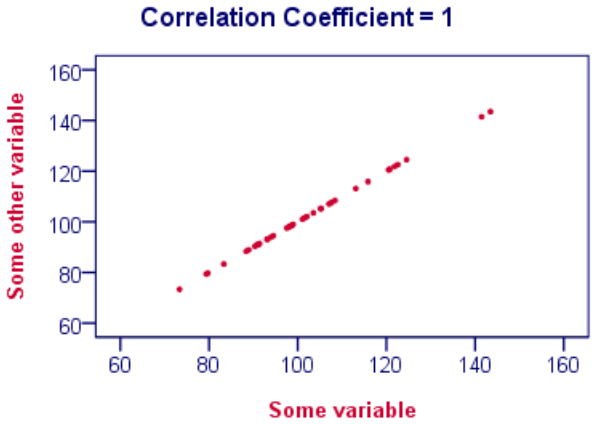

https://www.spss-tutorials.com/pearson-correlation-coefficient/

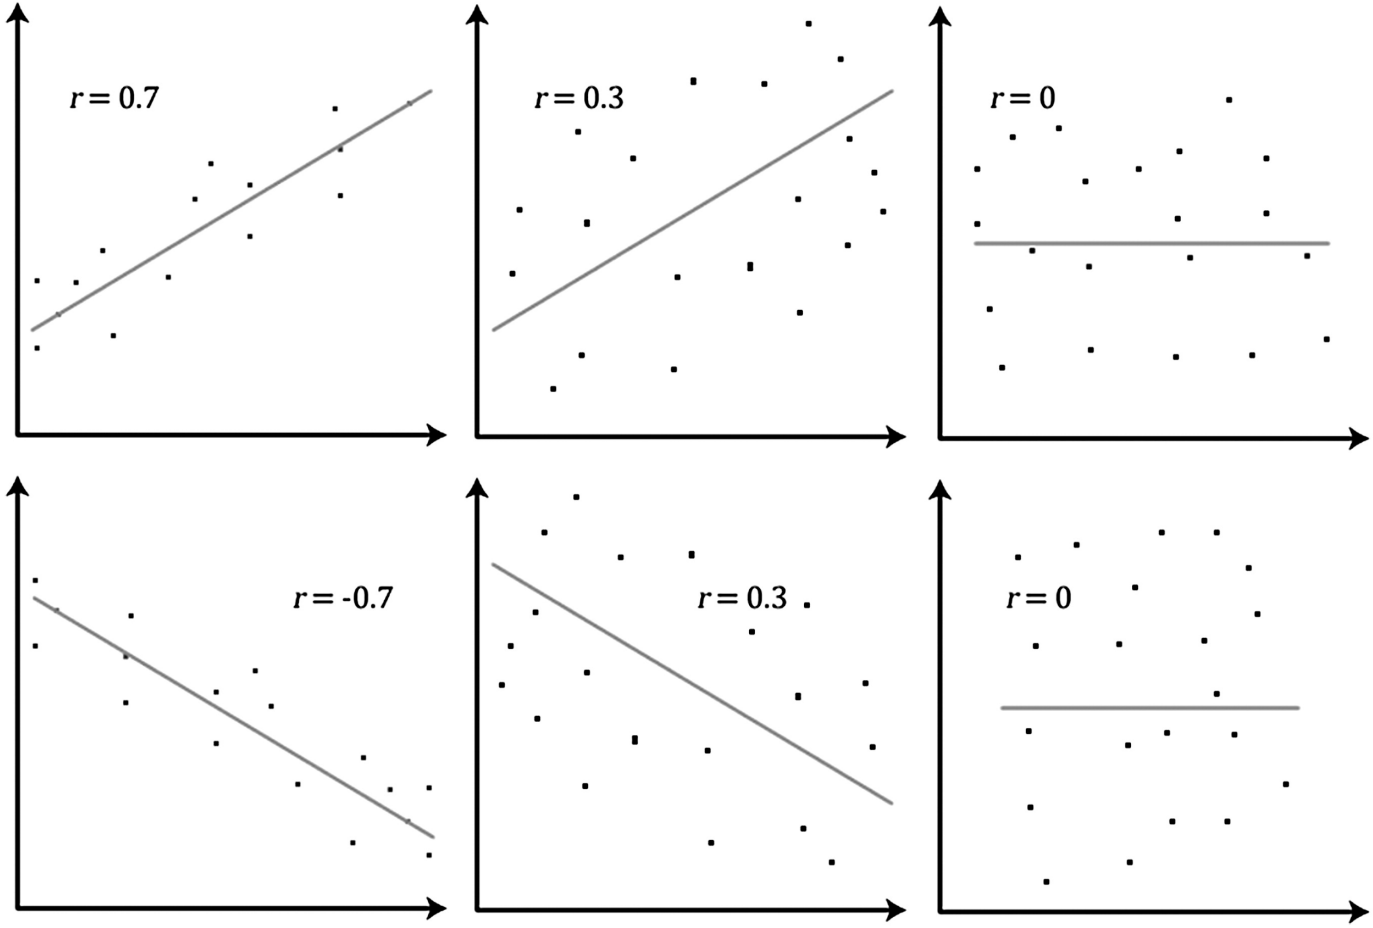

https://www.guessthecorrelation.com/

Which of the following is the Pearson correlation coefficient (*r*) for this relationship?

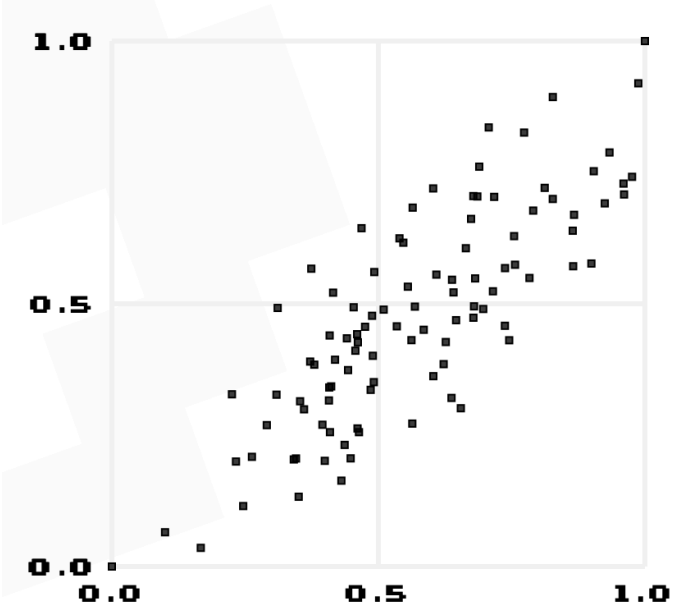

https://forms.gle/M4Q3hFkeoLKCixUX8

Correlation != Causation

Correlation establishes a relationship.

It does NOT establish causation.

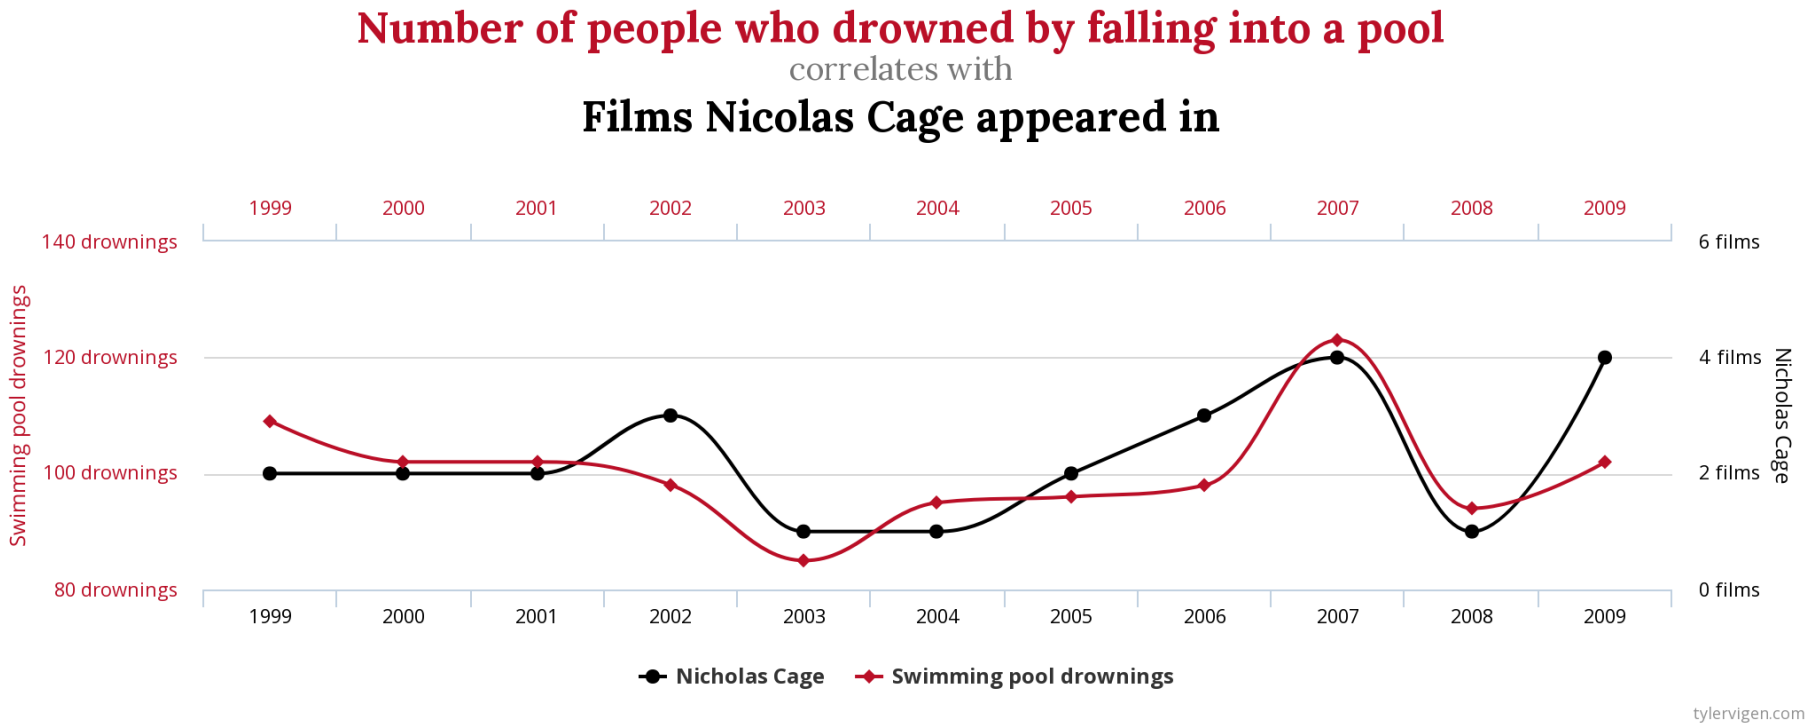

http://www.tylervigen.com/spurious-correlations

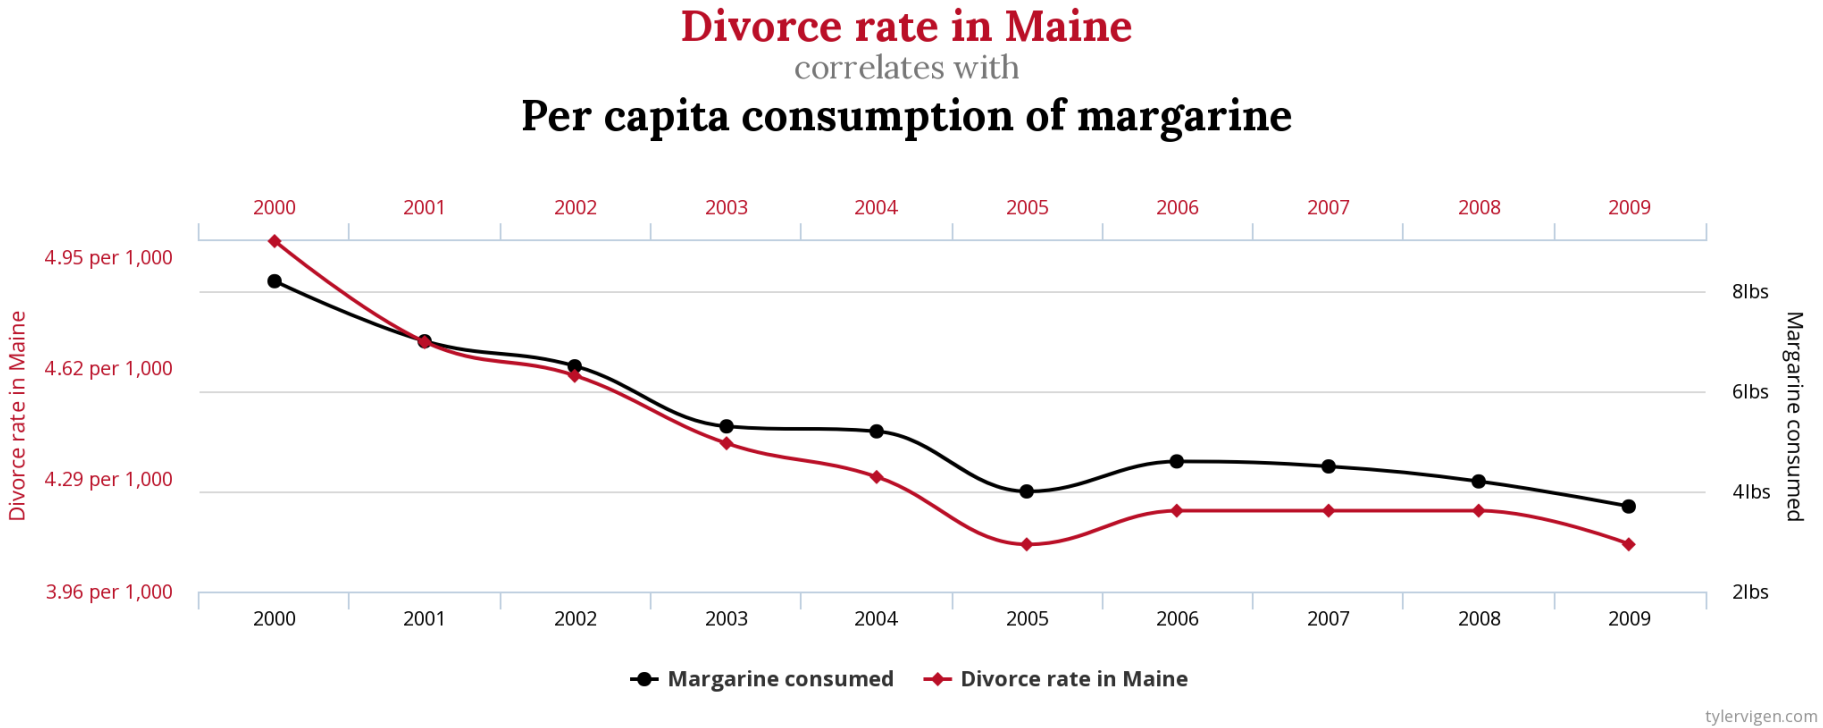

http://www.tylervigen.com/spurious-correlations

What about correlation between categorical variables?

- Make a contingency table …

Handed-ness / Sex	Right-handed	Left-handed	Total
Male	43	9	52
Female	44	4	48
Total	87	13	100

- Then calculate phi (based on chi-squared statistic) …

Phi coefficient [edit]

Main article: Phi coefficient

A simple measure, applicable only to the case of 2 × 2 contingency tables, is the phi coefficient (φ) defined by

$$\phi = \pm\sqrt{\frac{\chi^2}{N}},$$

where χ^2 is computed as in Pearson's chi-squared test, and N is the grand total of observations. φ varies from 0 (corresponding to no association between the variables) to 1 or −1 (complete association or complete inverse association), provided it is based on frequency data represented in 2 × 2 tables. Then its sign equals the sign of the product of the main diagonal elements of the table minus the product of the off–diagonal elements. φ takes on the minimum value −1.0 or the maximum value of +1.0 if and only if every marginal proportion is equal to 0.5 (and two diagonal cells are empty).[2]

- Yes there are ways to do this in pandas/numpy and other libraries!

CORRELATION

ASSOCIATION BETWEEN VARIABLES

i.e. Pearson Correlation, Spearman Correlation, chi-square test

COMPARISON OF MEANS

DIFFERENCE IN MEANS BETWEEN CONDITIONS

i.e. t-test, ANOVA

REGRESSION

DOES CHANGE IN ONE VARIABLE MEAN CHANGE IN ANOTHER?

I.e. simple regression, multiple regression

NON-PARAMETRIC TESTS

FOR WHEN ASSUMPTIONS IN THESE OTHER 3 CATEGORIES ARE NOT MET

i.e. Wilcoxon rank-sum test, Wilcoxon sign-rank test, sign test

t-test:

tests for difference in means between groups

William Sealy Gosset (13 June 1876 – 16 October 1937) was an English statistician, chemist and brewer who served as Head Brewer of Guinness and Head Experimental Brewer of Guinness and was a pioneer of modern statistics. He pioneered small sample experimental design and analysis with an economic approach to the logic of uncertainty. Gosset published under the pen name **Student** and developed most famously Student's t-distribution – originally called Student's "z" – and "Student's test of statistical significance".[1]

Contents [hide]

Life and career [edit]

Born in Canterbury, England the eldest son of Agnes Sealy Vidal and Colonel Frederic Gosset, R.E. Royal Engineers, Gosset attended Winchester College before matriculating as Winchester Scholar in natural sciences and mathematics at New College, Oxford. Upon graduating in 1899, he joined the brewery of Arthur Guinness & Son in Dublin, Ireland; he spent the rest of his 38-year career at Guinness.[1][2]

Gosset had three children with Marjory Gosset (née Phillpotts). Harry Gosset (1907–1965) was a consultant paediatrician; Bertha Marian Gosset (1909–2004) was a geographer and nurse; the youngest, Ruth Gosset (1911–1953) married the Oxford mathematician Douglas Roaf and had five children.

In his job as Head Experimental Brewer at Guinness, the self-trained Gosset developed new statistical methods – both in the brewery and on the farm – now central to the design of experiments, to proper use of significance testing on repeated trials, and to analysis of economic significance (an early instance of decision theory interpretation of statistics) and more, such as his small-sample, stratified, and repeated balanced experiments on barley for proving the best yielding varieties.[3] Gosset acquired that knowledge by study, by trial and error, by cooperating with others, and by spending two terms in 1906–1907 in the Biometrics laboratory of Karl Pearson.[4] Gosset and Pearson had a good relationship.[4] Pearson helped Gosset with the mathematics of his papers, including the 1908 papers, but had little appreciation of their importance. The papers addressed the brewer's concern with small samples; biometricians like Pearson, on the other hand, typically had hundreds of observations and saw no urgency in developing small-sample methods.[2]

Gosset's first publication came in 1907, "On the Error of Counting with a Haemacytometer," in which – unbeknownst to Gosset aka "Student" – he rediscovered the Poisson distribution.[3] Another researcher at Guinness had previously published a paper containing trade secrets of the Guinness

William Sealy Gosset

William Sealy Gosset (aka *Student*) in 1908 (age 32)

Born	13 June 1876 Canterbury, Kent, England
Died	16 October 1937 (aged 61) Beaconsfield, Buckinghamshire, England
Other names	Student
Alma mater	New College, Oxford, Winchester College
Known for	Student's t-distribution, statistical significance, design of experiments, Monte Carlo method, quality control, Modern synthesis, agricultural economics, econometrics
Children	5, including Isaac Henry Gosset
Scientific career	

Do the heights between males and females differ?

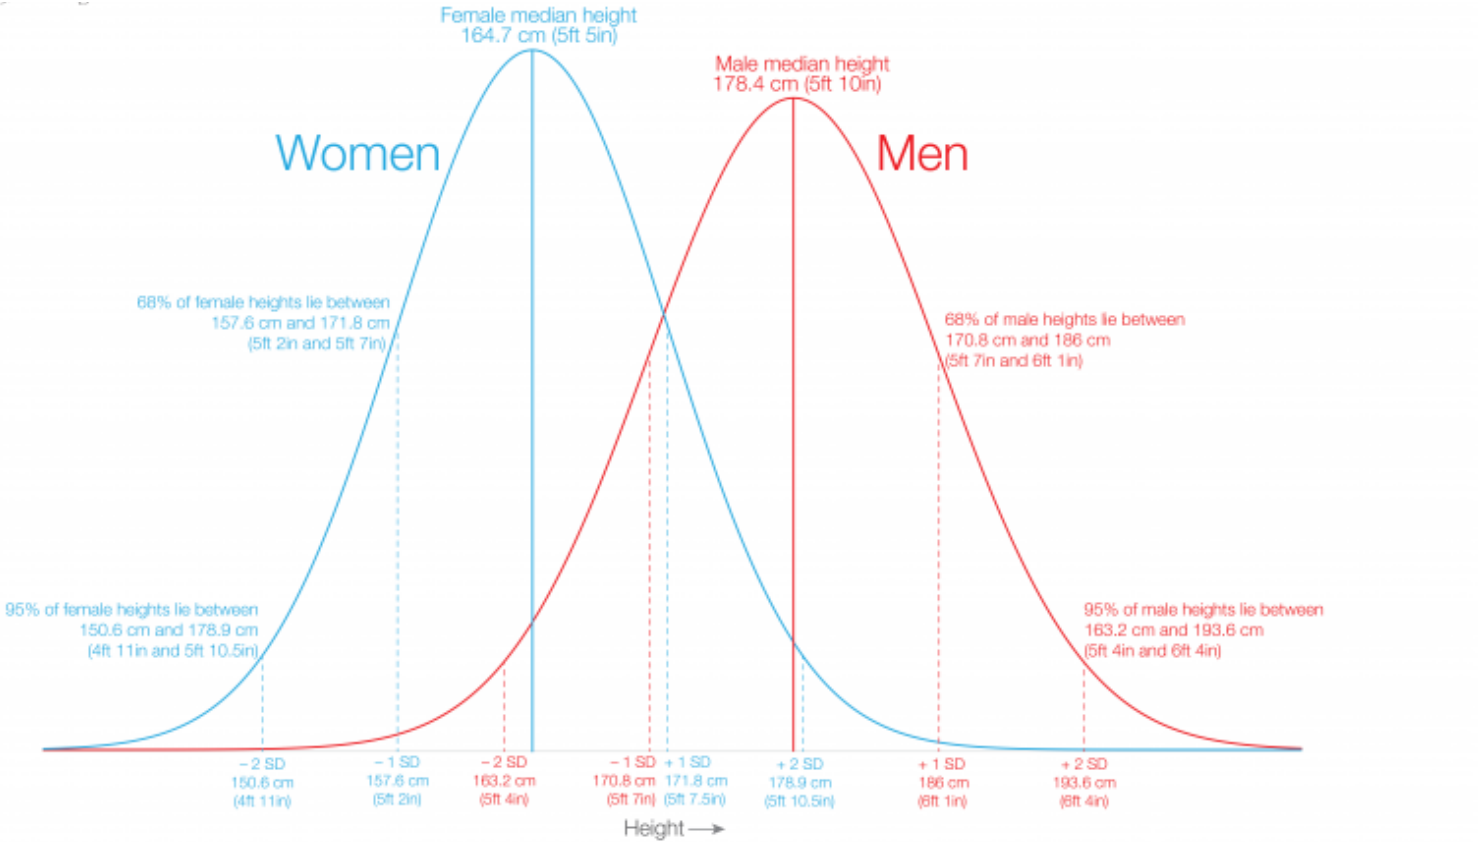

https://ourworldindata.org/human-height

t-test Assumptions

1. Data are continuous
2. Normally distributed
3. Equal variance b/w groups (but can use Welch's test!)
4. Not paired (will talk more about this later)

Do the heights between males and females differ?

variance is equal(ish) between groups (57 vs 50cm)

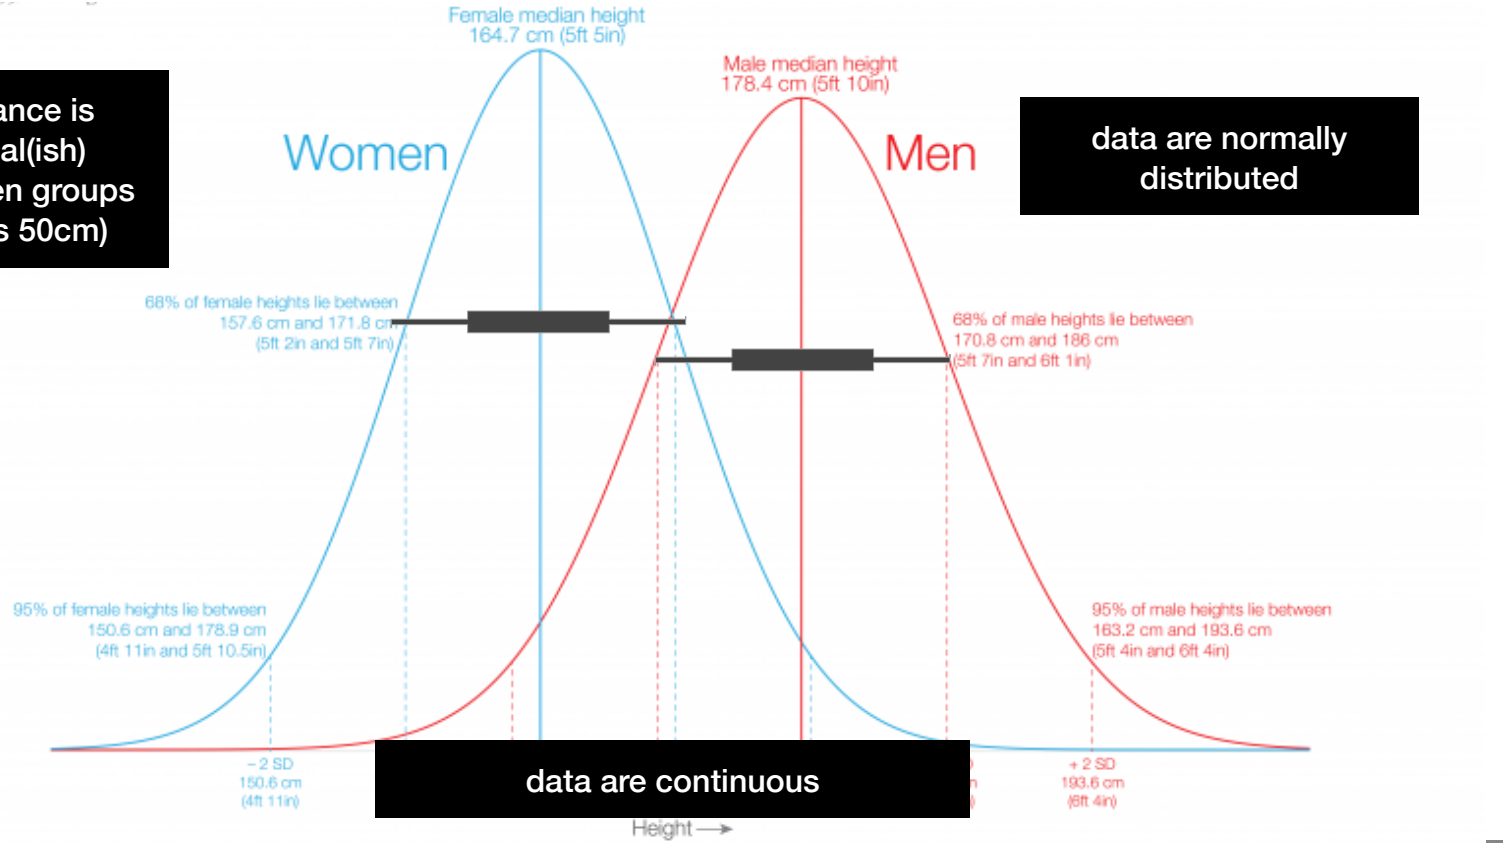

data are normally distributed

sample size affects statistic

N=10,000

data are continuous

Do the heights between males and females differ?

t-statistic: -95.6

p-value << 0.001

95% CI for true difference in means
[-5.43, -5.21]

Yes.

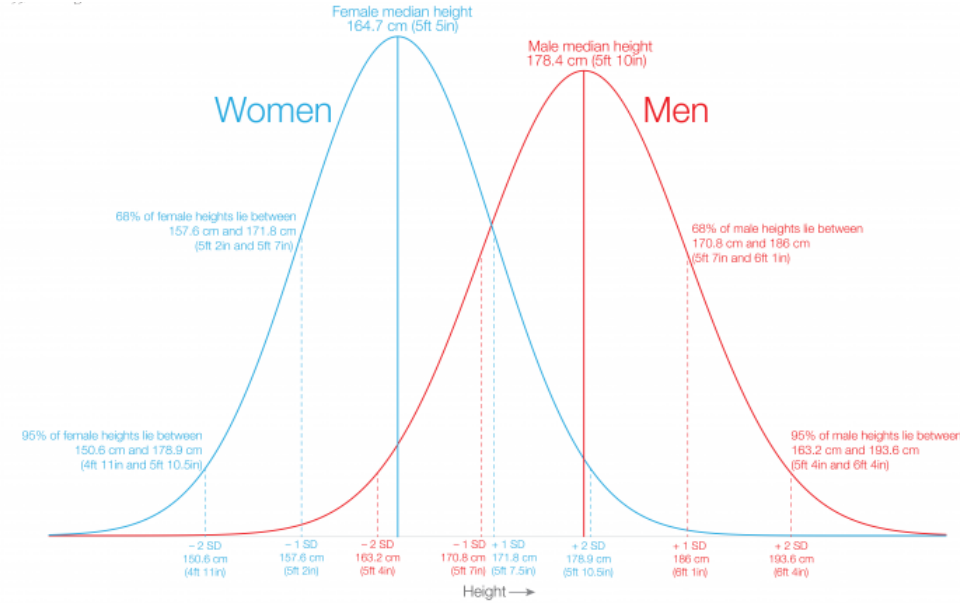

p-value : the probability under the null hypothesis of getting measurements as extreme as the observed results by chance alone

THIS IS NOT TYPE 1 ERROR RATE

Confidence interval : a range of values calculated from a sample statistic, such that there is a specified probability that the value of the true value of the population (parameter) lies within it.

https://forms.gle/6MCyp7qFsaHgGKi5A

If we flip a coin 16 times and record the number of heads....

….and then do that 1M times

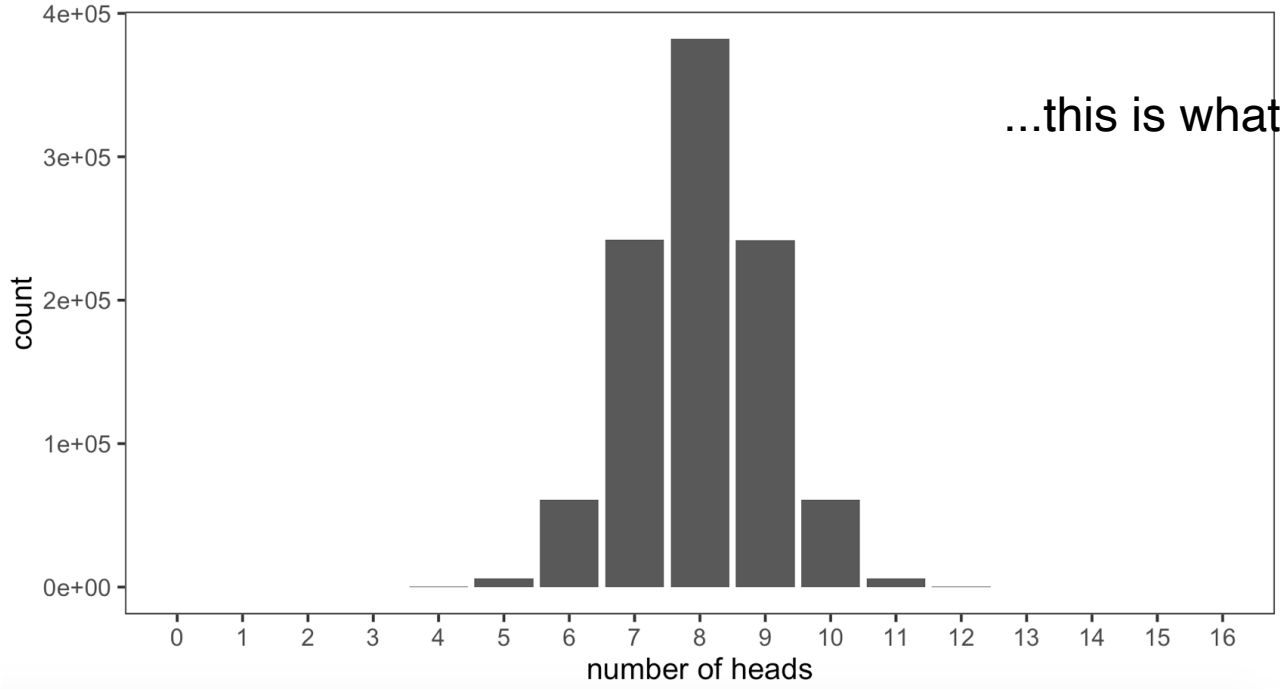

...this is what happens by chance alone.

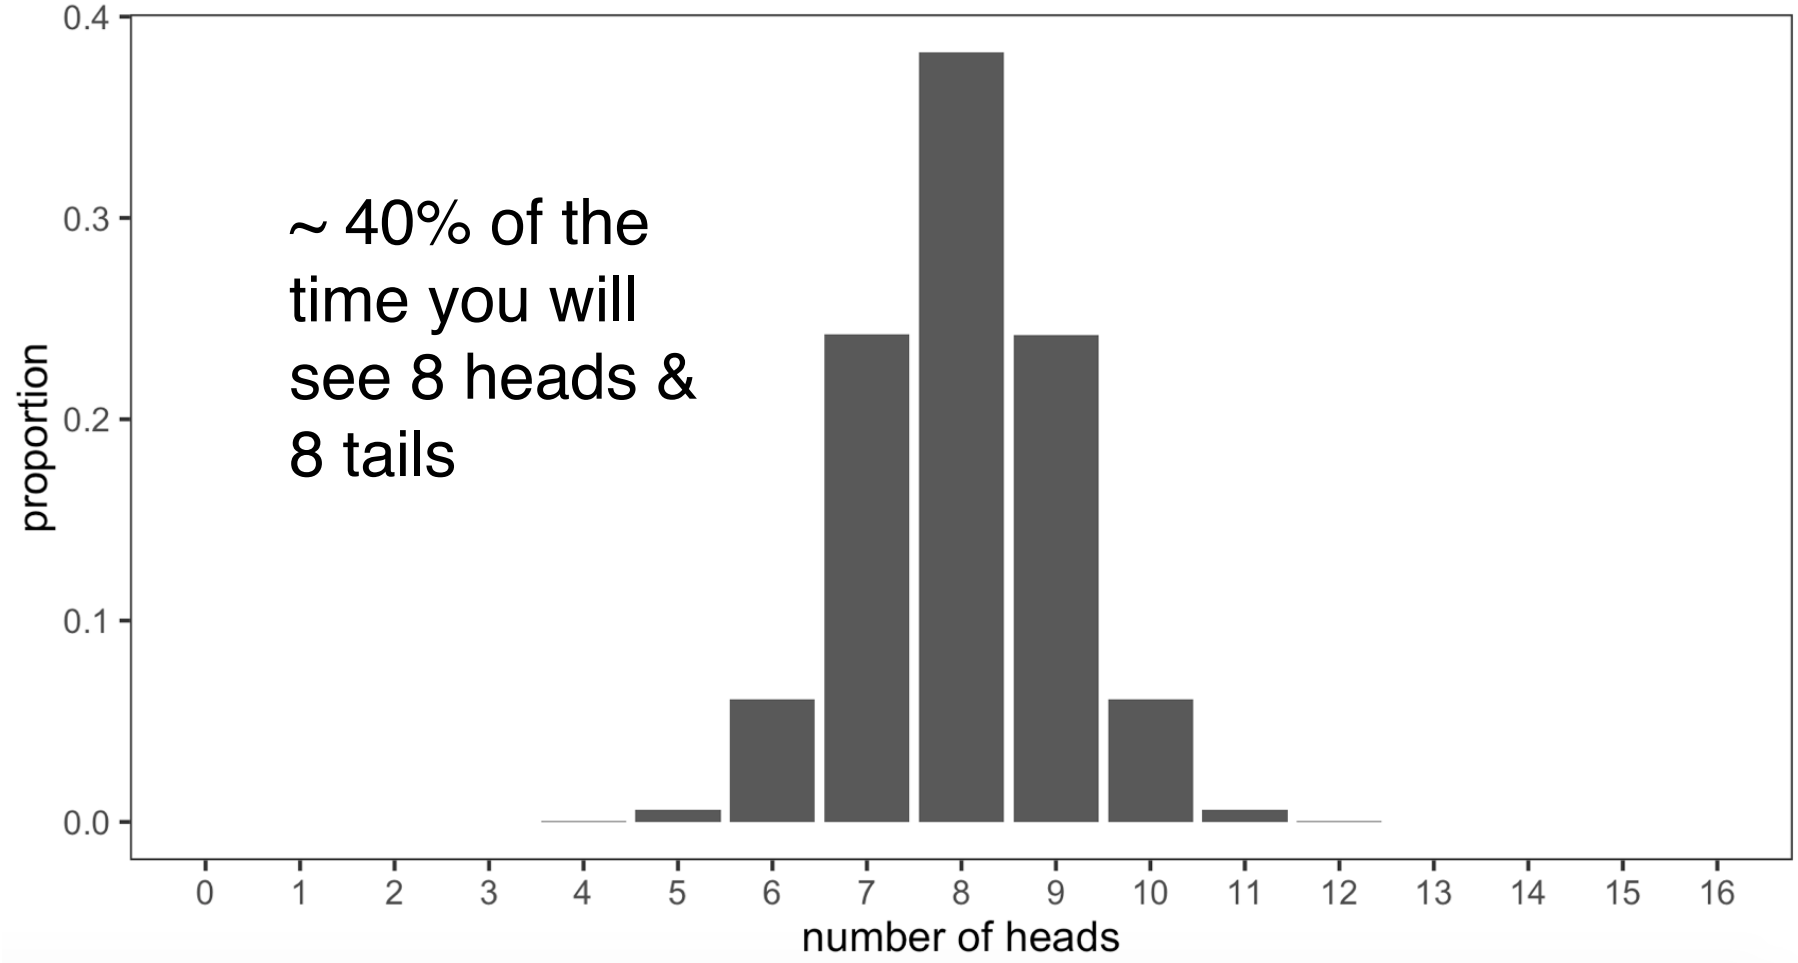
~ 40% of the time you will see 8 heads & 8 tails
proportion
0.4
0.3
0.2
0.1
0.0
0 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16
number of heads

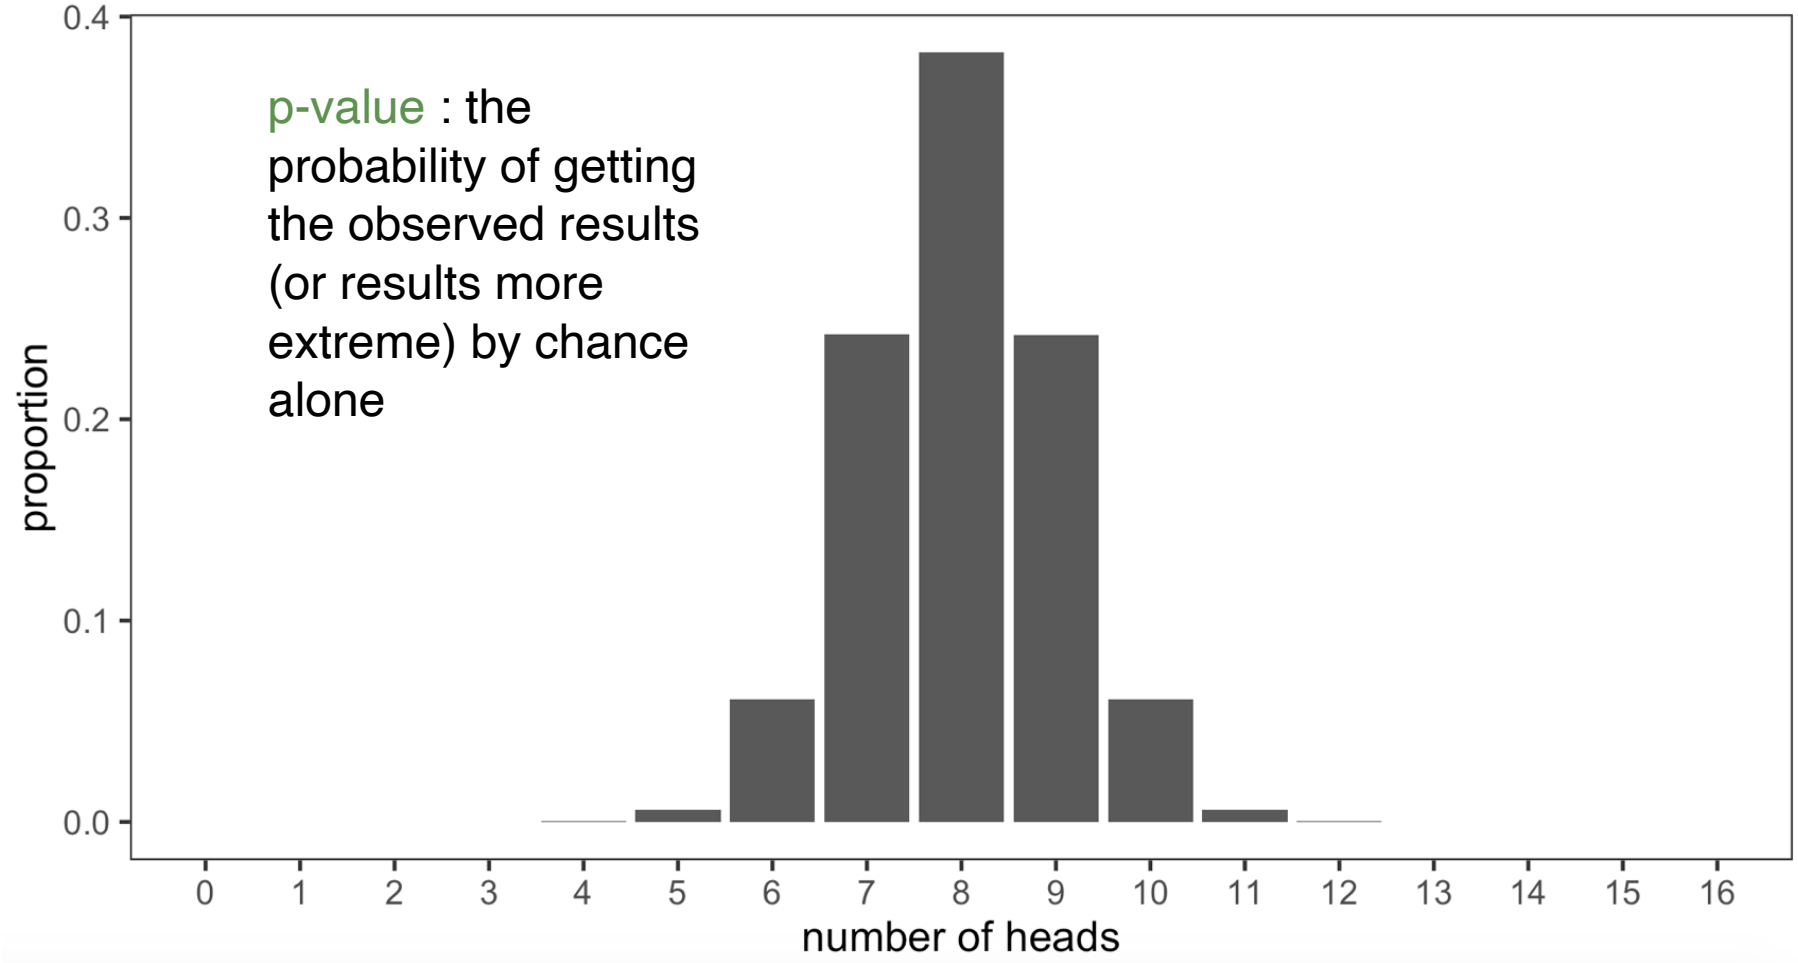
p-value : the probability of getting the observed results (or results more extreme) by chance alone
proportion
0.4
0.3
0.2
0.1
0.0
0
1
2
3
4
5
6
7
8
9
10
11
12
13
14
15
16
number of heads

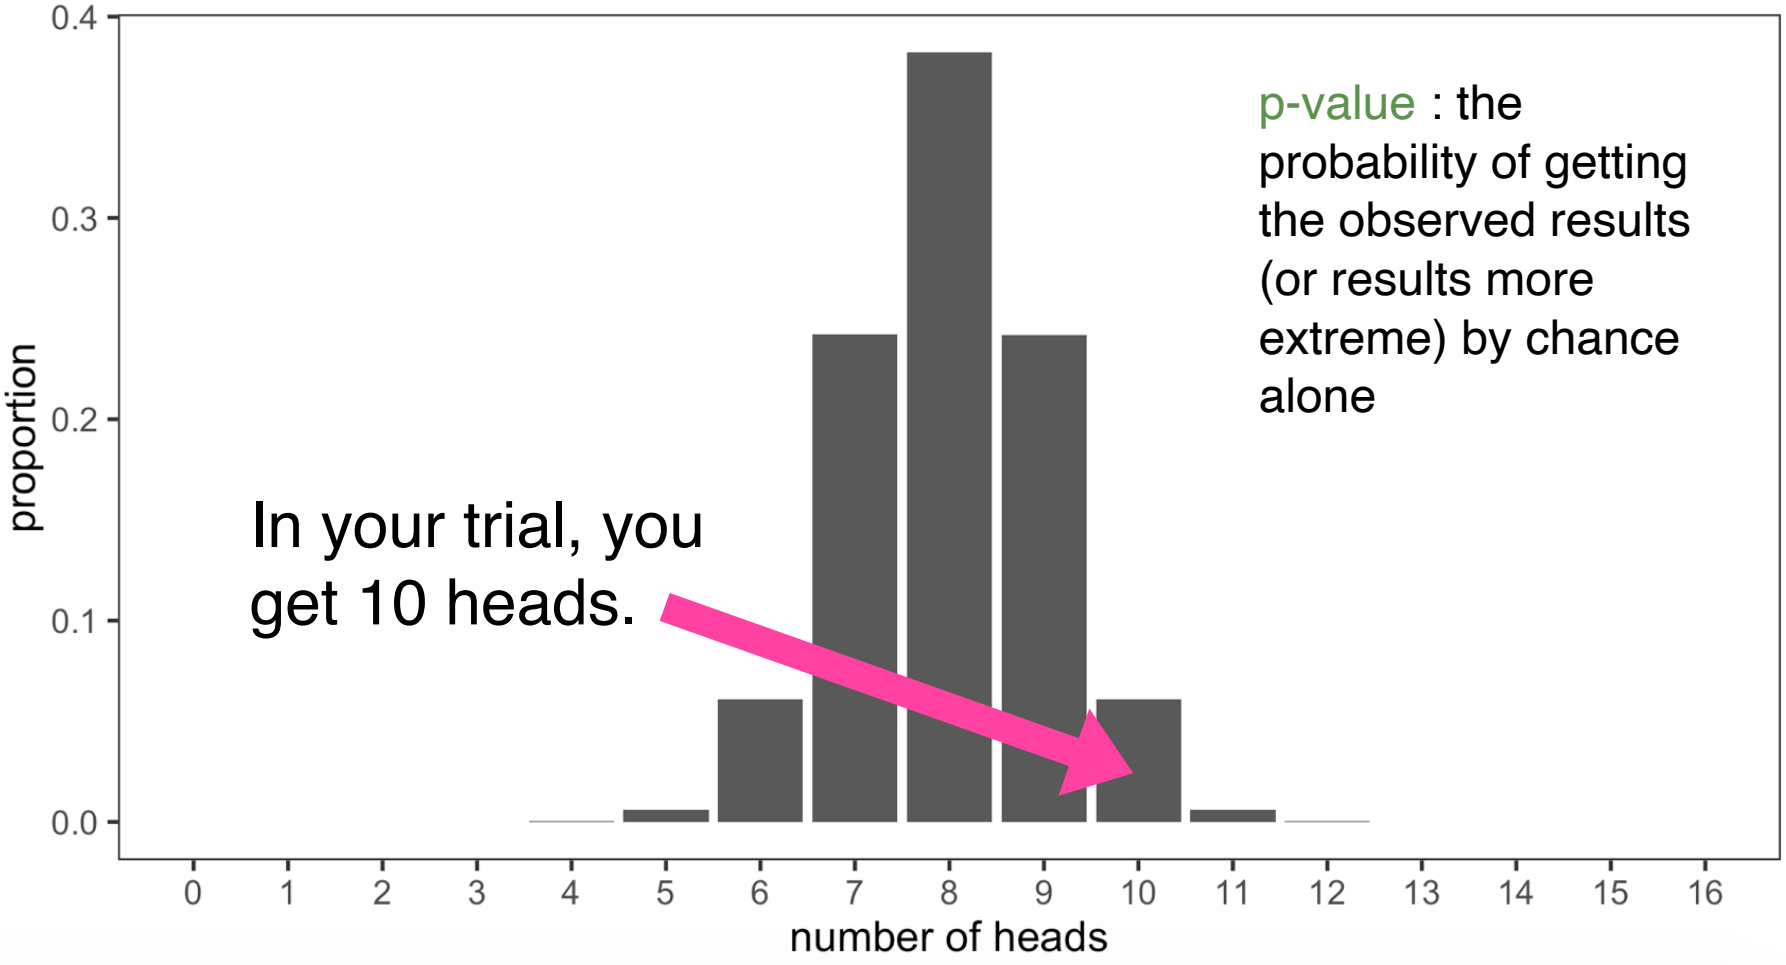
0.4
0.3
0.2
0.1
0.0
proportion
0
1
2
3
4
5
6
7
8
9
10
11
12
13
14
15
16
number of heads
In your trial, you get 10 heads.
p-value : the probability of getting the observed results (or results more extreme) by chance alone

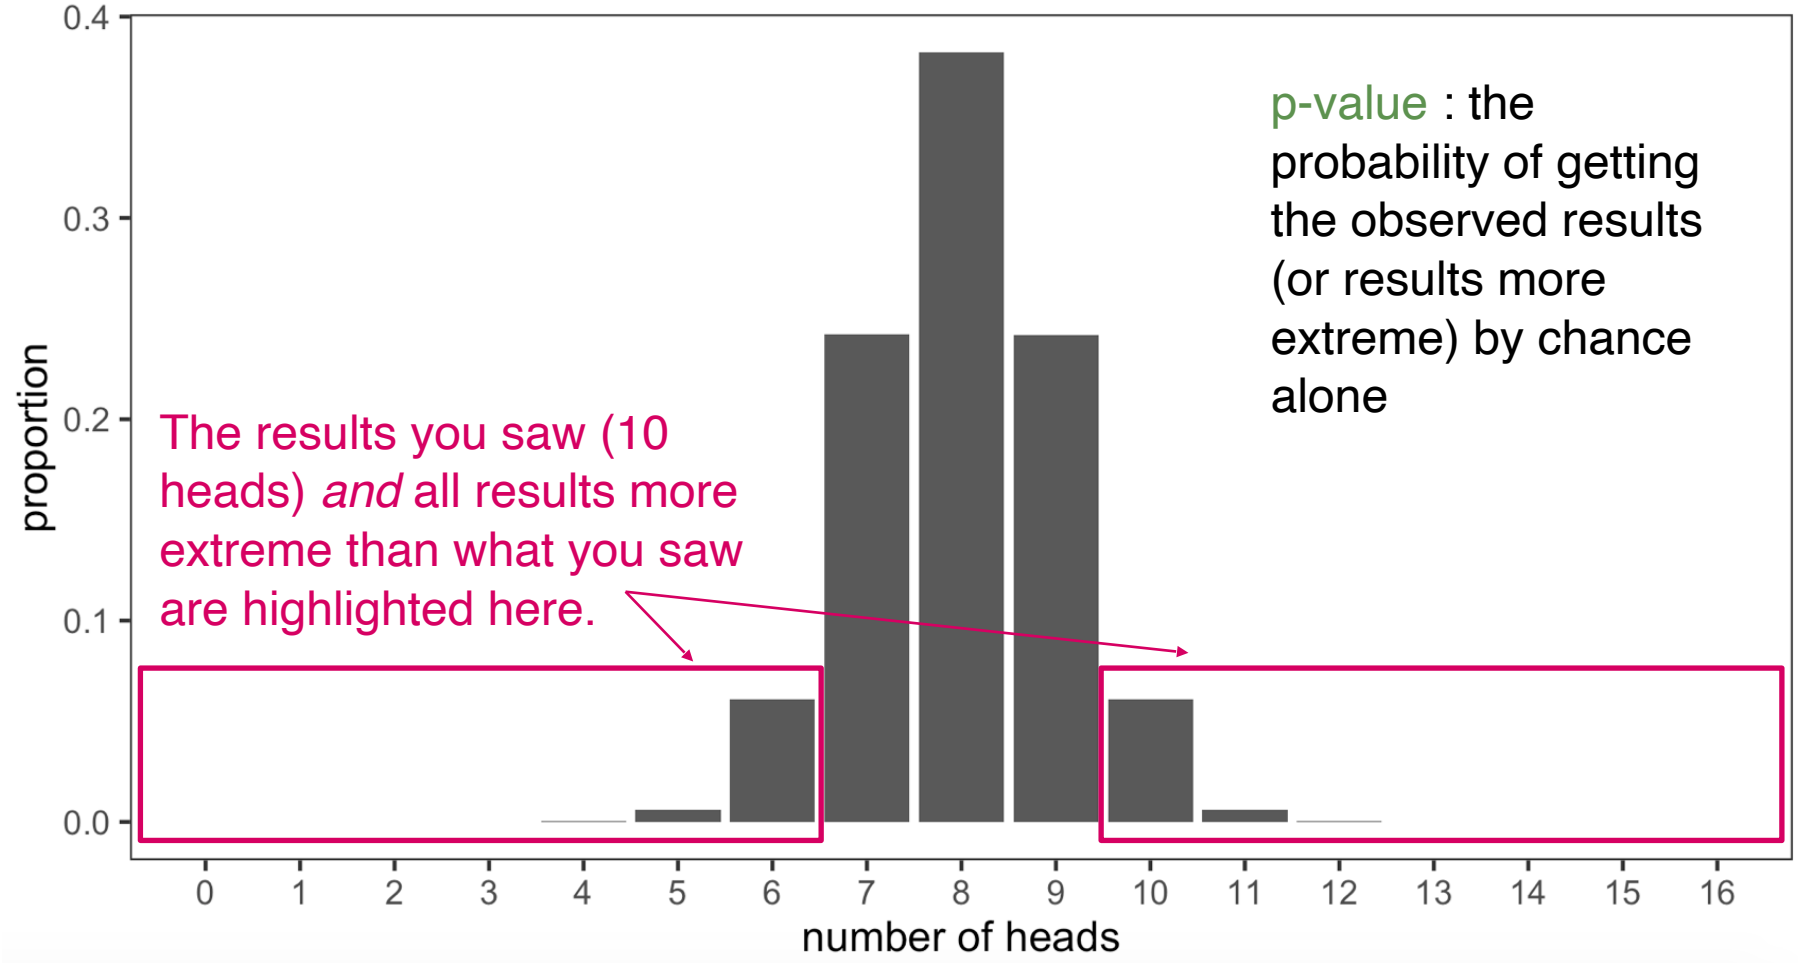
0.4
0.3
0.2
0.1
0.0
proportion
0
1
2
3
4
5
6
7
8
9
10
11
12
13
14
15
16
number of heads
p-value : the probability of getting the observed results (or results more extreme) by chance alone
The results you saw (10 heads) *and* all results more extreme than what you saw are highlighted here.

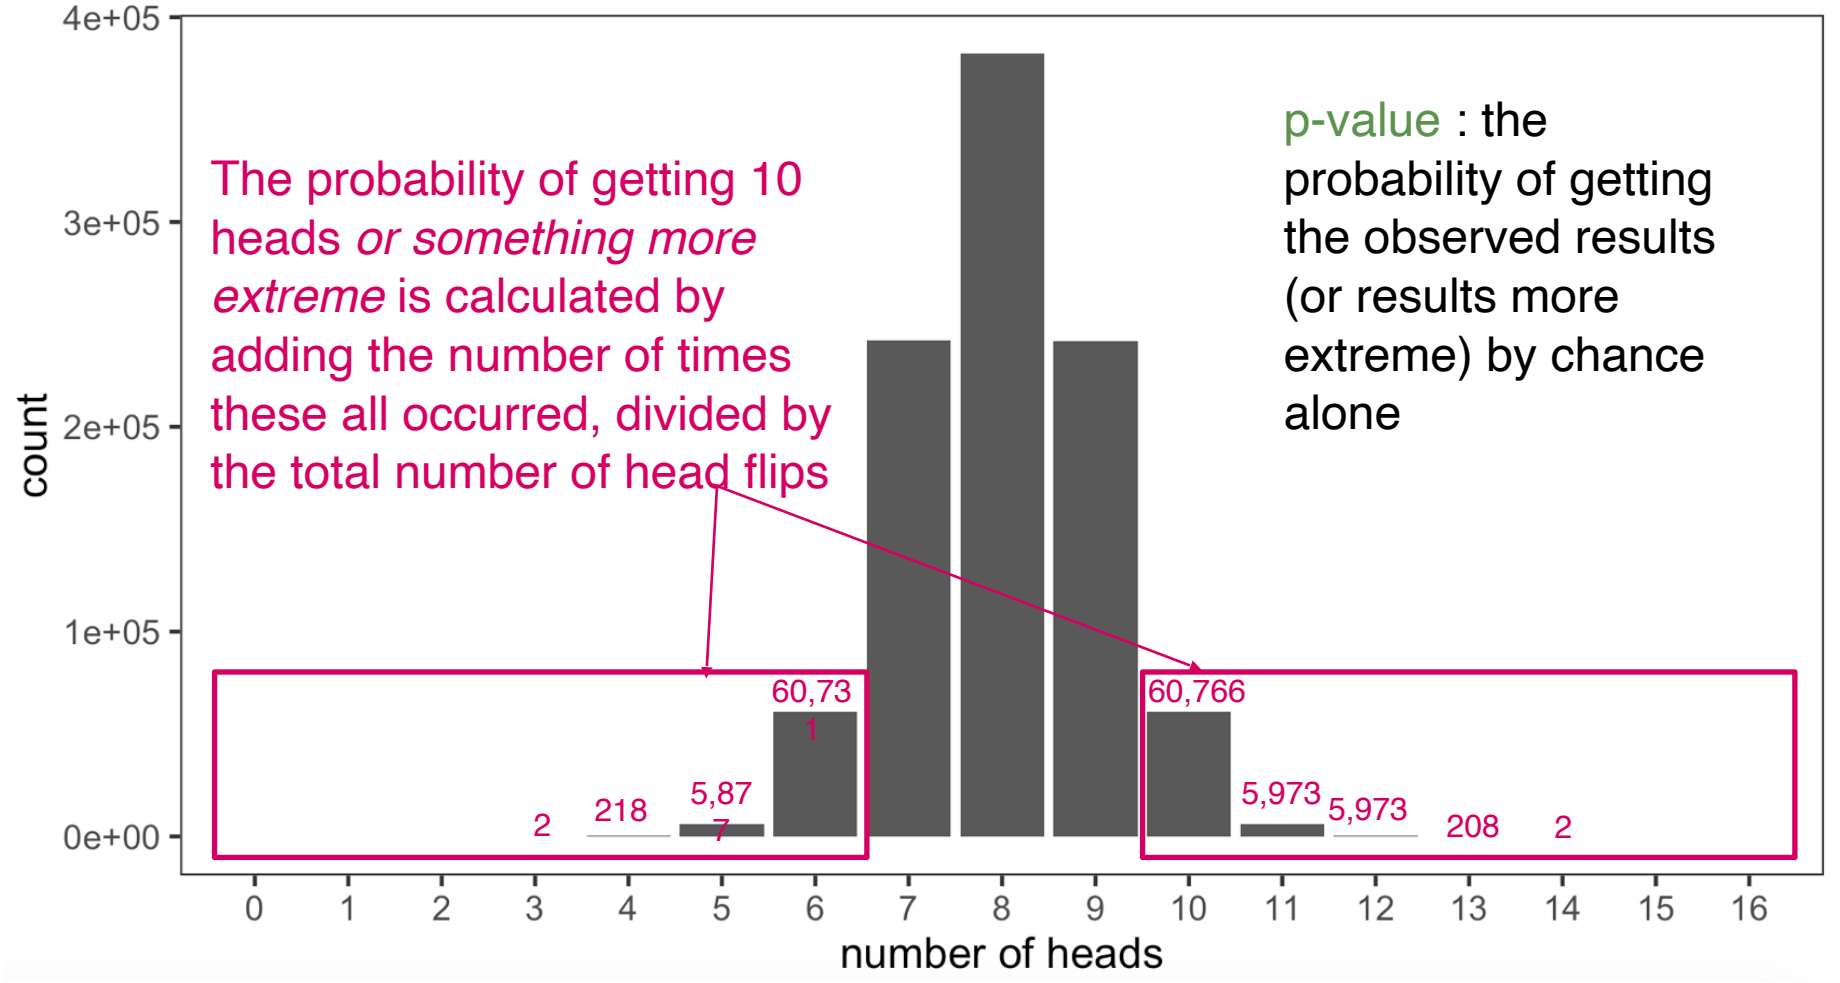
The probability of getting 10 heads *or something more extreme* is calculated by adding the number of times these all occurred, divided by the total number of head flips
p-value : the probability of getting the observed results (or results more extreme) by chance alone
count
4e+05
3e+05
2e+05
1e+05
0e+00
0
1
2
3
4
5
6
7
8
9
10
11
12
13
14
15
16
number of heads
2
218
5,87
7
60,73
1
60,766
5,973
5,973
208
2

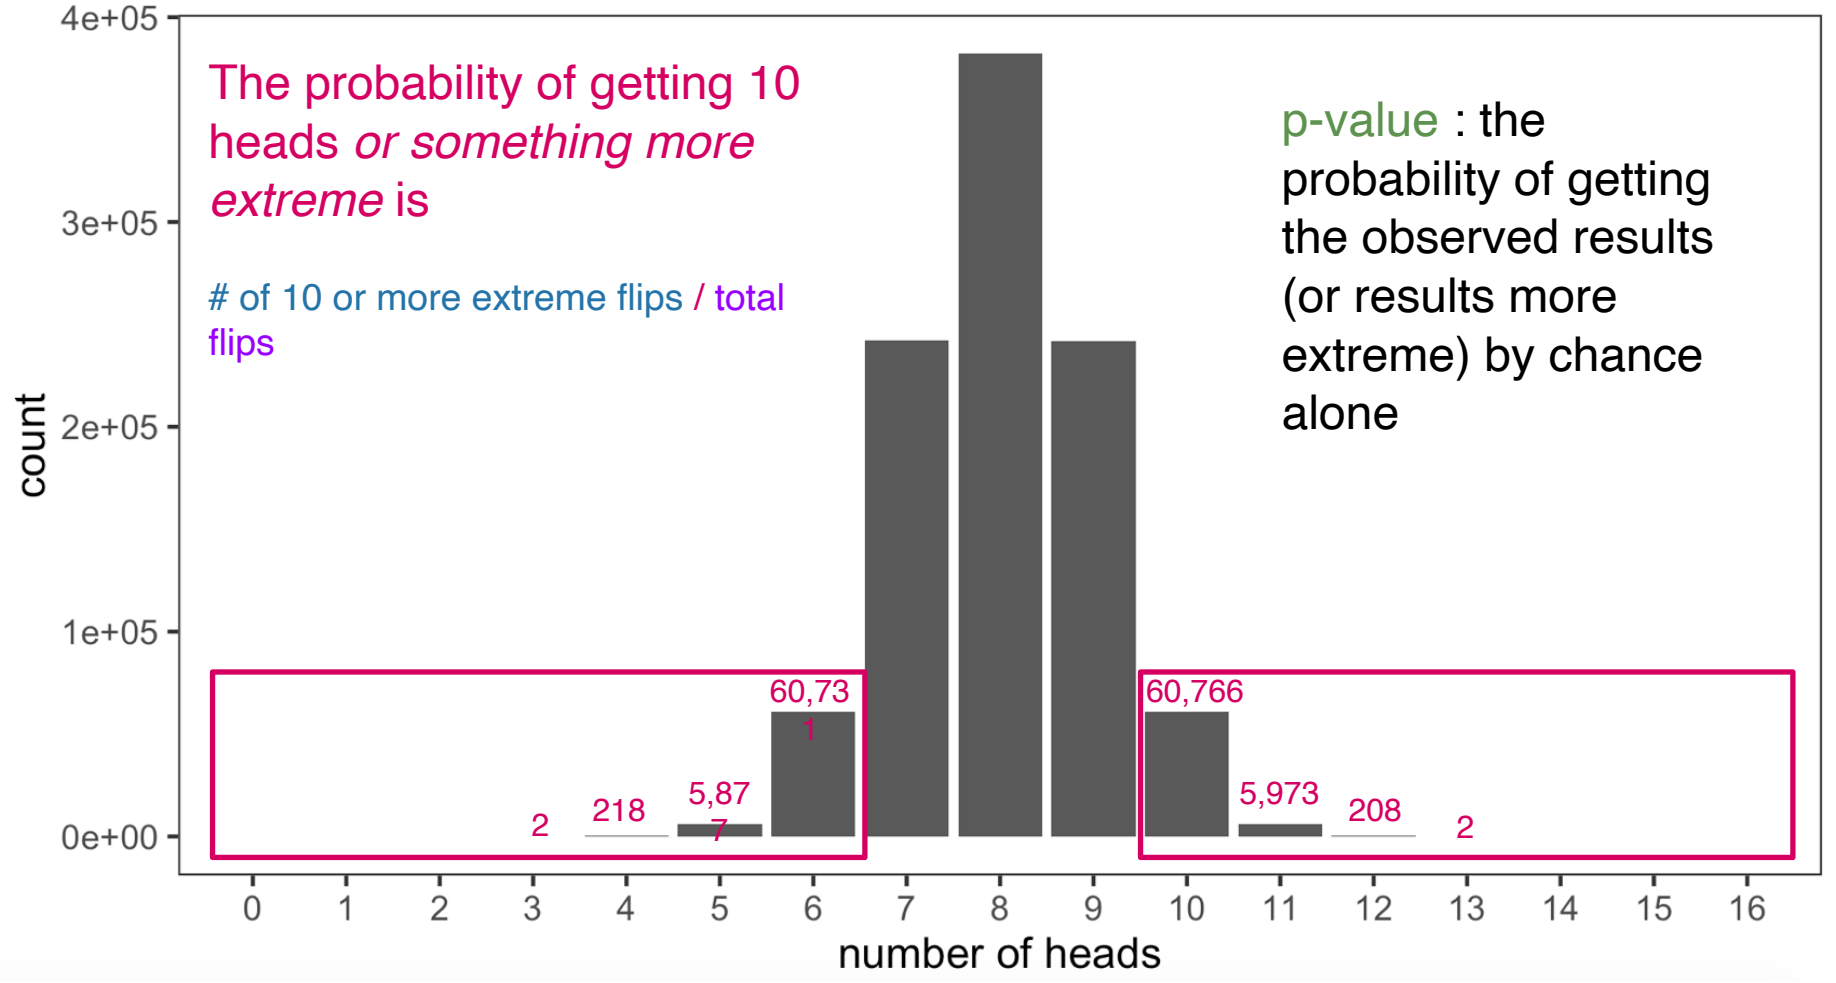
The probability of getting 10 heads *or something more extreme* is
# of 10 or more extreme flips / total flips
p-value : the probability of getting the observed results (or results more extreme) by chance alone
count
4e+05
3e+05
2e+05
1e+05
0e+00
2
218
5,87
7
60,73
1
60,766
5,973
208
2
0
1
2
3
4
5
6
7
8
9
10
11
12
13
14
15
16
number of heads

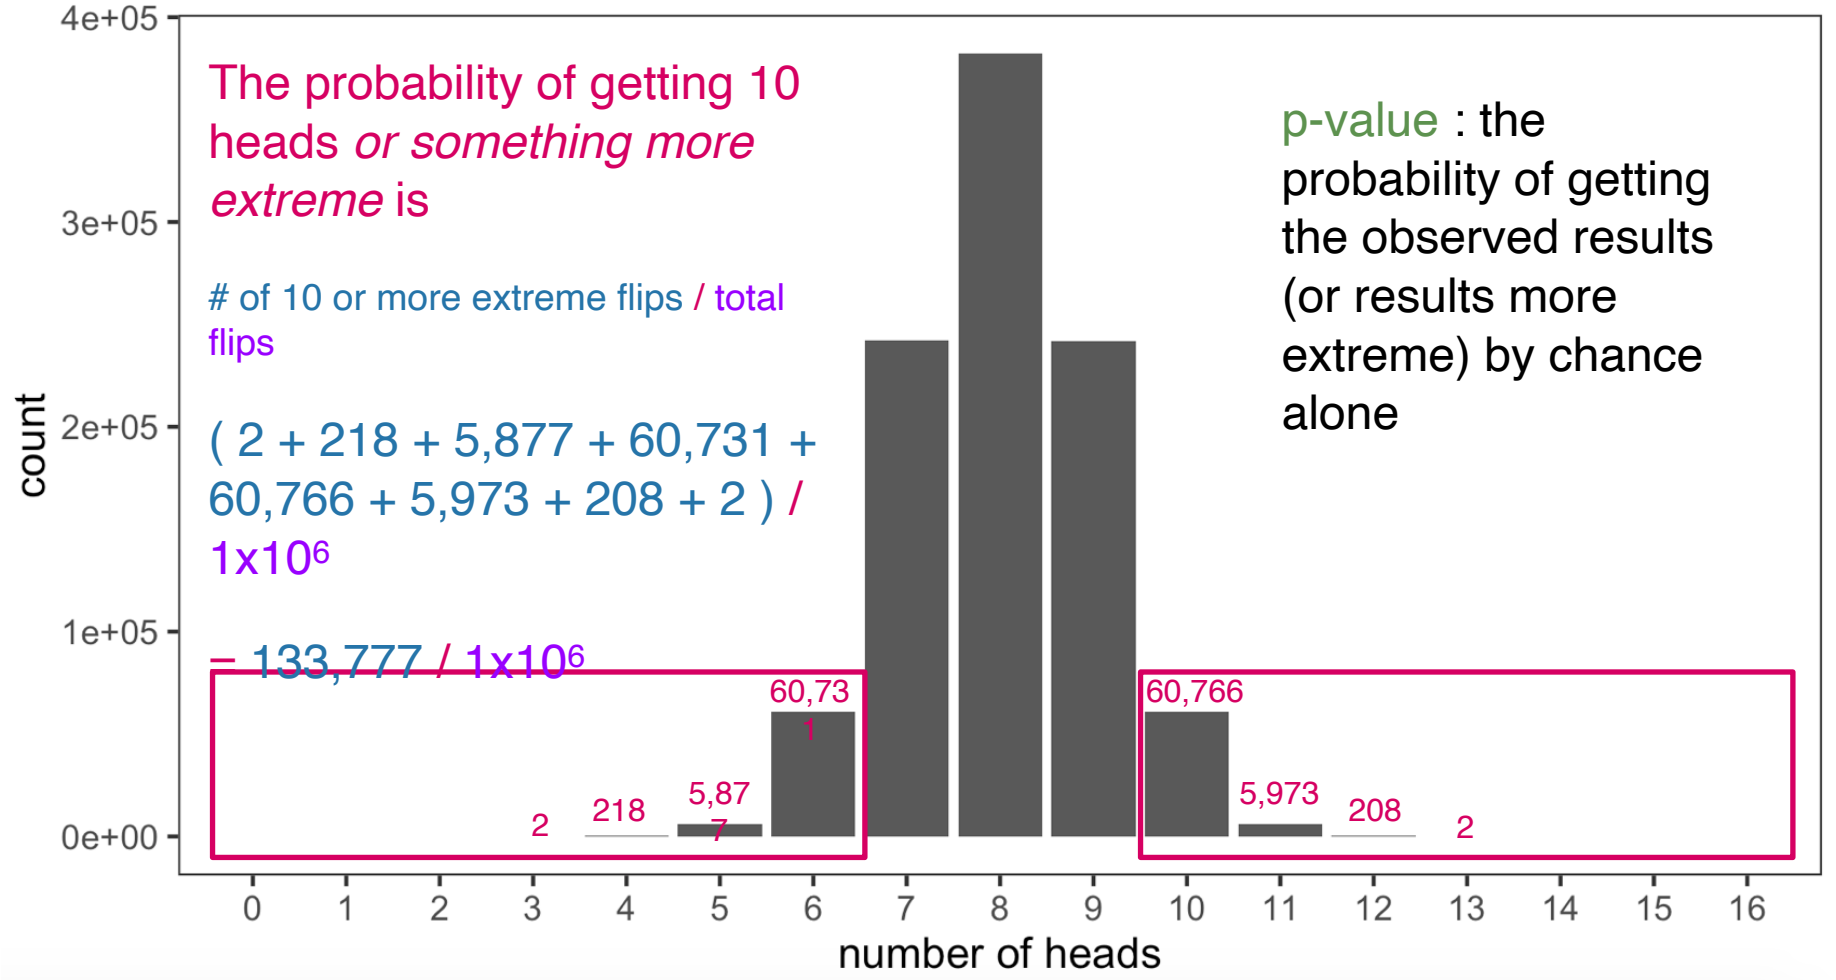
The probability of getting 10 heads *or something more extreme* is
# of 10 or more extreme flips / total flips
(2 + 218 + 5,877 + 60,731 + 60,766 + 5,973 + 208 + 2) / 1x10^6
= 133,777 / 1x10^6
p-value : the probability of getting the observed results (or results more extreme) by chance alone
4e+05
3e+05
2e+05
1e+05
0e+00
count
2
218
5,87
7
60,73
1
60,766
5,973
208
2
0
1
2
3
4
5
6
7
8
9
10
11
12
13
14
15
16
number of heads

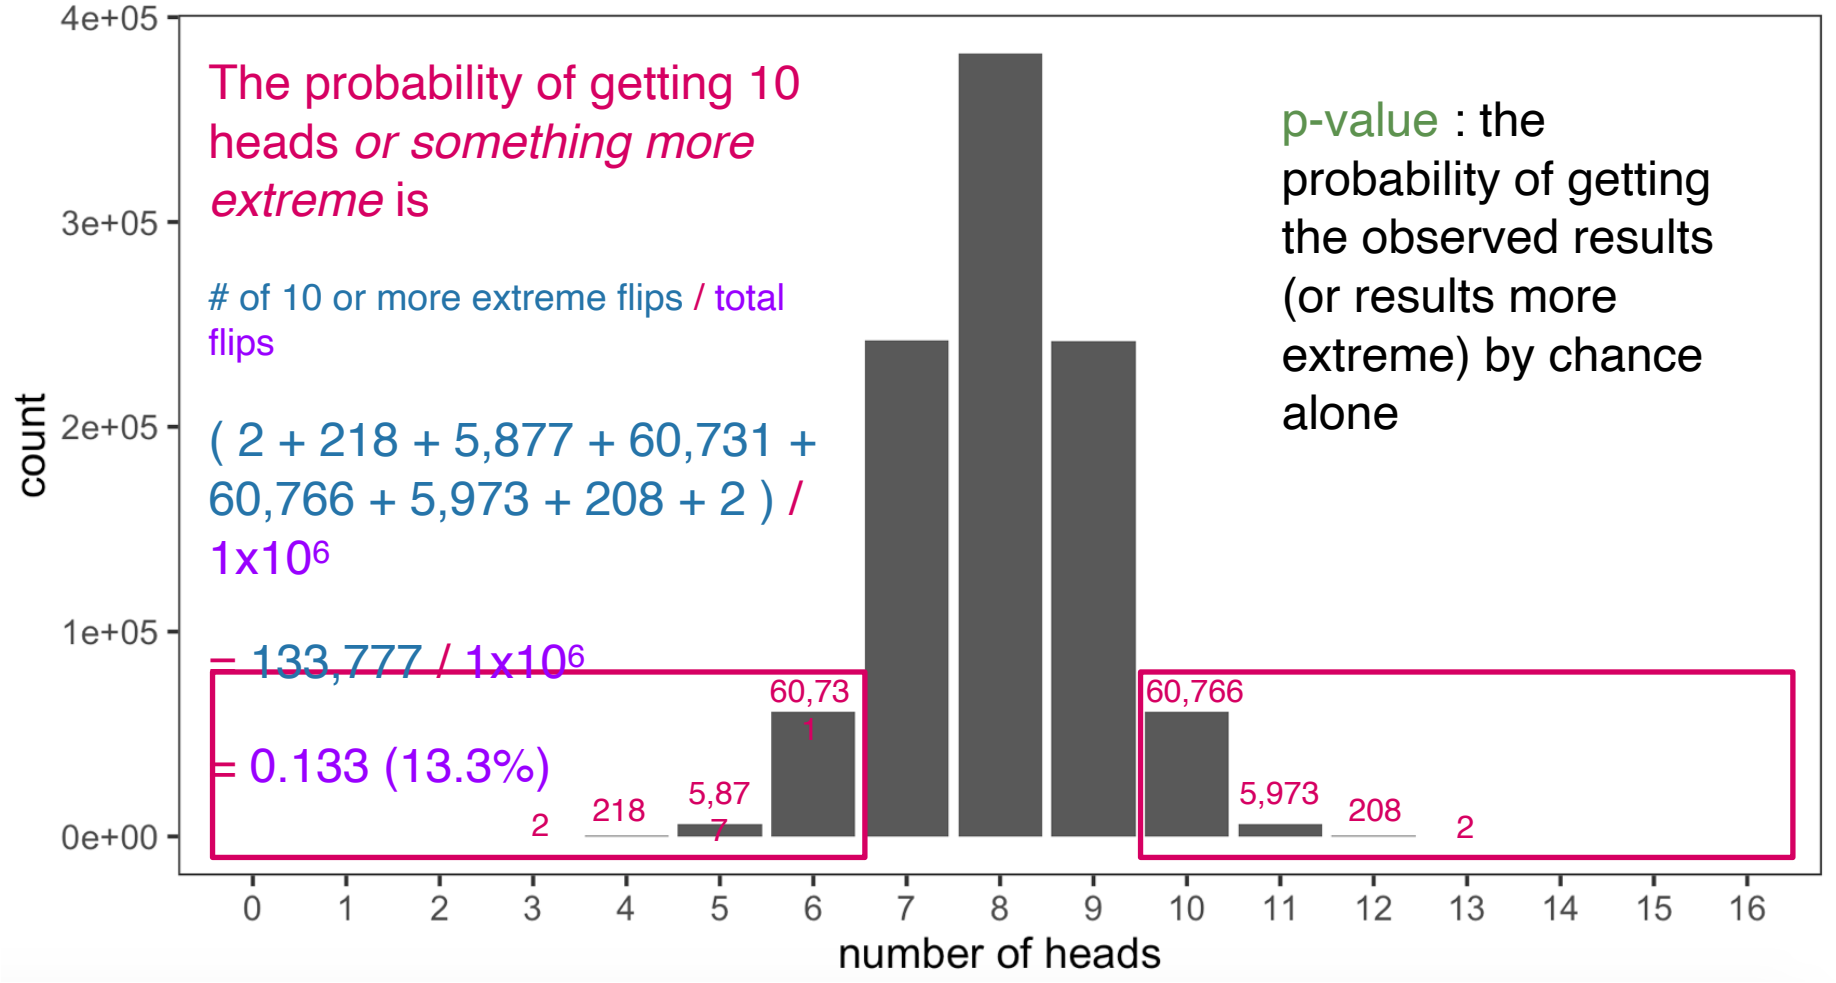
The probability of getting 10 heads *or something more extreme* is
# of 10 or more extreme flips / total flips
(2 + 218 + 5,877 + 60,731 + 60,766 + 5,973 + 208 + 2) / 1x10^6
= 133,777 / 1x10^6
= 0.133 (13.3%)
p-value : the probability of getting the observed results (or results more extreme) by chance alone
count
4e+05
3e+05
2e+05
1e+05
0e+00
2
218
5,87
7
60,73
1
60,766
5,973
208
2
0
1
2
3
4
5
6
7
8
9
10
11
12
13
14
15
16
number of heads

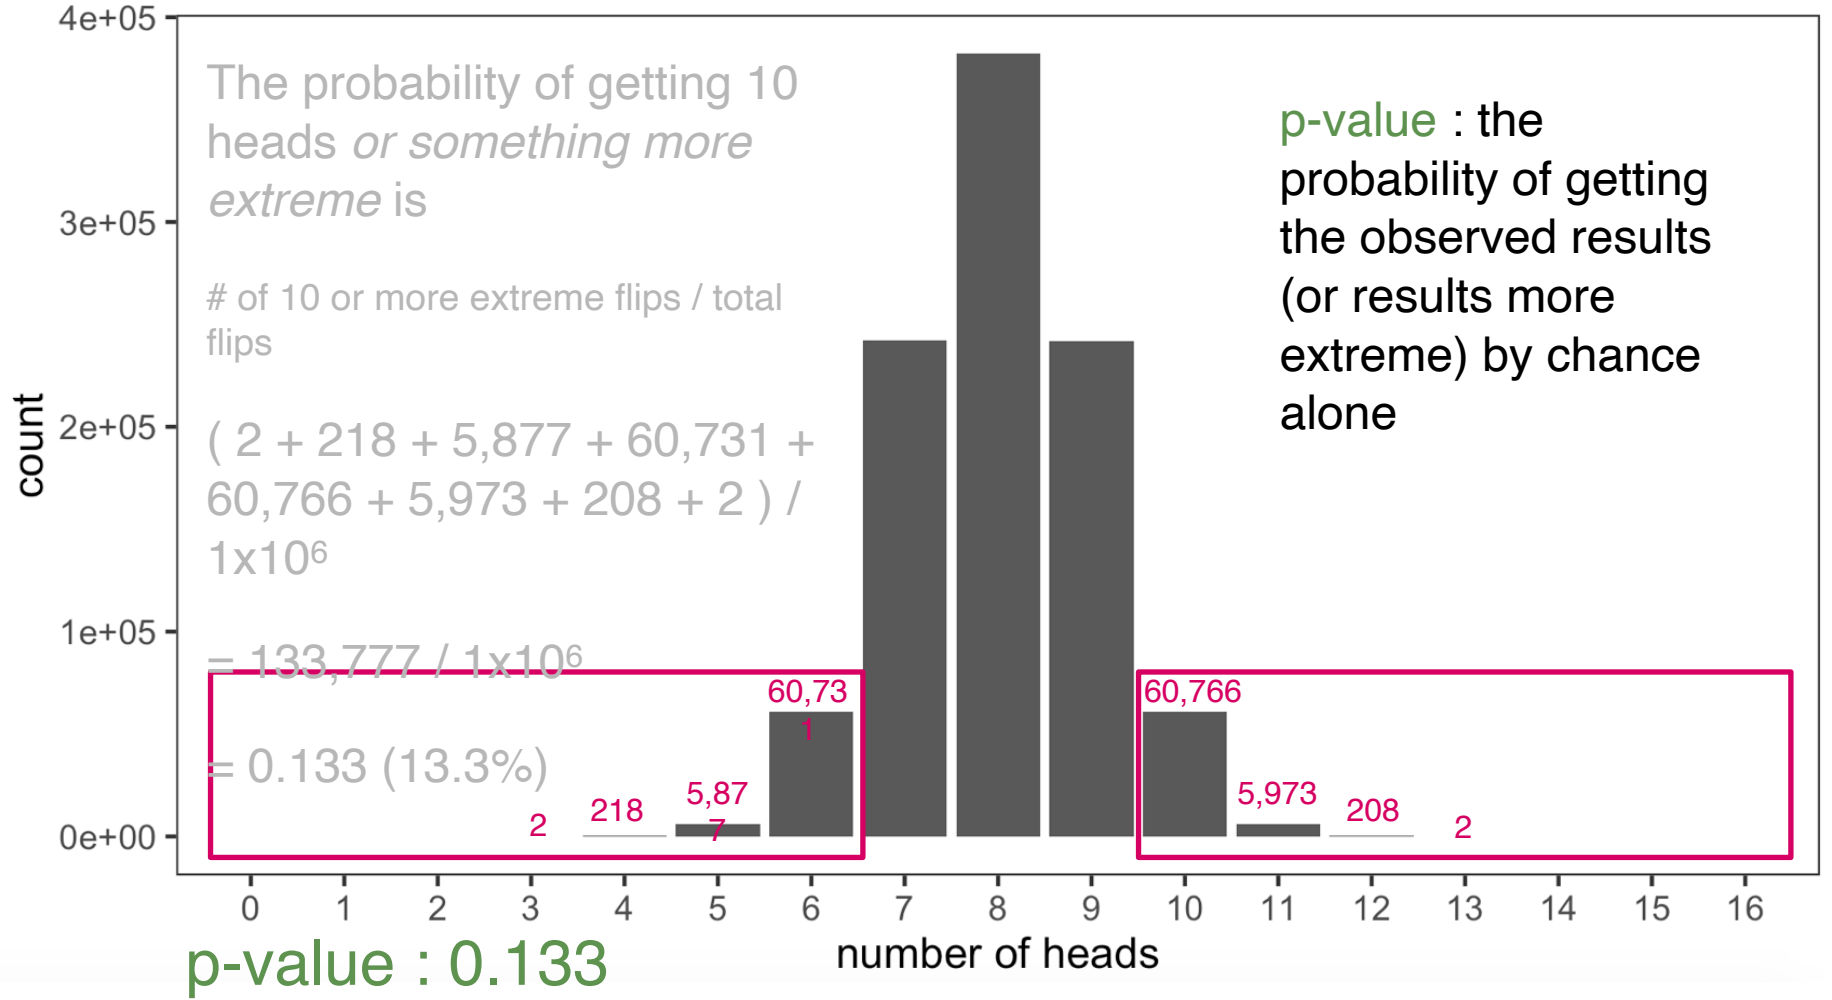
The probability of getting 10 heads *or something more extreme* is
# of 10 or more extreme flips / total flips
(2 + 218 + 5,877 + 60,731 + 60,766 + 5,973 + 208 + 2) / 1x10^6
= 133,777 / 1x10^6
= 0.133 (13.3%)
p-value : the probability of getting the observed results (or results more extreme) by chance alone
4e+05
3e+05
2e+05
1e+05
0e+00
count
2
218
5,87
7
60,73
1
60,766
5,973
208
2
0 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16
number of heads
p-value : 0.133

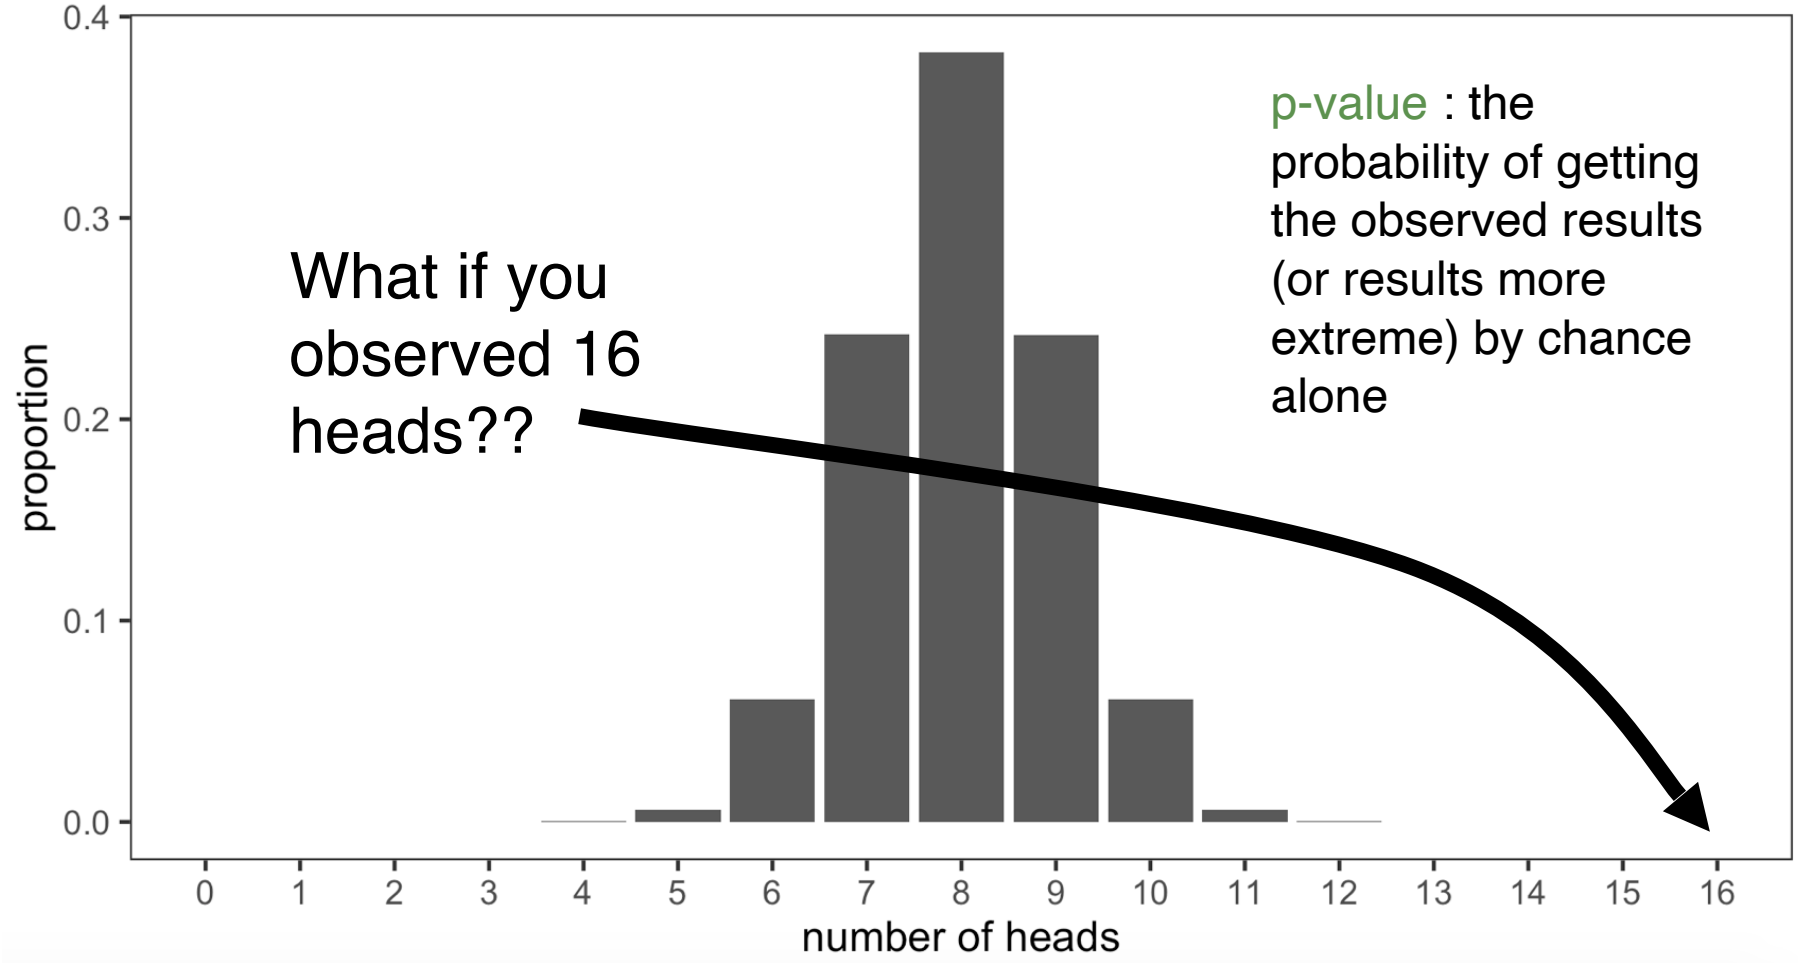
What if you observed 16 heads??
p-value : the probability of getting the observed results (or results more extreme) by chance alone
proportion
0.4
0.3
0.2
0.1
0.0
0 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16
number of heads

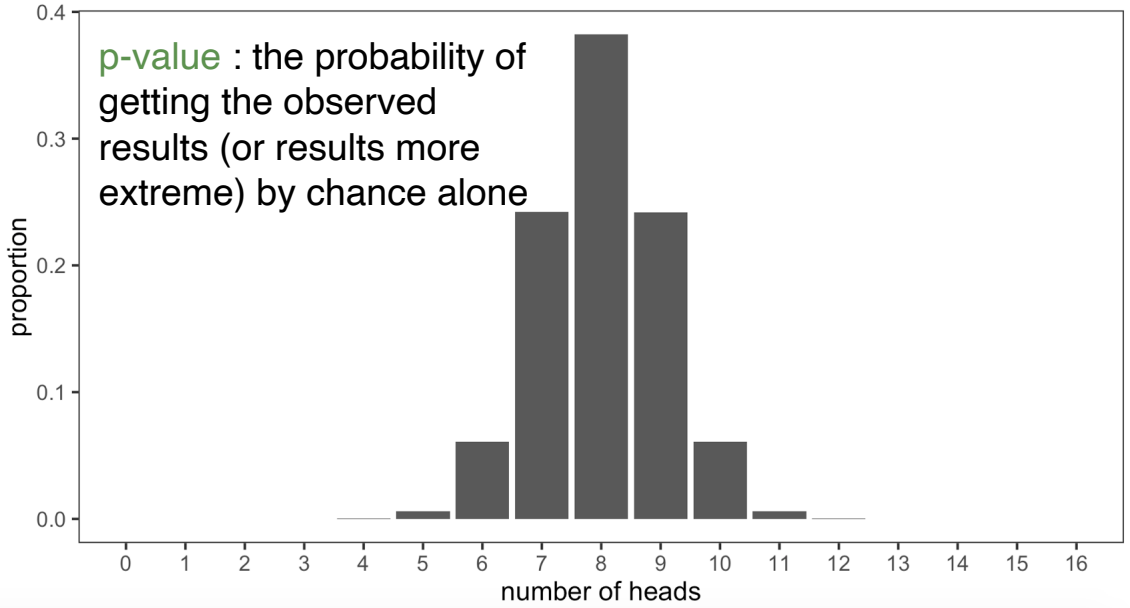

What would be the p-value of you flipping 16 heads?

A
< 0.13

B
> 0.13

The interpretation of p-values in NHST

- Null hypothesis statistical testing means we assume there is no difference
- P-value represents the evidence against the null
  - What's the likelihood of getting values this weird **given an assumption that A and B come from the same distribution**?
- Historically because some old dead guy said so, we take a value of p=0.05 as a threshold of "significance"
- This represents a 1/20 chance of this data under the null assumption
- What happens if you ask 20 questions using NHST?
  - Take a look at this: https://xkcd.com/882/

The interpretation of p-values in NHST

- When your data are very inconsistent with the null hypothesis, P values can’t determine which of the following two possibilities is more probable:
  - The null hypothesis is true, but your sample is unusual due to random sampling error.
  - The null hypothesis is false
- When your data show the null has a high probability, then we can’t determine which of the following is more probable:
  - The null hypothesis is false, but your sample is unusual due to random sampling error.
  - The null hypothesis is true

The interpretation of p-values in NHST

- P-value is NOT the chance of a Type I error! This is a common mistake!
  - Type 1 error: rejecting the null, when it is really true
    - AKA a false positive result
  - P speaks about the assumption of the null being true… why did we make that assumption?
  - Bayesian viewpoint: What's the prior probability of the null?
    - If this is the first study ever on a topic then we know less about the truth than if this is the 36th study on a topic
  - Empirically a p=0.05 is usually a **MUCH BIGGER Type 1** error rate
    - Can't know in reality, but simulation / math studies tell us that we can expect something like for p=0.05 a Type 1 error rate of somewhere like 23 to 50%

Difference in Means

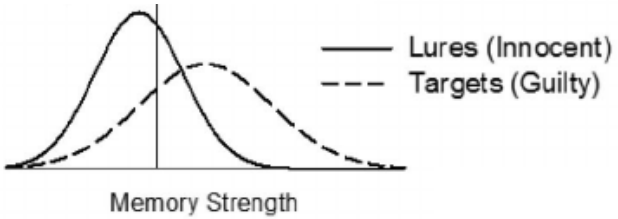

Why would a t-test *not* be appropriate for these data?

A
Not normally distributed

B
Unequal variances

C
Small sample size

D
Data are not continuous

https://www.researchgate.net

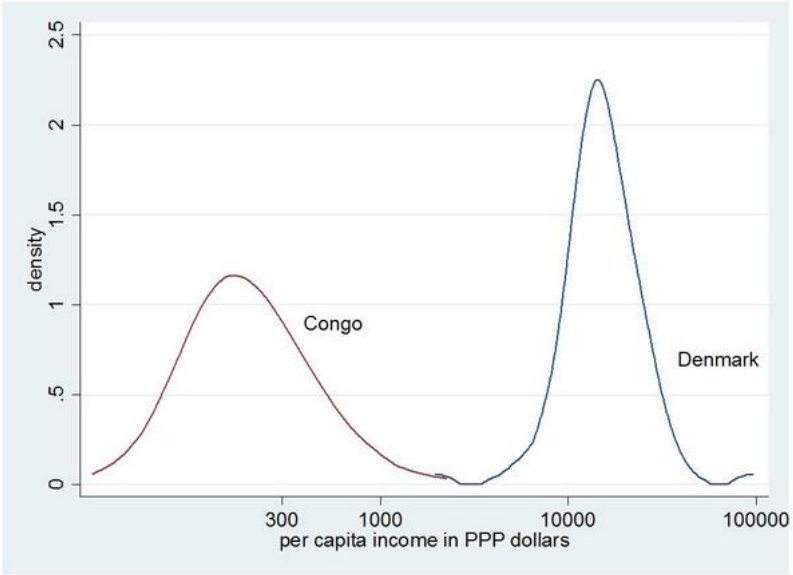

Would a t-test find a significant difference in means?

A
t-test not appropriate

B
Yes

C
No

D
Need more information

https://twitter.com/brankomilan/status/406105955080757249

Paired data

Independent samples t-test

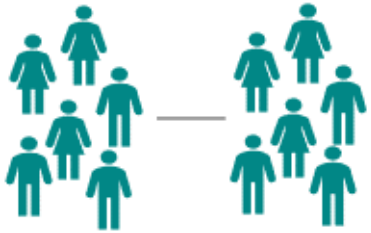

Is there a **difference** between **two groups**

Paired samples t-test

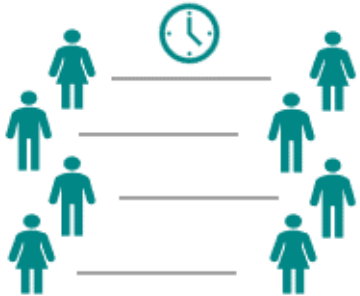

Is there a **difference** in a **group** between **two points in time**

https://datatab.net/tutorial/paired-t-test

Do the heights between males and females differ?

variance is equal(ish) between groups (57 vs 50cm)

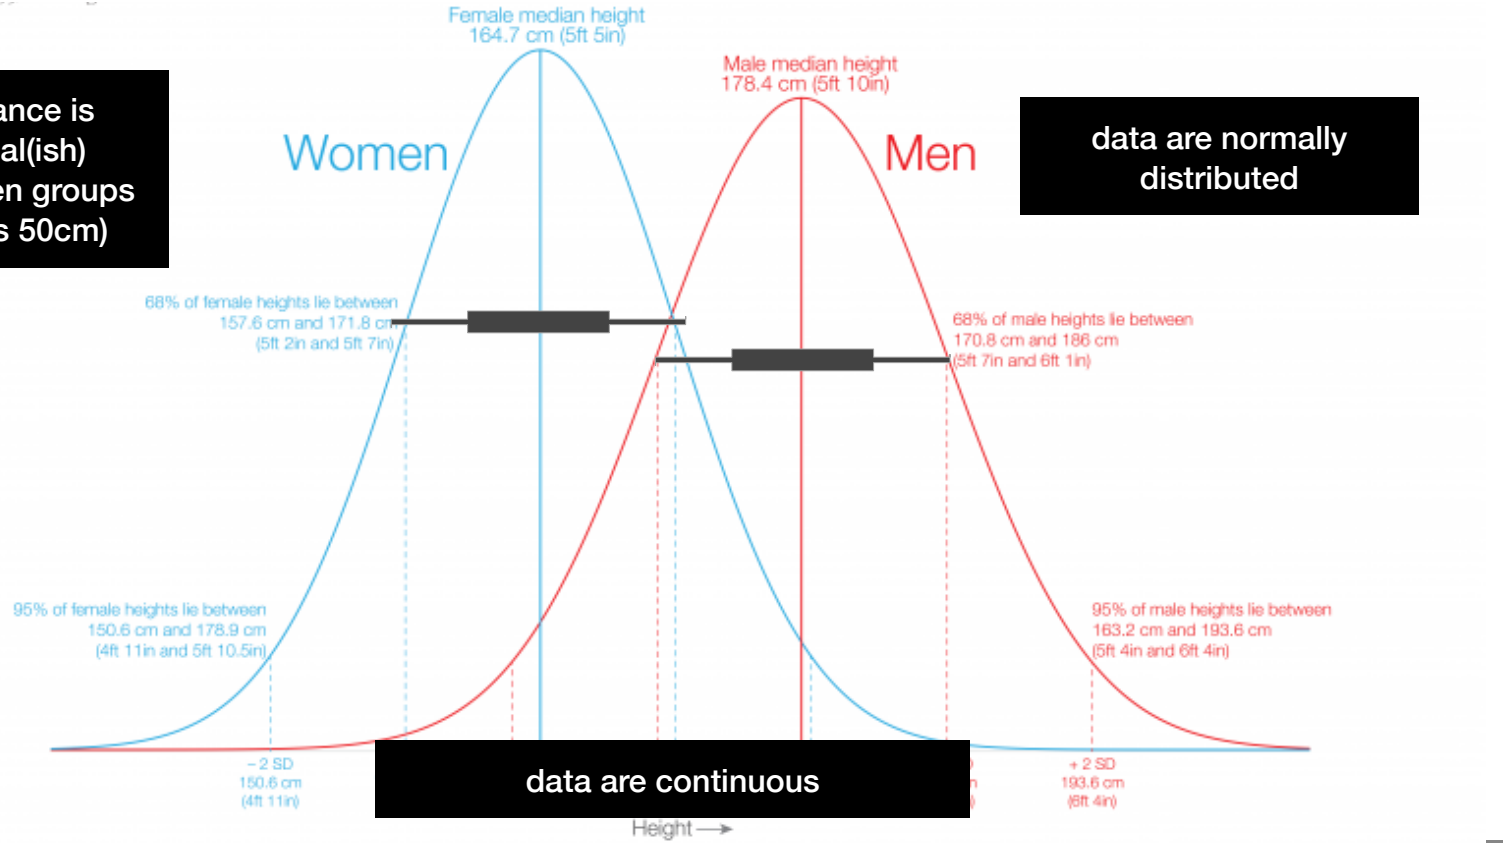

data are normally distributed

sample size affects statistic

N=10,000

data are continuous

Cohen's *d*

Cohen's *d* is defined as the difference between two means divided by a standard deviation for the data

Effect sizes!

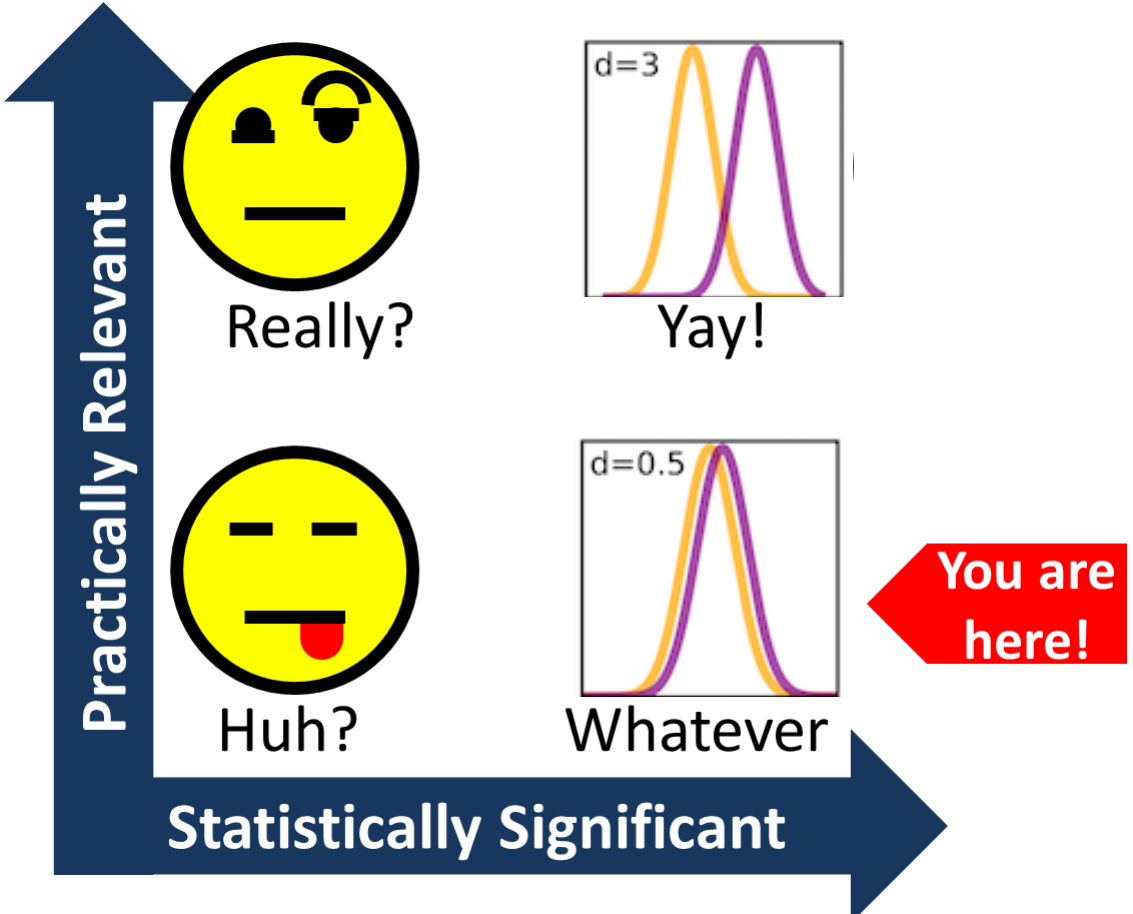

CORRELATION

ASSOCIATION BETWEEN VARIABLES

i.e. Pearson Correlation, Spearman Correlation, chi-square test

COMPARISON OF MEANS

DIFFERENCE IN MEANS BETWEEN VARIABLES

i.e. t-test, ANOVA

REGRESSION

DOES CHANGE IN ONE VARIABLE MEAN CHANGE IN ANOTHER?

I.e. simple regression, multiple regression

NON-PARAMETRIC TESTS

FOR WHEN ASSUMPTIONS IN THESE OTHER 3 CATEGORIES ARE NOT MET

i.e. Wilcoxon rank-sum test, Wilcoxon sign-rank test, sign test

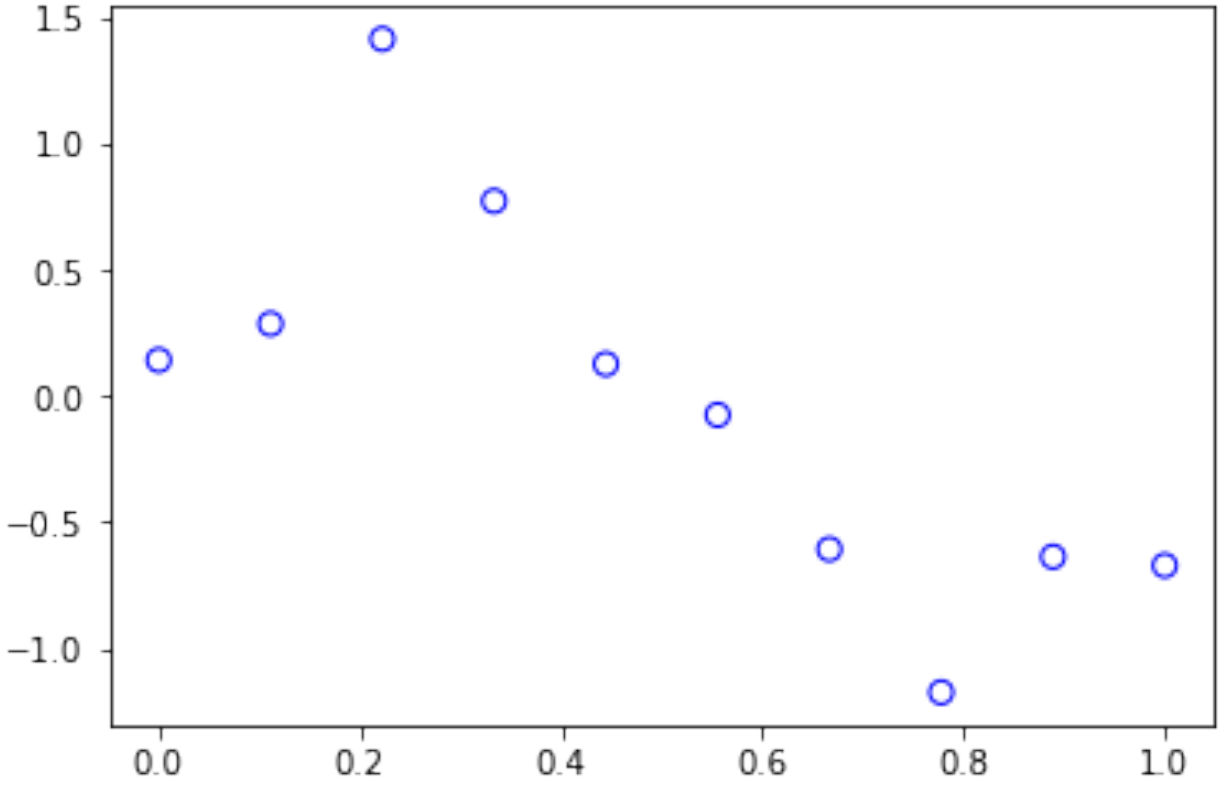
1.5
1.0
0.5
0.0
−0.5
−1.0
0.0
0.2
0.4
0.6
0.8
1.0

CORRELATION

ASSOCIATION BETWEEN VARIABLES

i.e. Pearson Correlation, Spearman Correlation, chi-square test

COMPARISON OF MEANS

DIFFERENCE IN MEANS BETWEEN VARIABLES

i.e. t-test, ANOVA

REGRESSION

DOES CHANGE IN ONE VARIABLE MEAN CHANGE IN ANOTHER?

I.e. simple regression, multiple regression

NON-PARAMETRIC TESTS

FOR WHEN ASSUMPTIONS IN THESE OTHER 3 CATEGORIES ARE NOT MET

i.e. Wilcoxon rank-sum test, Wilcoxon sign-rank test, sign test

Linear regression can be used to describe this relationship

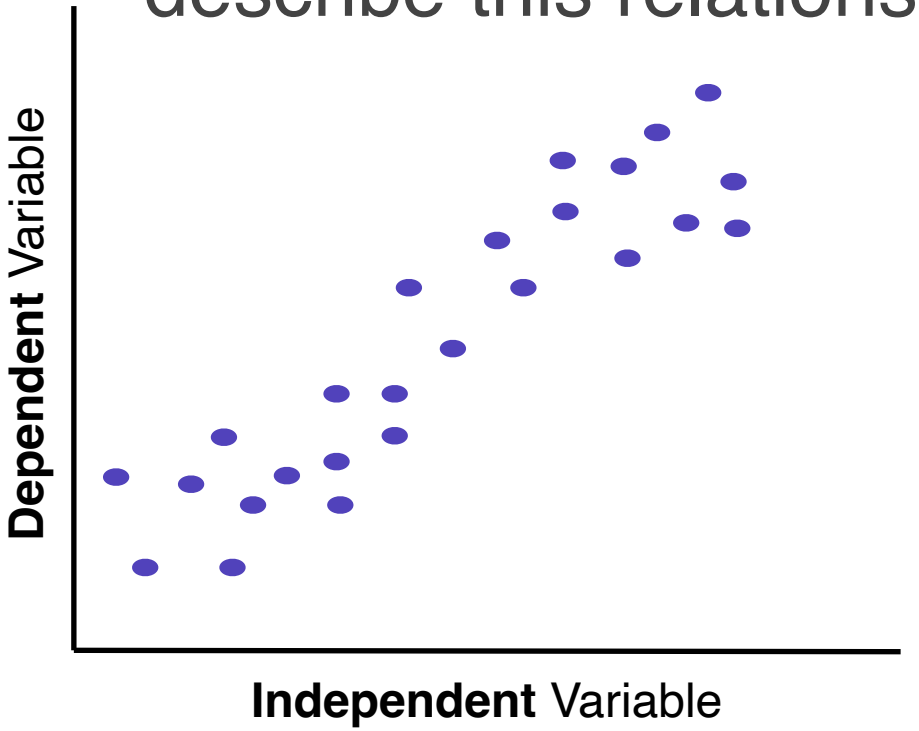

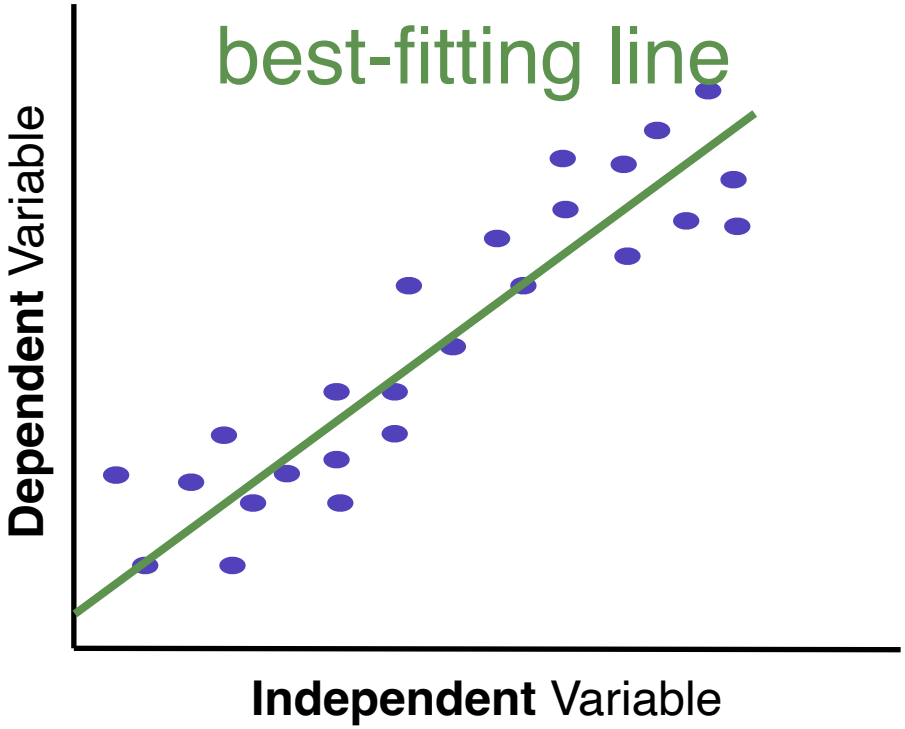
best-fitting line
Dependent Variable
Independent Variable

Best-fitting line

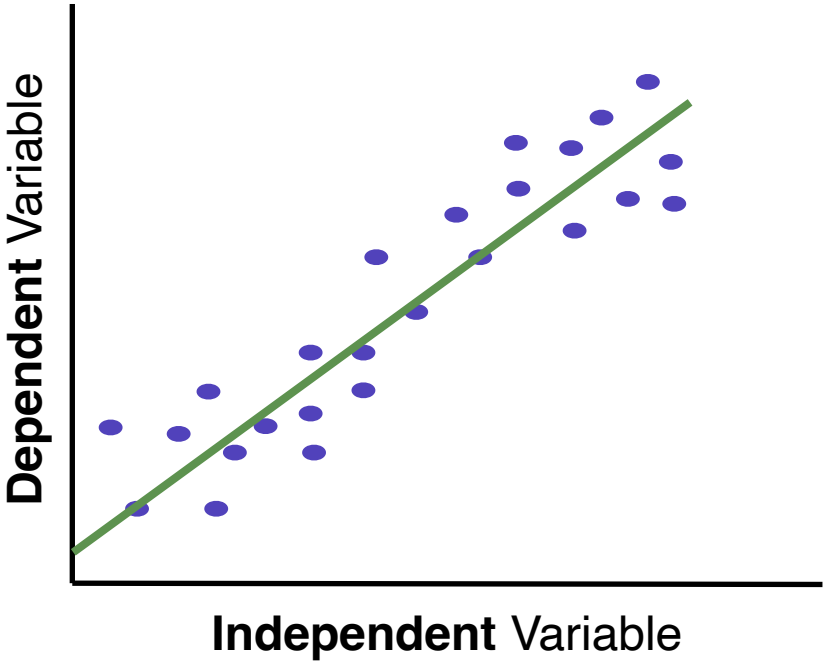

NOT a best-fitting line

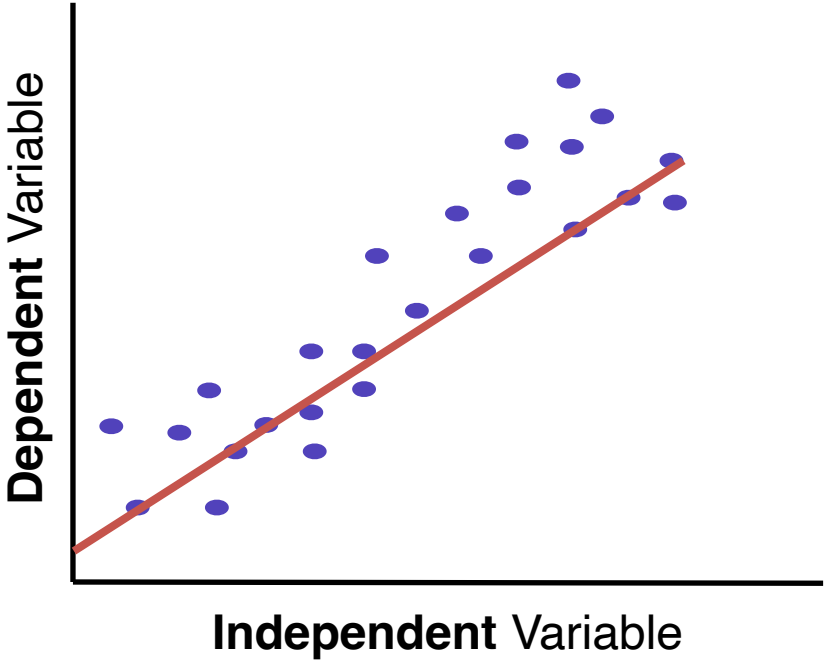

This line is a model of the data

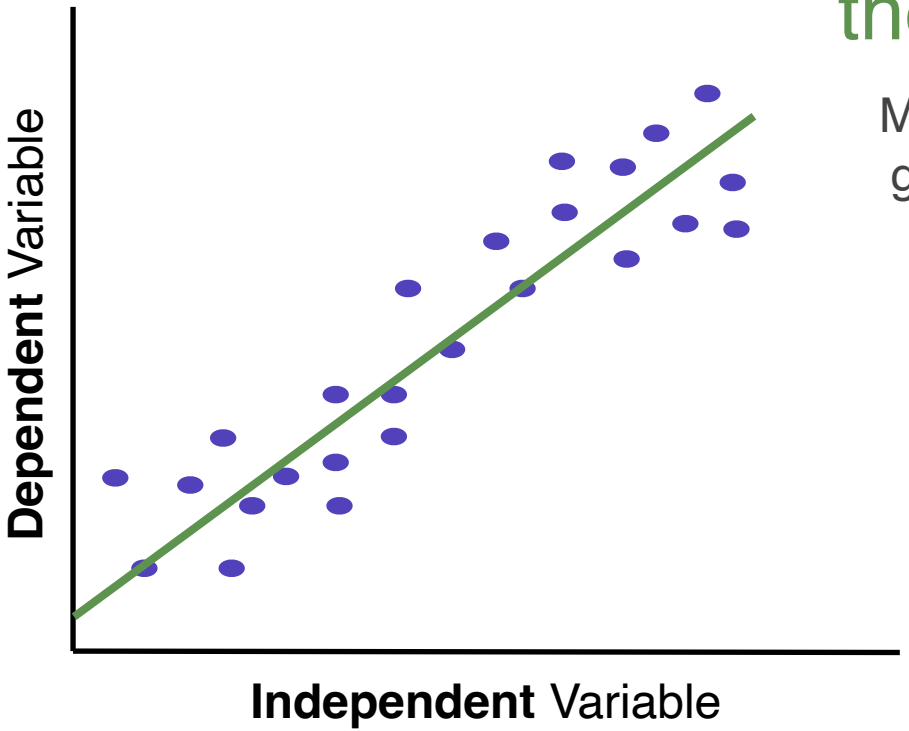

Models are mathematical equations generated to *represent* the real life situation

This line is a model of the data

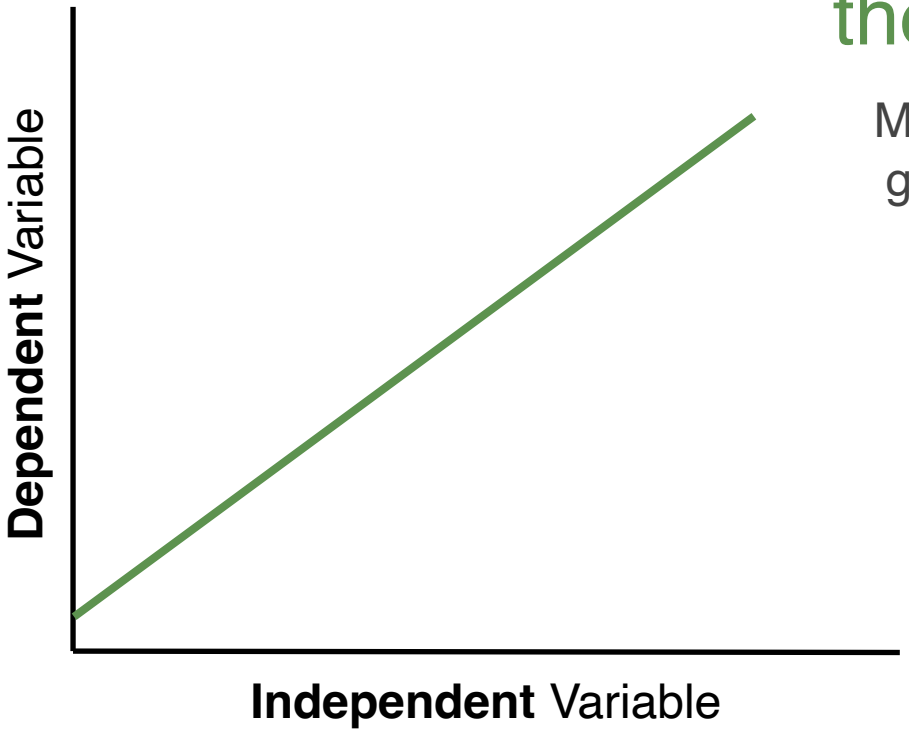

Models are mathematical equations generated to *represent* the real life situation

2.3 Parsimony

Since all models are wrong the scientist cannot obtain a "correct" one by excessive elaboration. On the contrary following William of Occam he should seek an economical description of natural phenomena. Just as the ability to devise simple but evocative models is the signature of the great scientist so overelaboration and overparameterization is often the mark of mediocrity.

2.4 Worrying Selectively

Since all models are wrong the scientist must be alert to what is importantly wrong. It is inappropriate to be concerned about mice when there are tigers abroad.

George Box

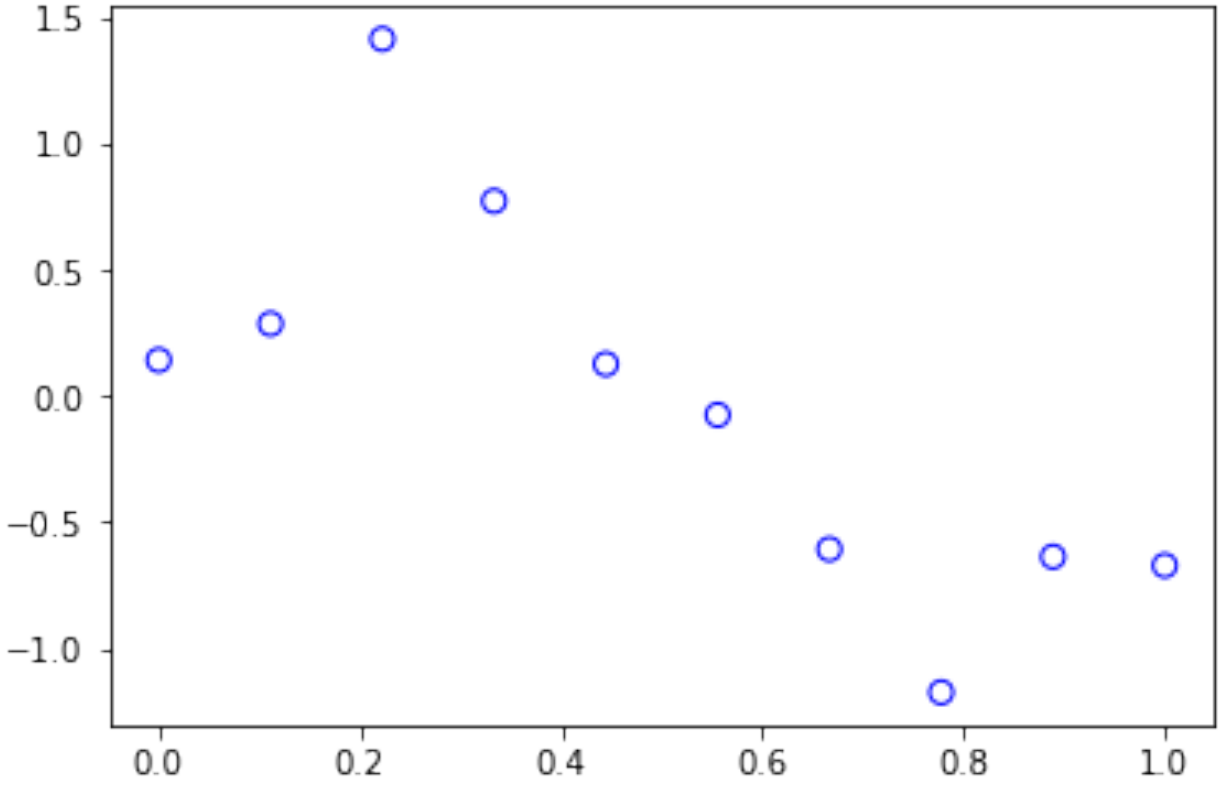
1.5
1.0
0.5
0.0
−0.5
−1.0
0.0
0.2
0.4
0.6
0.8
1.0

weaker relationship

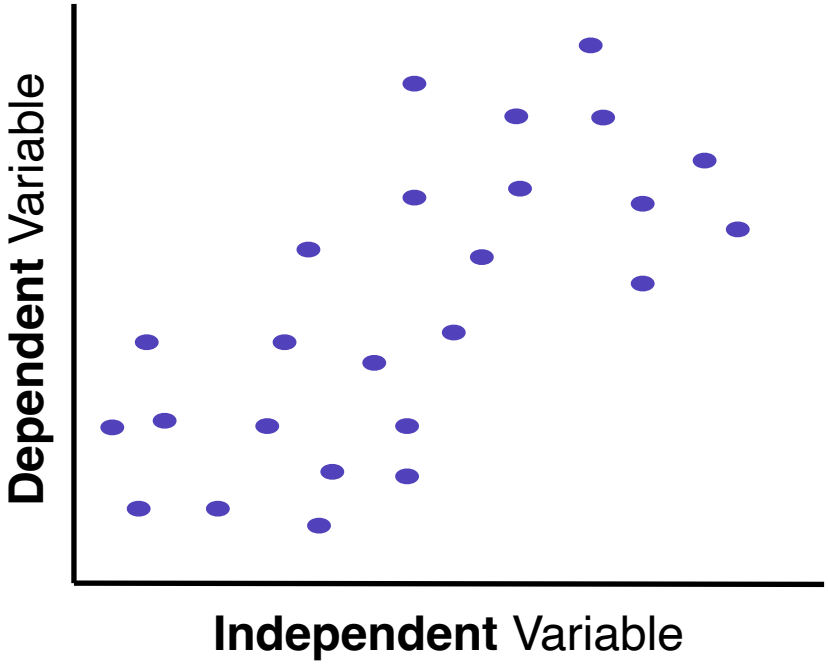

stronger relationship

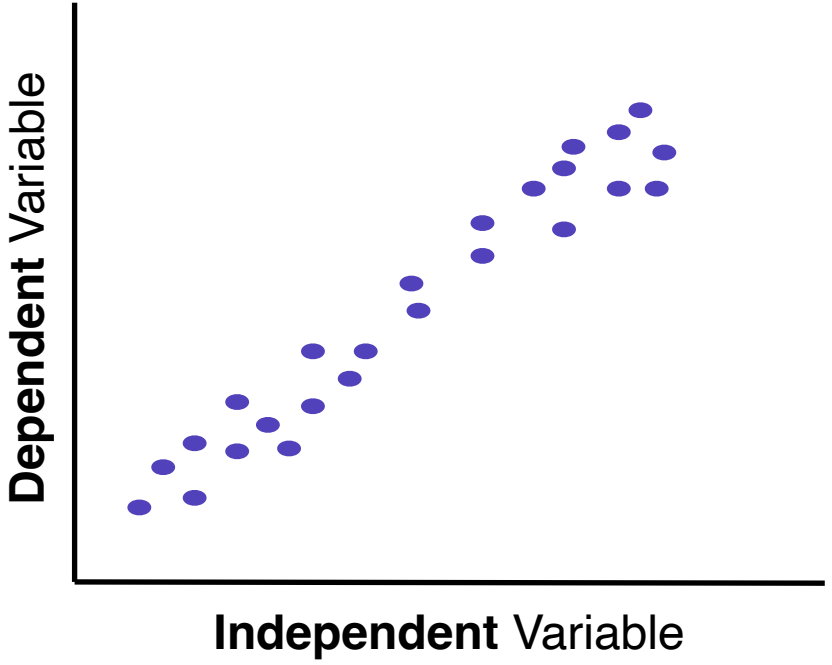

stronger relationship = higher correlation

This is a kind of effect size

The *closer* the points are to the regression line, the *less uncertain* we are in our estimate

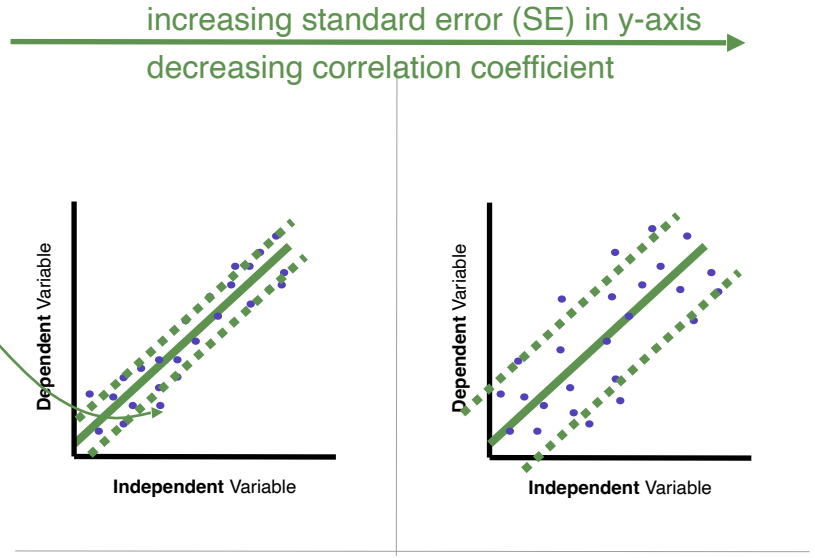

Standard error is standard deviation / sqrt(n)

Linear regression can be used to determine whether a change in one variable is related to the change in the other variable

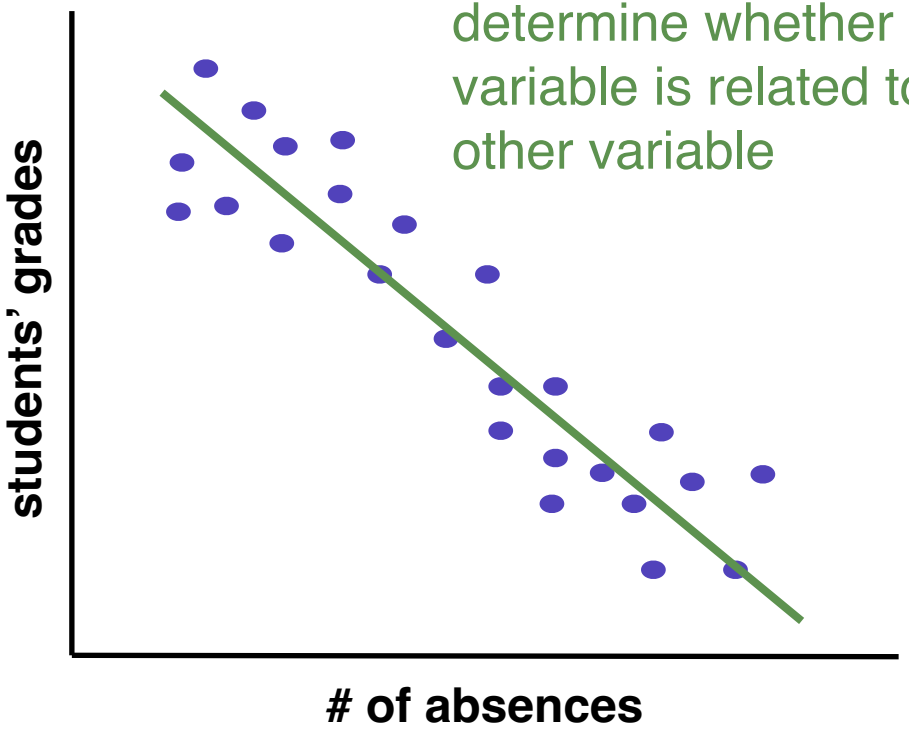

Linear regression can be used to determine whether a change in one variable is related to the change in the other variable

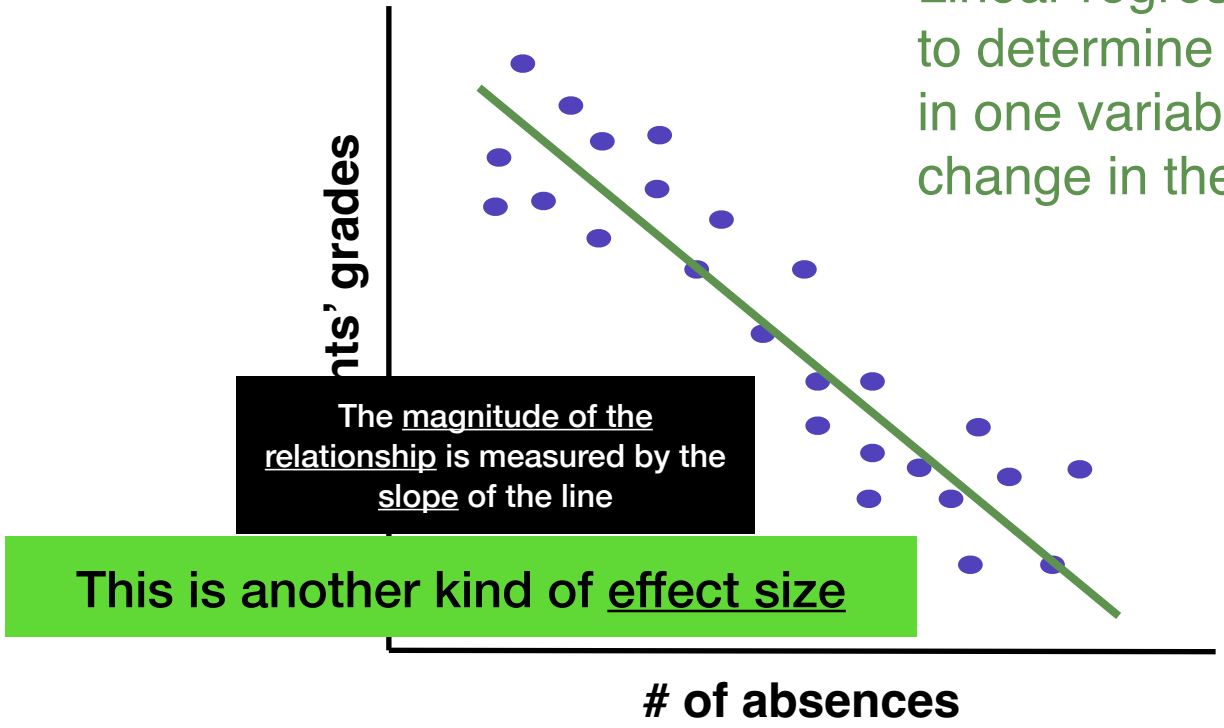

The magnitude of the relationship is measured by the slope of the line

This is another kind of effect size

Effect size (β) can be estimated using the slope of the line

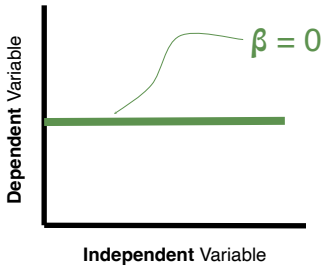

Effect size (β) can be estimated using the slope of the line

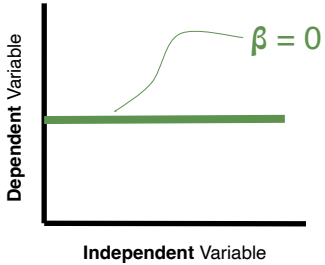

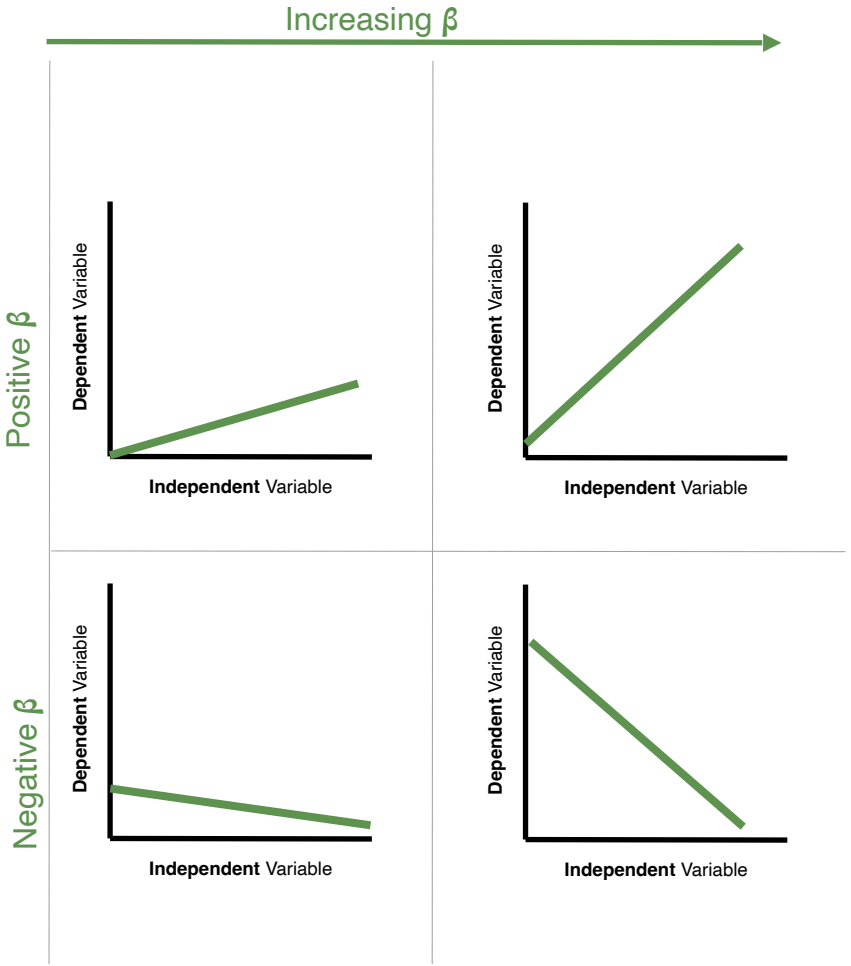

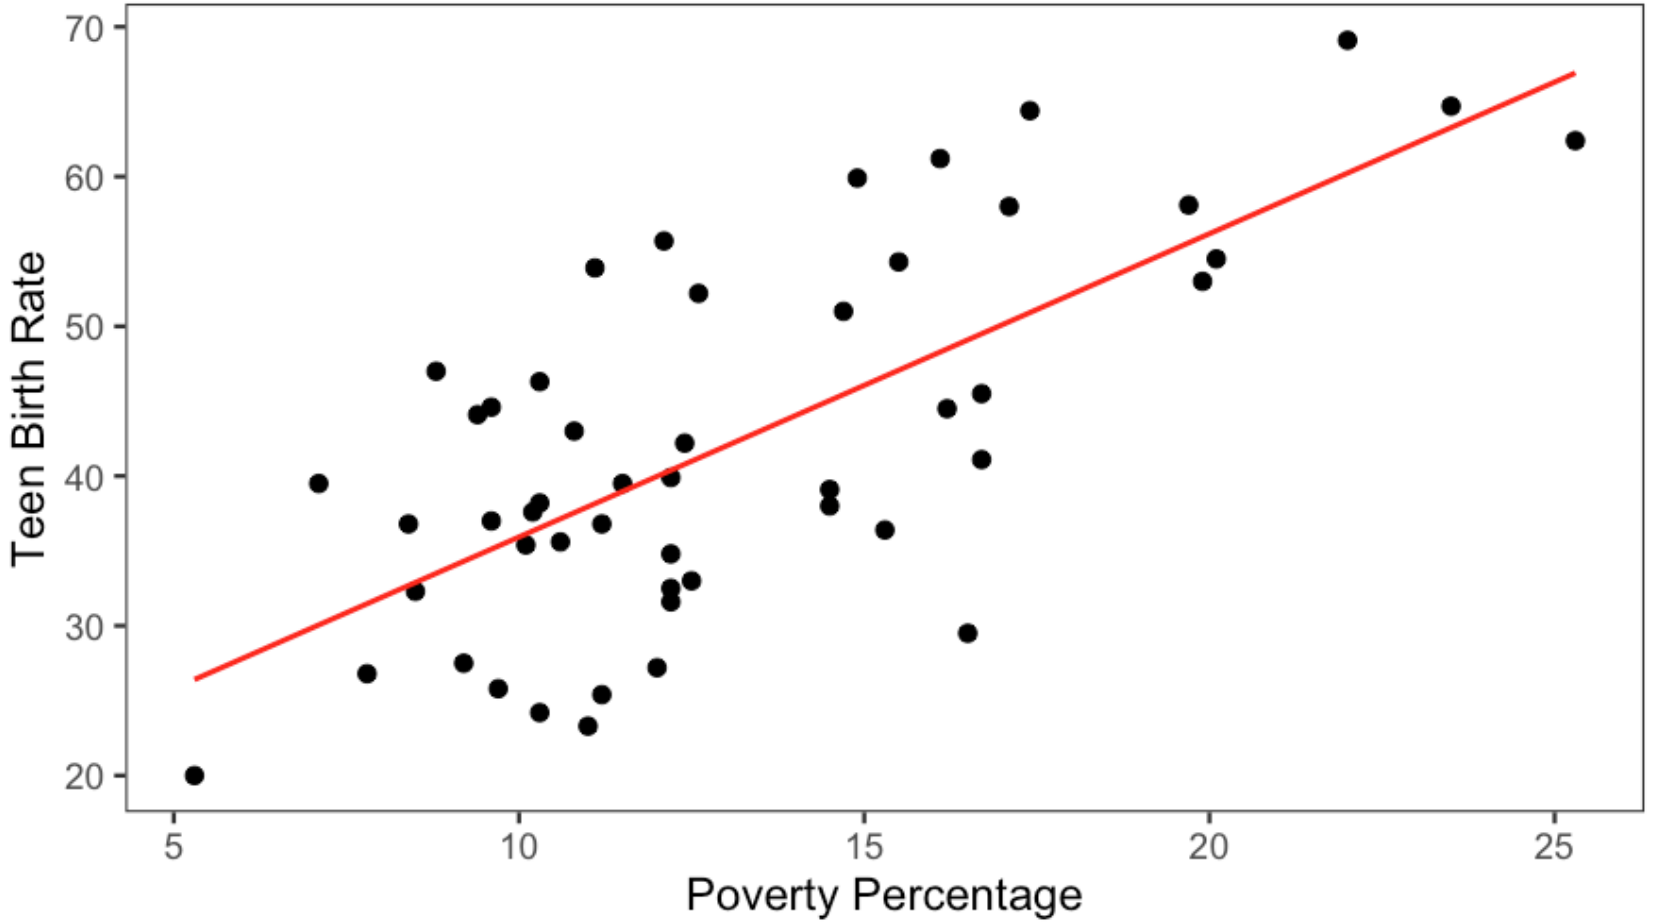
Teen Birth Rate
70
60
50
40
30
20
5
10
15
20
25
Poverty Percentage

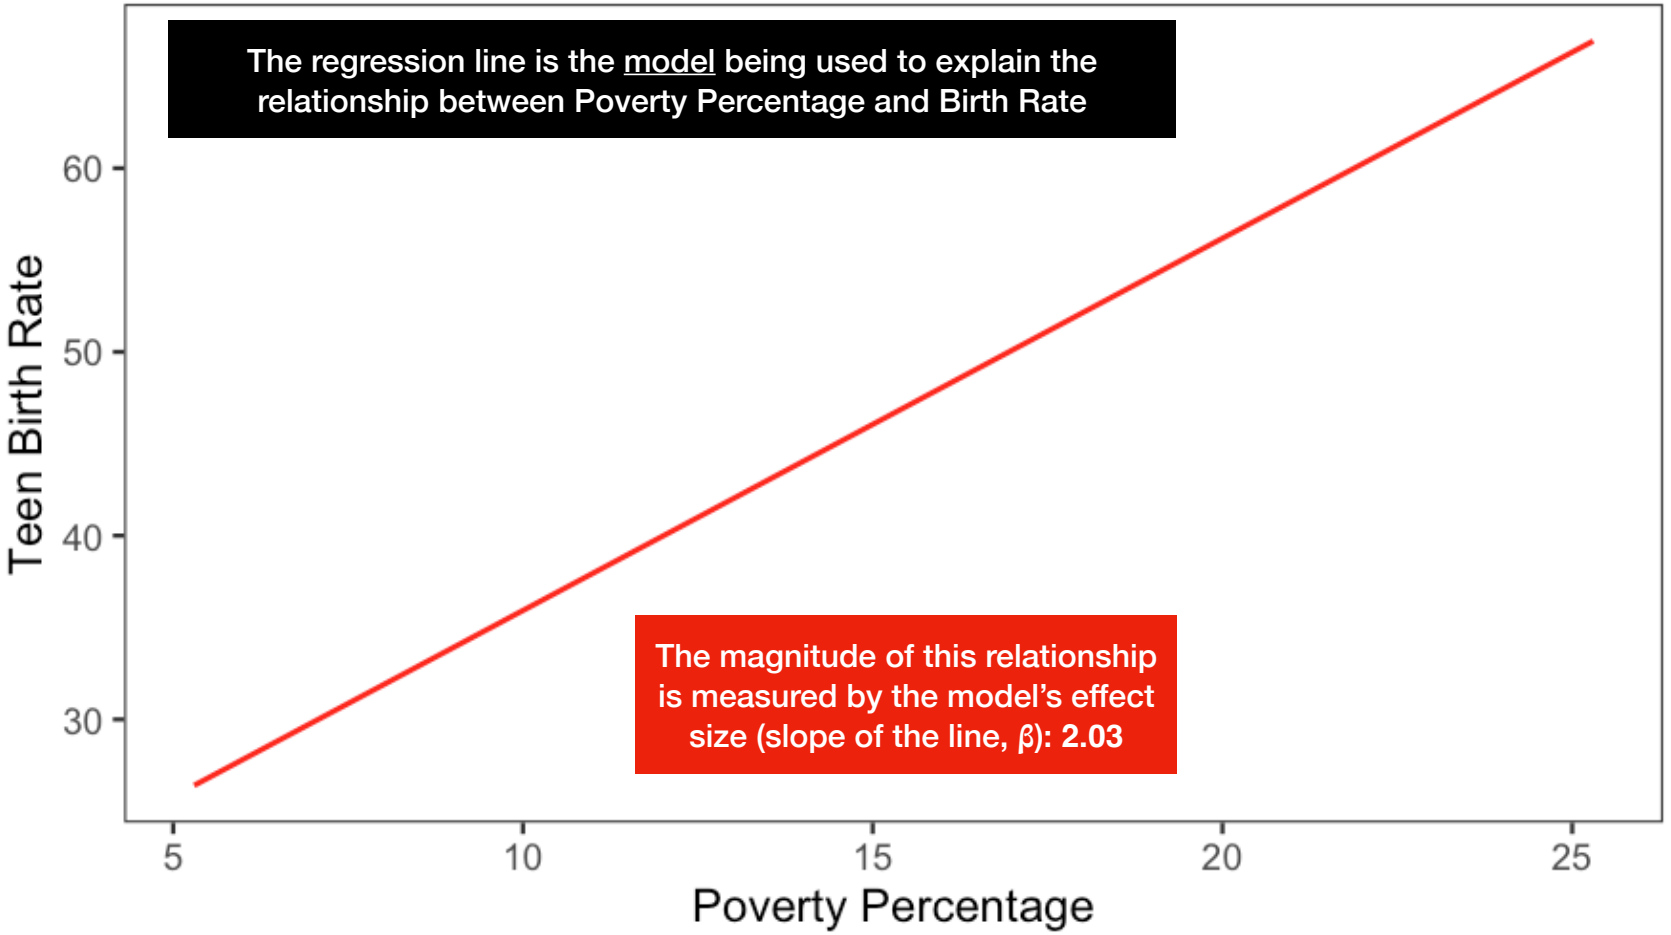
The regression line is the model being used to explain the relationship between Poverty Percentage and Birth Rate
The magnitude of this relationship is measured by the model's effect size (slope of the line, β): 2.03
Teen Birth Rate
60
50
40
30
5
10
15
20
25
Poverty Percentage

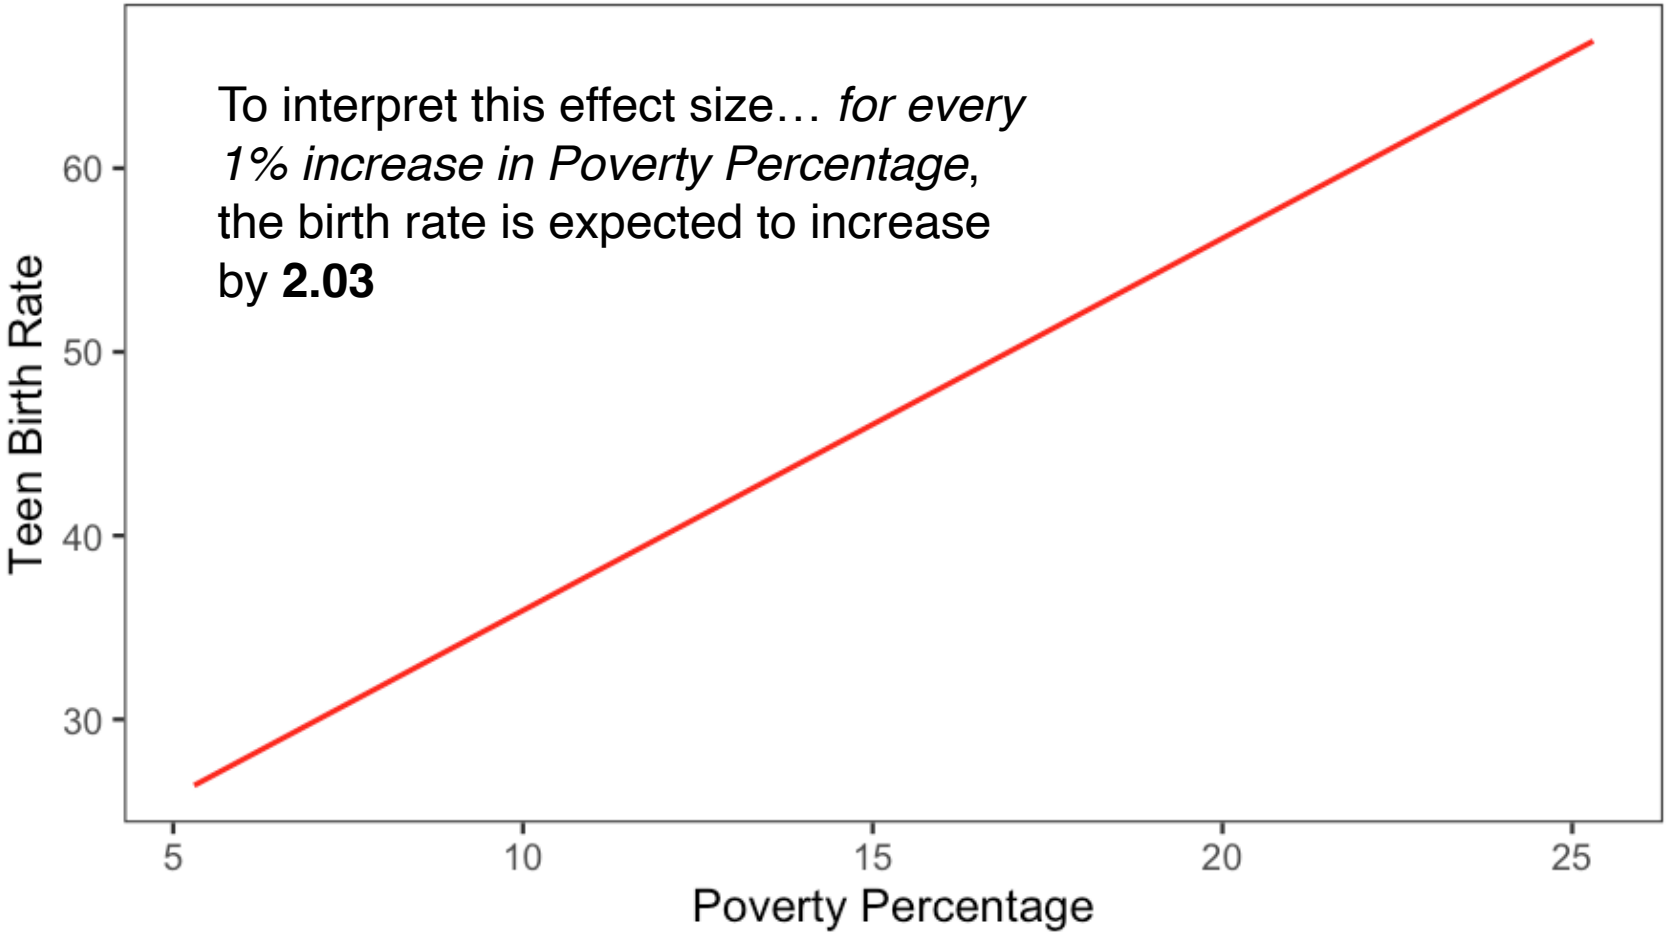
To interpret this effect size… *for every 1% increase in Poverty Percentage*, the birth rate is expected to increase by **2.03**
Teen Birth Rate
60
50
40
30
5
10
15
20
25
Poverty Percentage

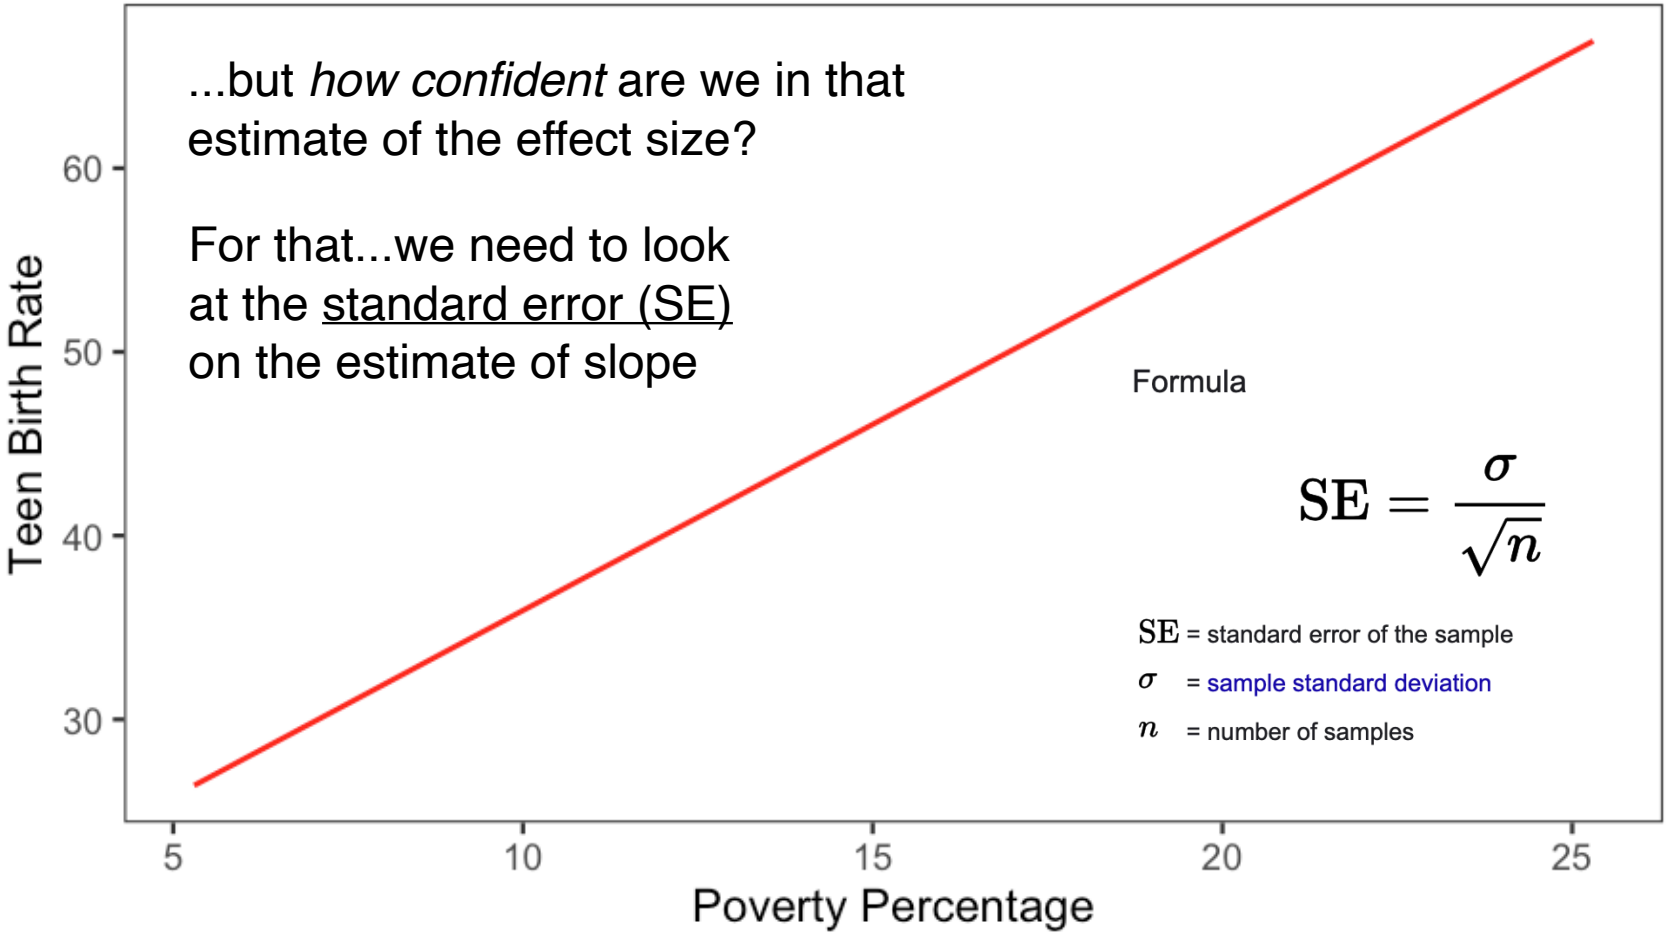
...but *how confident* are we in that estimate of the effect size?
For that...we need to look at the standard error (SE) on the estimate of slope
Formula
$$SE = \frac{\sigma}{\sqrt{n}}$$
SE = standard error of the sample
σ = sample standard deviation
n = number of samples
Teen Birth Rate
60
50
40
30
5
10
15
20
25
Poverty Percentage

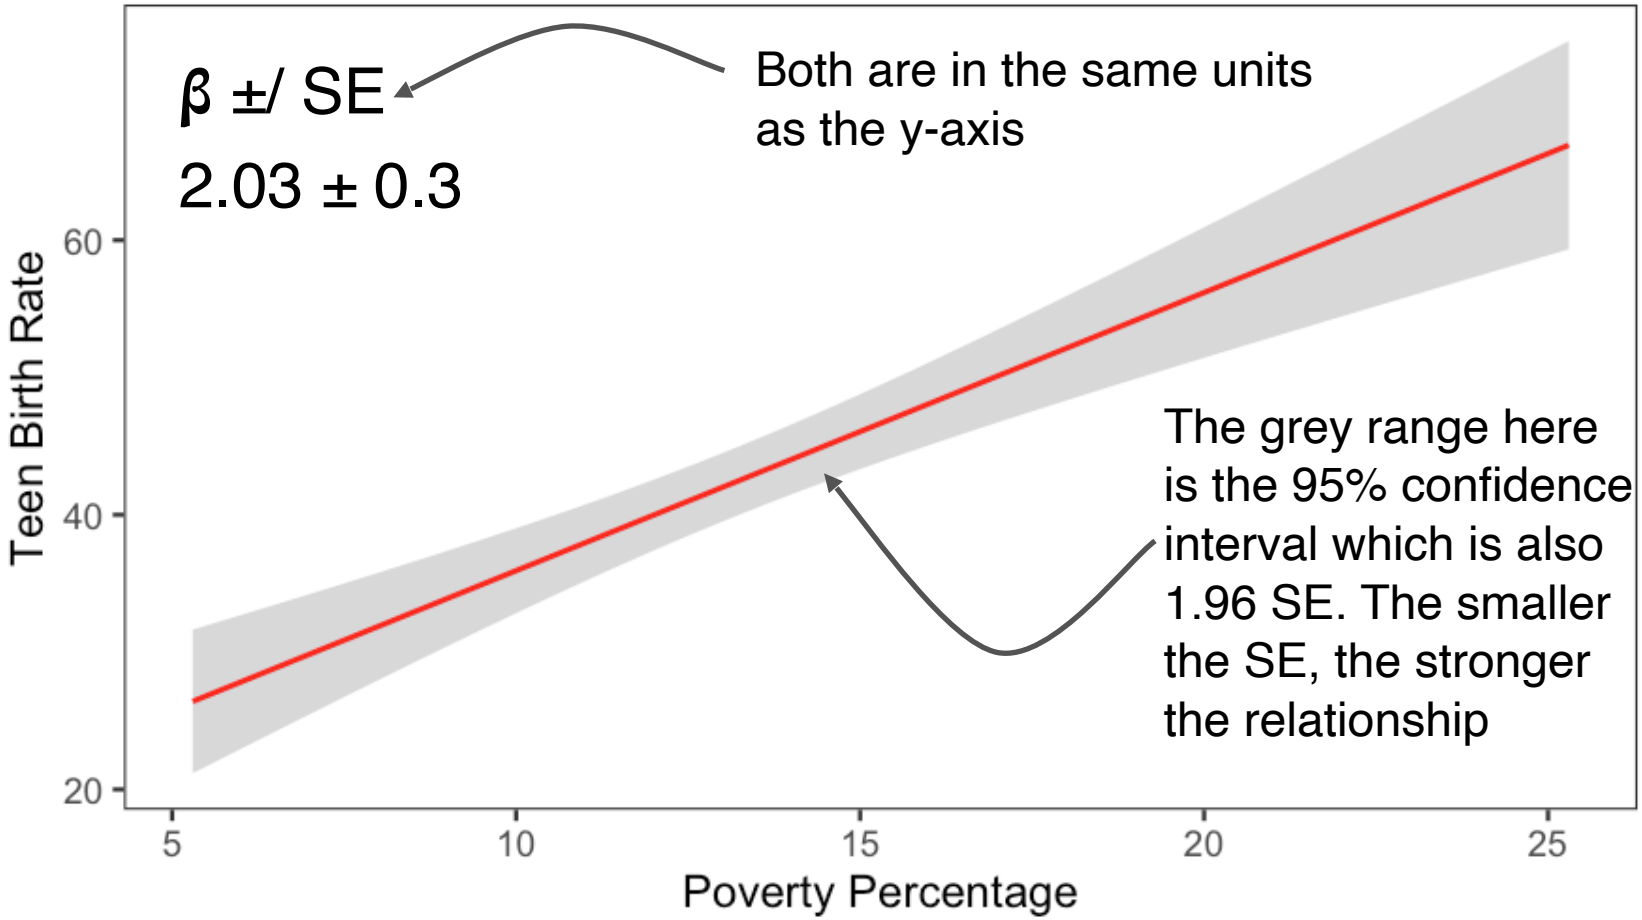
β ±/ SE
2.03 ± 0.3
Both are in the same units as the y-axis
The grey range here is the 95% confidence interval which is also 1.96 SE. The smaller the SE, the stronger the relationship
Teen Birth Rate
60
40
20
5
10
15
20
25
Poverty Percentage

p-value : the probability of getting the observed results (or results more extreme) by chance alone

Takes into account the effect size (β) and the SE

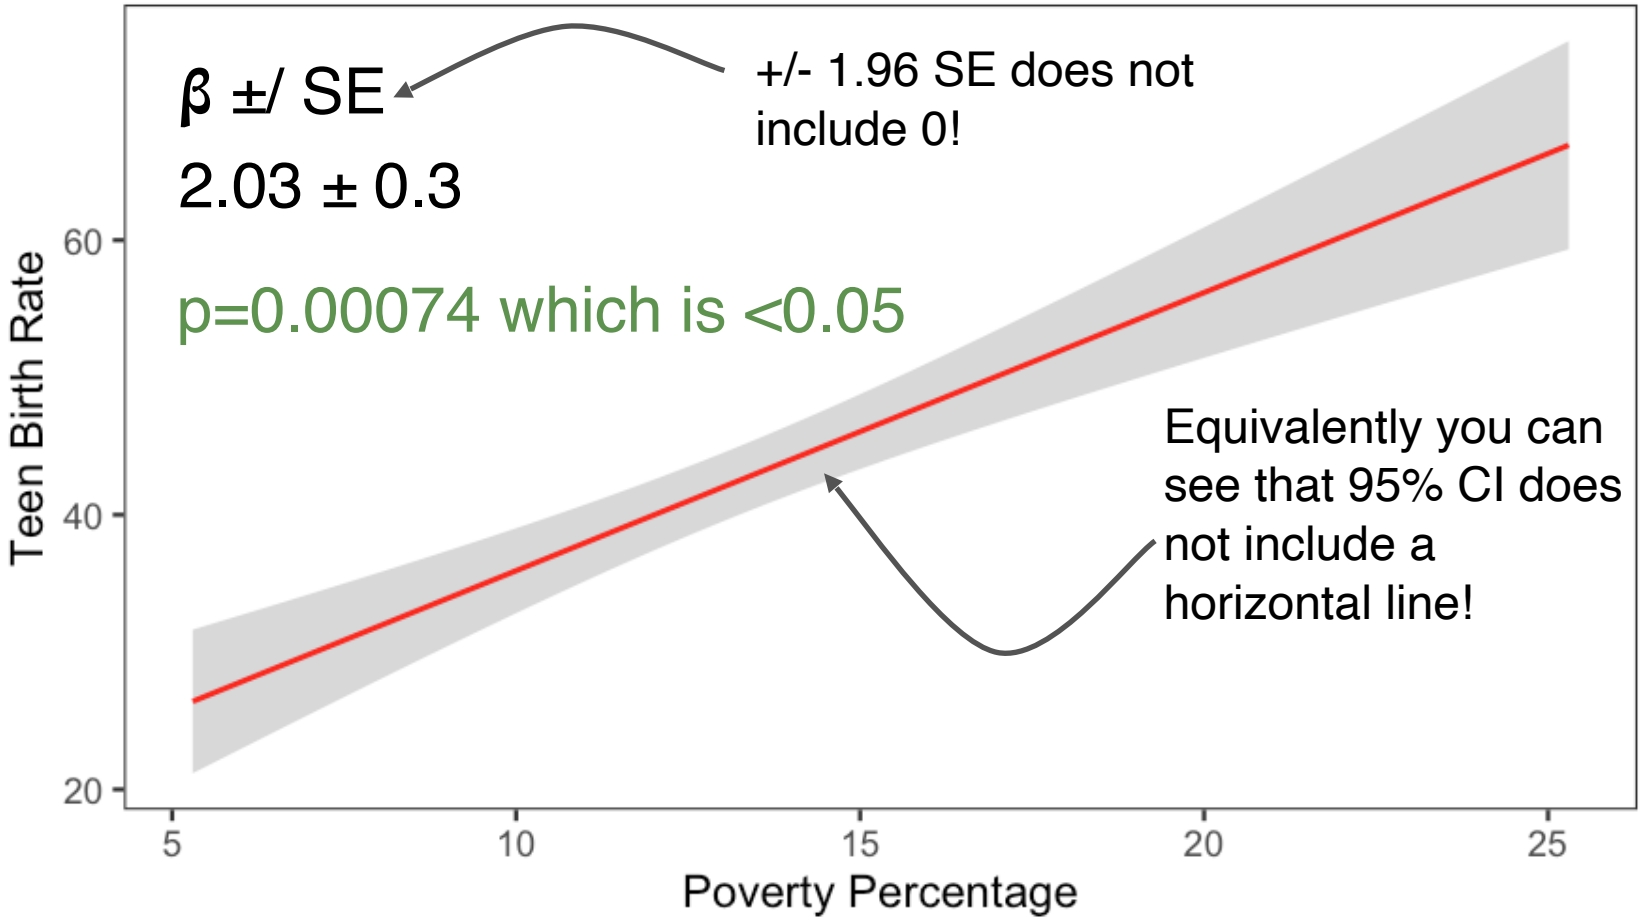
β ±/ SE
2.03 ± 0.3
+/- 1.96 SE does not include 0!
p=0.00074 which is <0.05
Equivalently you can see that 95% CI does not include a horizontal line!
Teen Birth Rate
60
40
20
5
10
15
20
25
Poverty Percentage

Assumptions of linear regression

1. Linear relationship
2. Multivariate normality
3. No multicollinearity
4. No autocorrelation
5. Homoscedasticity

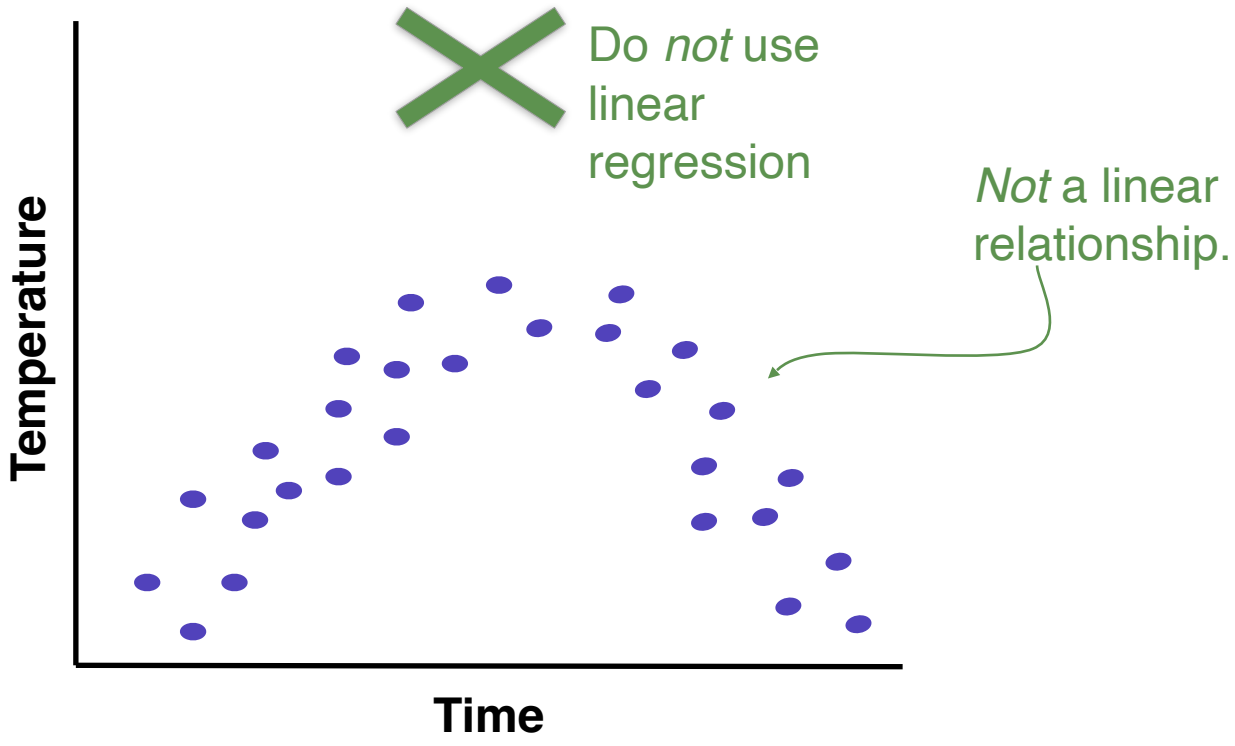
Do *not* use linear regression
Not a linear relationship.
Temperature
Time

A multivariate normal probability distribution (joint normal)

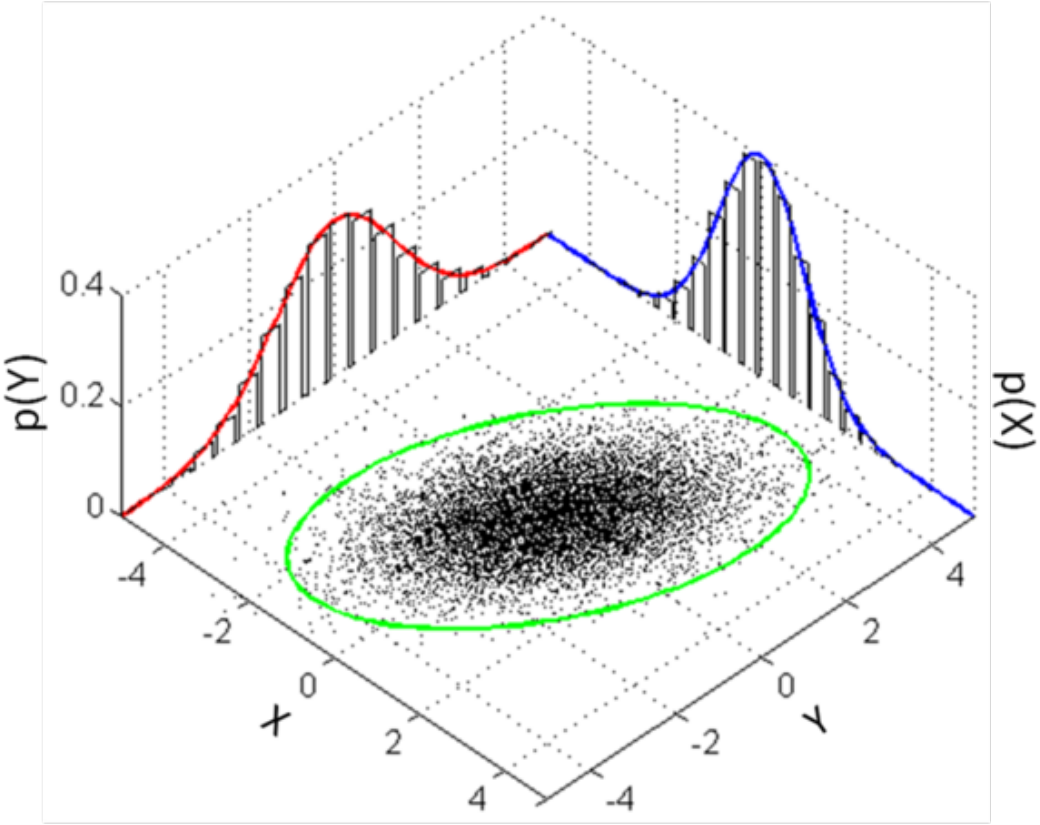

Linear regression assumes no multicollinearity. Multicollinearity occurs when the independent variables (in multiple linear regression) are too highly correlated with each other.

Daily returns are
$\ln(\text{price}_t / \text{price}_{t-1})$

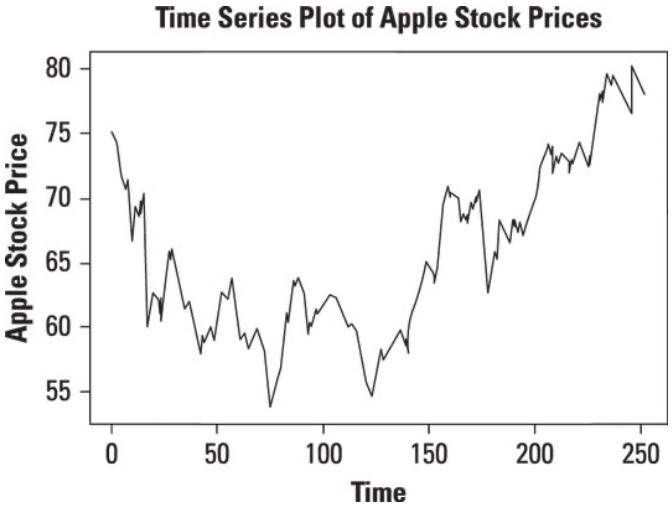

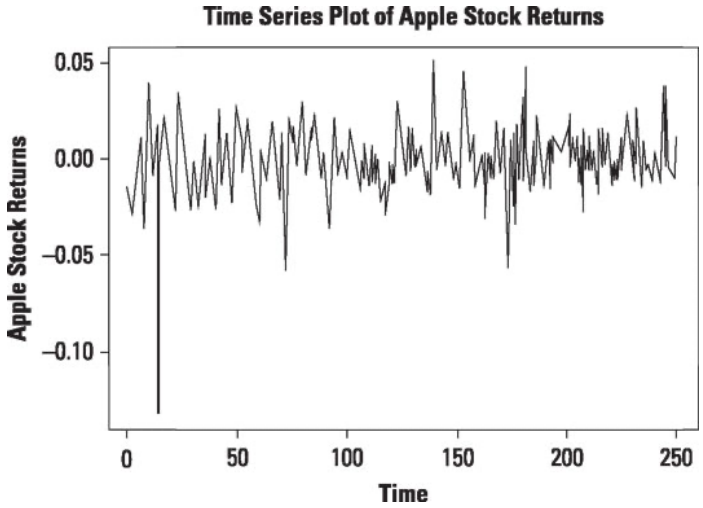

Autocorrelation occurs when the observations are *not* independent of one another (i.e. stock prices)

https://www.dummies.com/programming/big-data/data-science/autocorrelation-plots-graphical-technique-for-statistical-data//

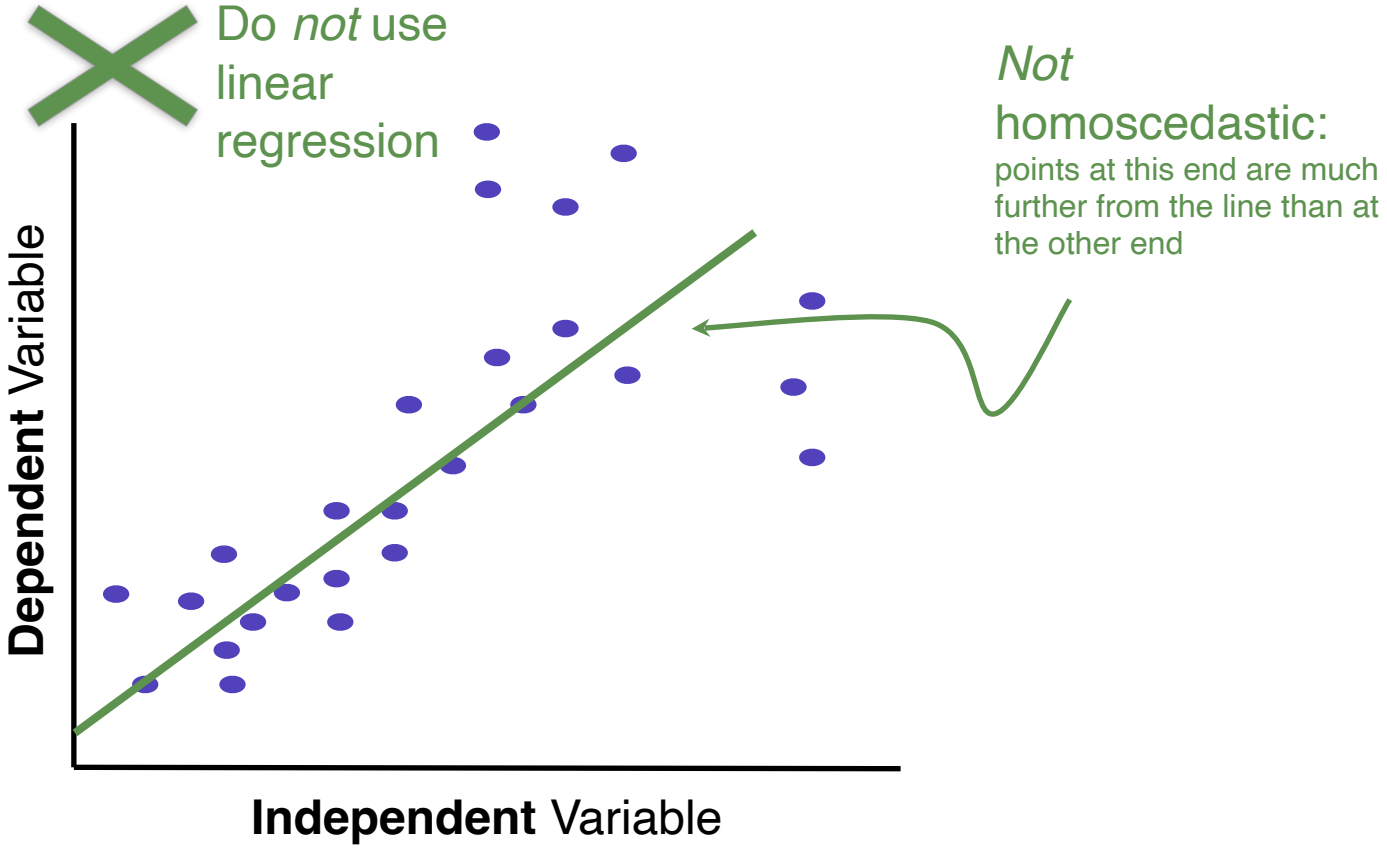
Do *not* use linear regression
Not homoscedastic:
points at this end are much further from the line than at the other end
Dependent Variable
Independent Variable

Does Poverty Percentage affect Teen Birth Rate?

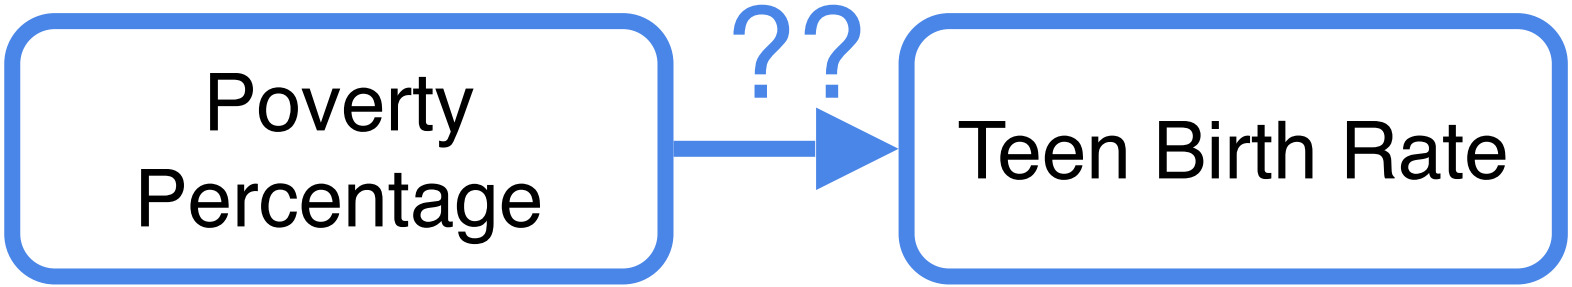
Poverty Percentage
??
Teen Birth Rate

What is the relationship between Poverty Percentage & Teen Birth Rate?

What's your hypothesis?

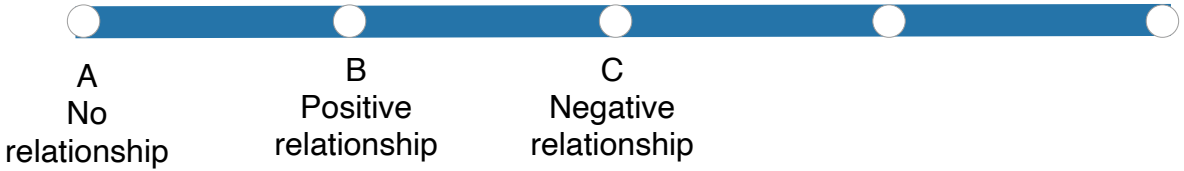

	Location	PovPct	Brth15to17	Brth18to19	ViolCrime	TeenBrth
1	Alabama	20.1	31.5	88.7	11.2	54.5
2	Alaska	7.1	18.9	73.7	9.1	39.5
3	Arizona	16.1	35.0	102.5	10.4	61.2
4	Arkansas	14.9	31.6	101.7	10.4	59.9
5	California	16.7	22.6	69.1	11.2	41.1
6	Colorado	8.8	26.2	79.1	5.8	47.0
7	Connecticut	9.7	14.1	45.1	4.6	25.8
8	Delaware	10.3	24.7	77.8	3.5	46.3
9	District_of_Columbia	22.0	44.8	101.5	65.0	69.1
10	Florida	16.2	23.2	78.4	7.3	44.5
11	Georgia	12.1	31.4	92.8	9.5	55.7
12	Hawaii	10.3	17.7	66.4	4.7	38.2
13	Idaho	14.5	18.4	69.1	4.1	39.1
14	Illinois	12.4	23.4	70.5	10.3	42.2
15	Indiana	9.6	22.6	78.5	8.0	44.6
16	Iowa	12.2	16.4	55.4	1.8	32.5
17	Kansas	10.8	21.4	74.2	6.2	43.0

Normal(ish) distributions

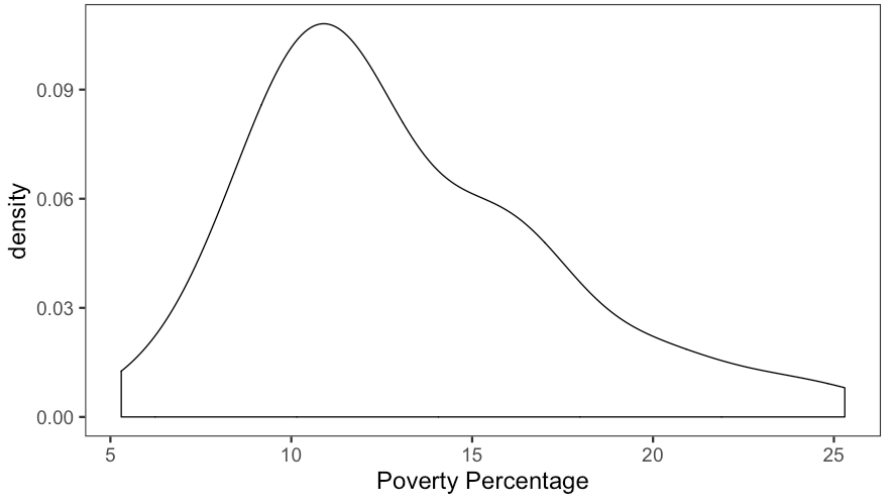

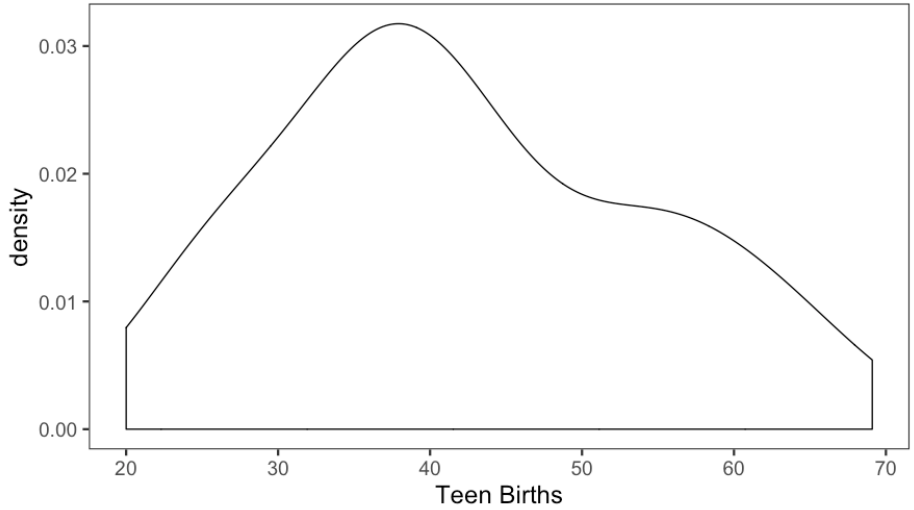

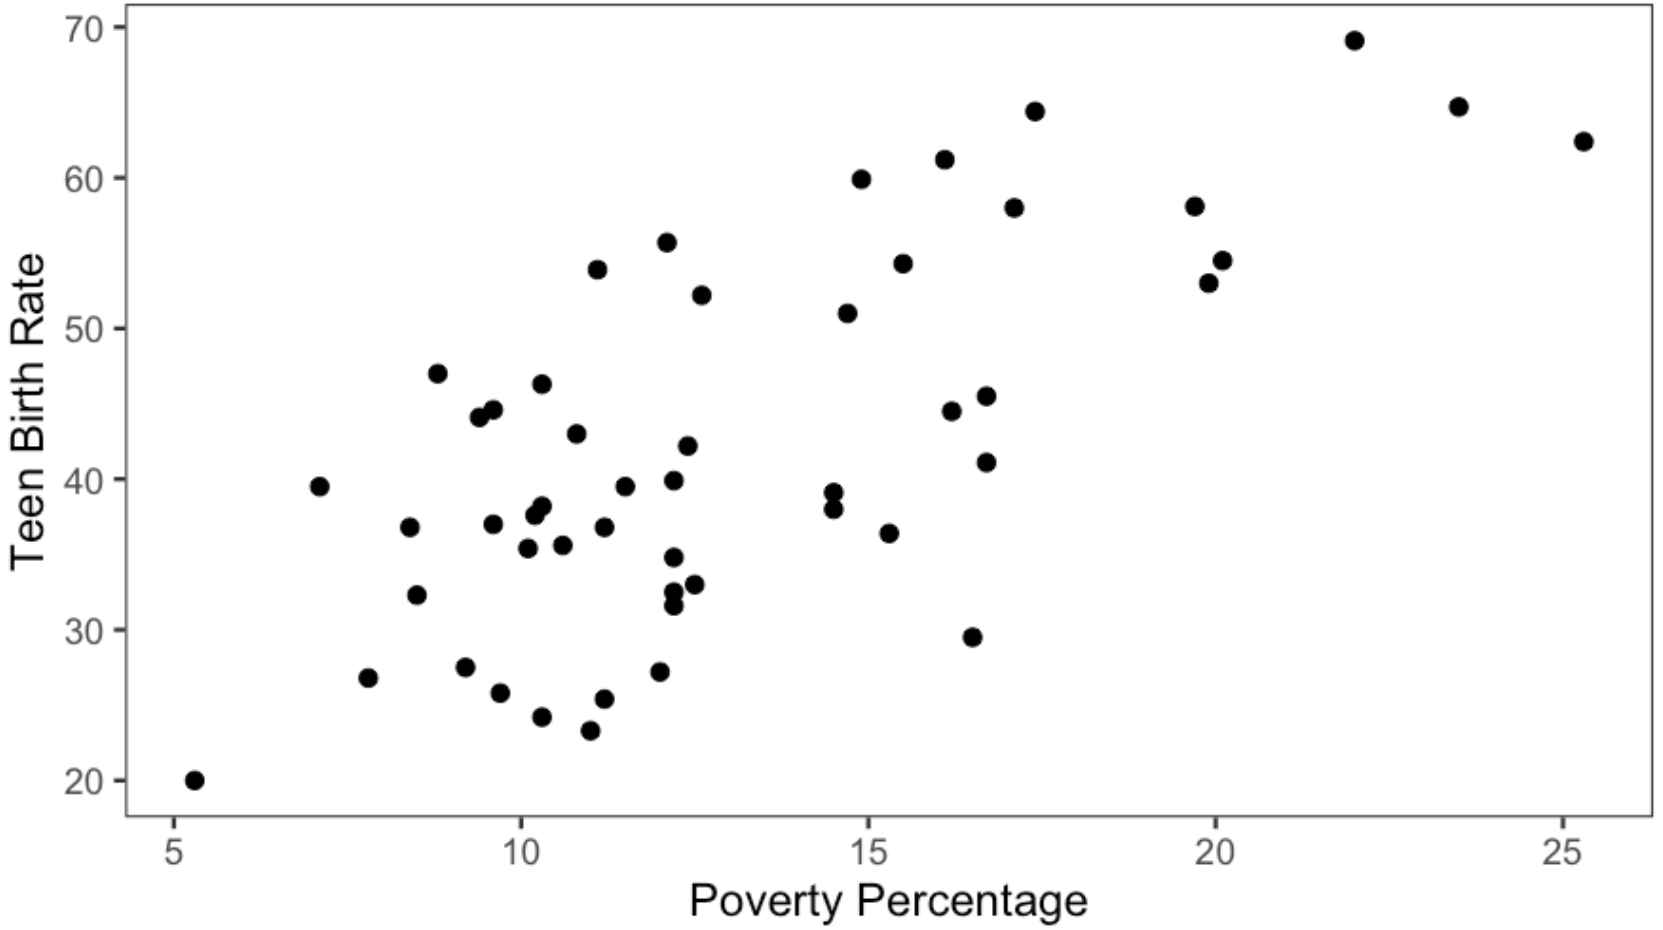
Teen Birth Rate
70
60
50
40
30
20
5
10
15
20
25
Poverty Percentage

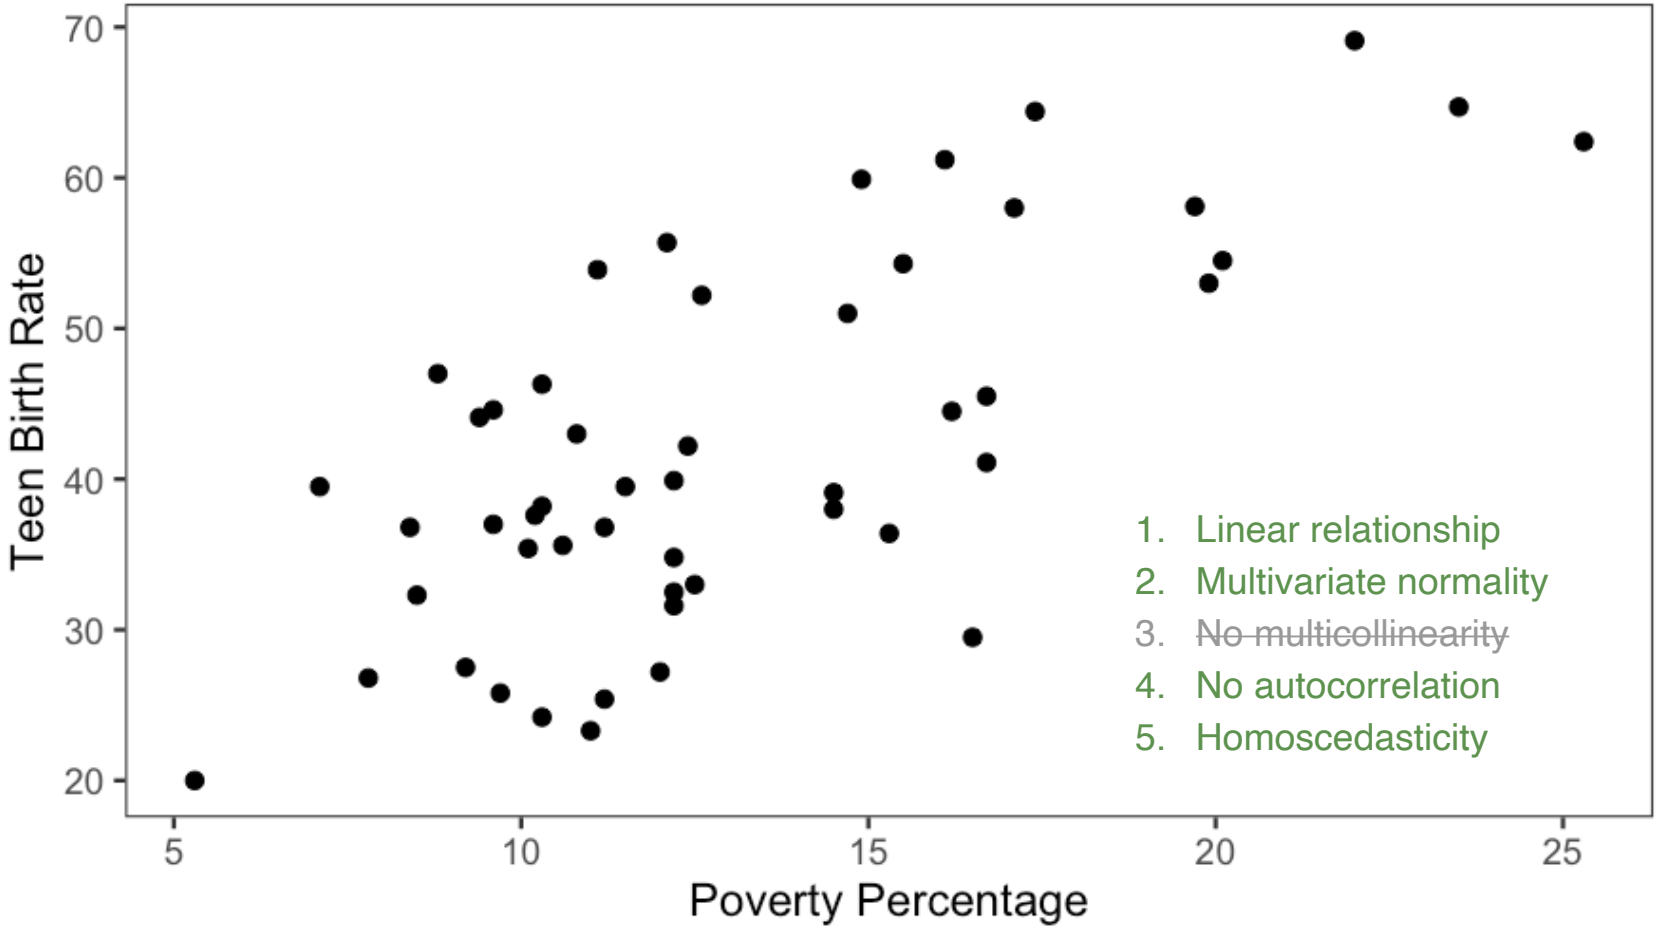
Teen Birth Rate
70
60
50
40
30
20
5
10
15
20
25
Poverty Percentage
1. Linear relationship
2. Multivariate normality
3. ~~No multicollinearity~~
4. No autocorrelation
5. Homoscedasticity

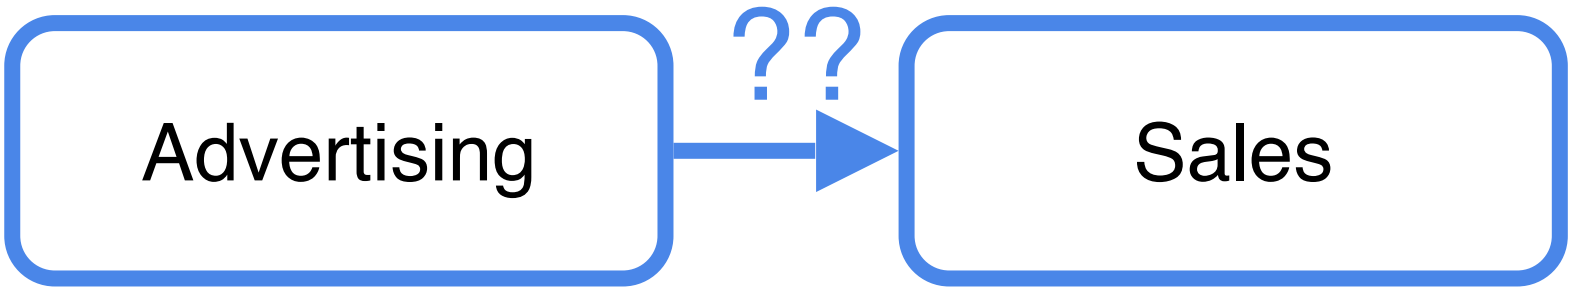
Advertising
??
Sales

Effect size interpretation

Sales (Million Euro)	Advertising (Million Euro)
651	23
762	26
856	30
1,063	34
1,190	43
1,298	48
1,421	52
1,440	57
1,518	58

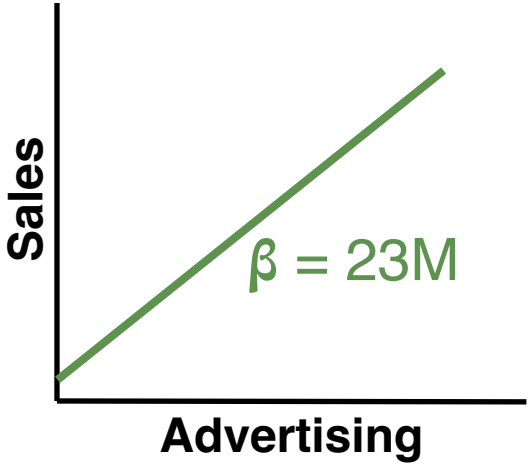

The effect size (β) between the advertising and sales is 23M. What does this mean?

A
For every 1M Euro spent on advertising, the company sees 23M more in sales

B
For every 1M Euro spent in sales, the company spends 23M more in advertising

C
For every 1M Euro spent on advertising, the company sees 24M less in sales

D
For every 1M Euro spent in sales, the company spends 23M less in advertising

Model selection problems

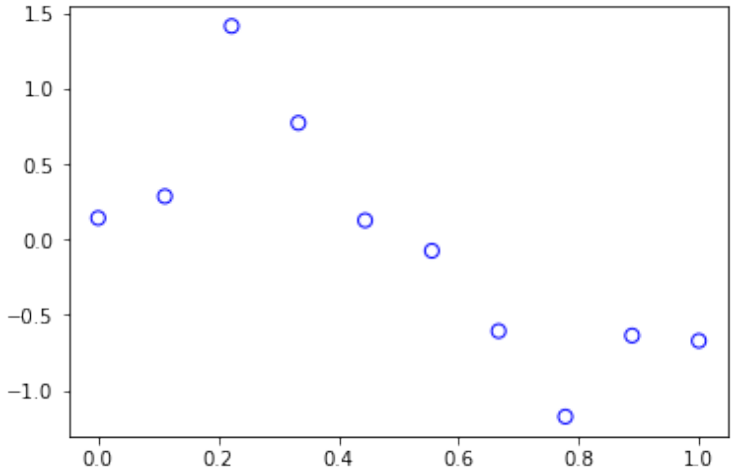
1.5
1.0
0.5
0.0
−0.5
−1.0
0.0
0.2
0.4
0.6
0.8
1.0

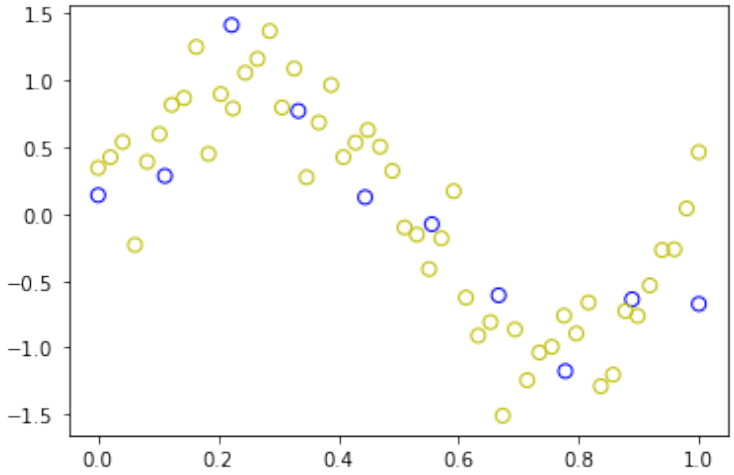
1.5
1.0
0.5
0.0
−0.5
−1.0
−1.5
0.0
0.2
0.4
0.6
0.8
1.0

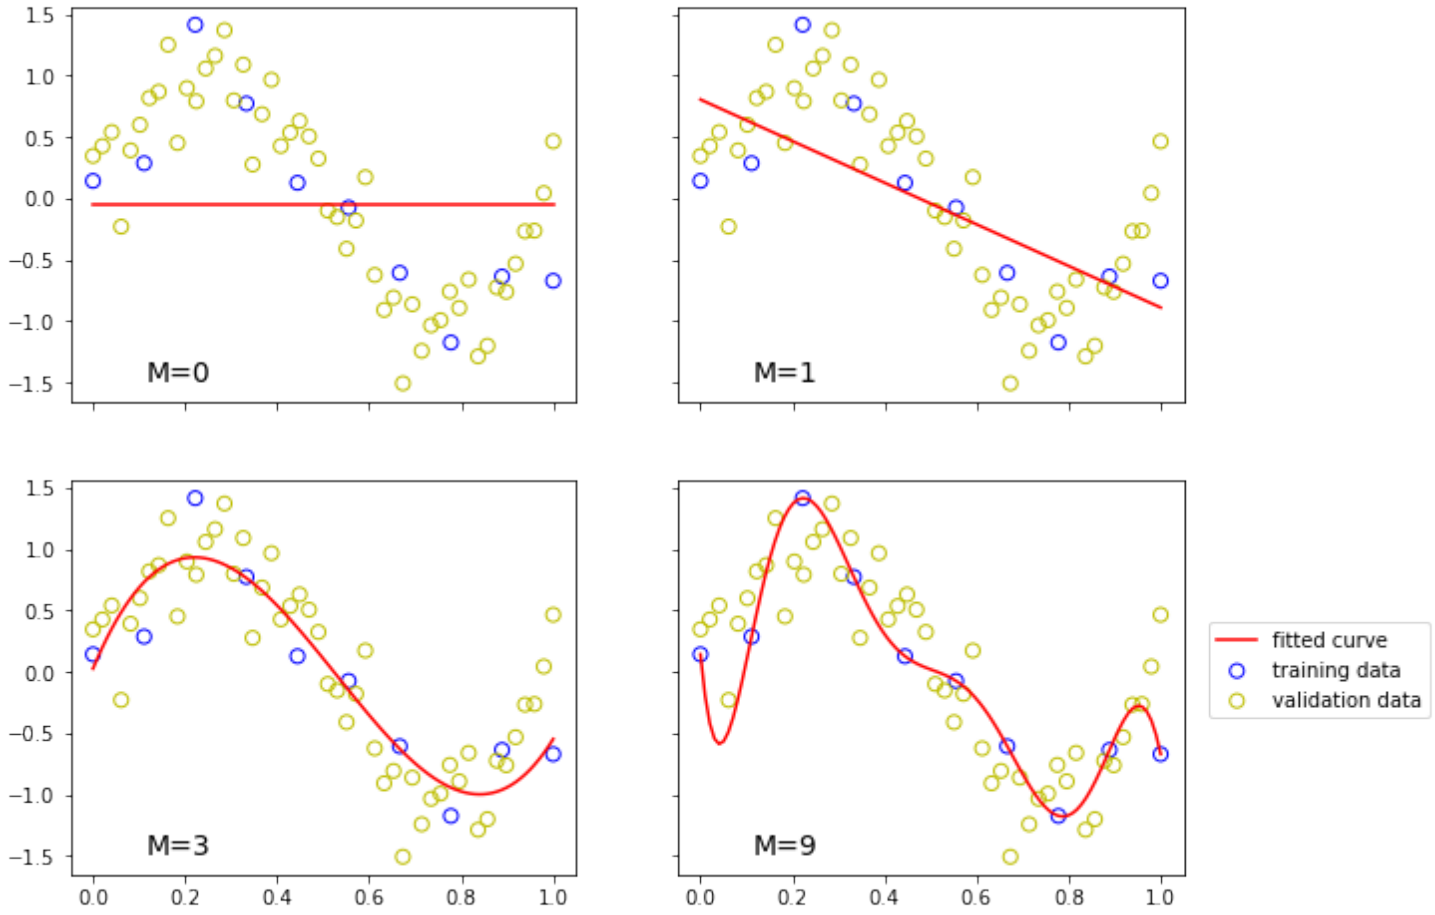
M=0
M=1
M=3
M=9
fitted curve
training data
validation data

Confounding

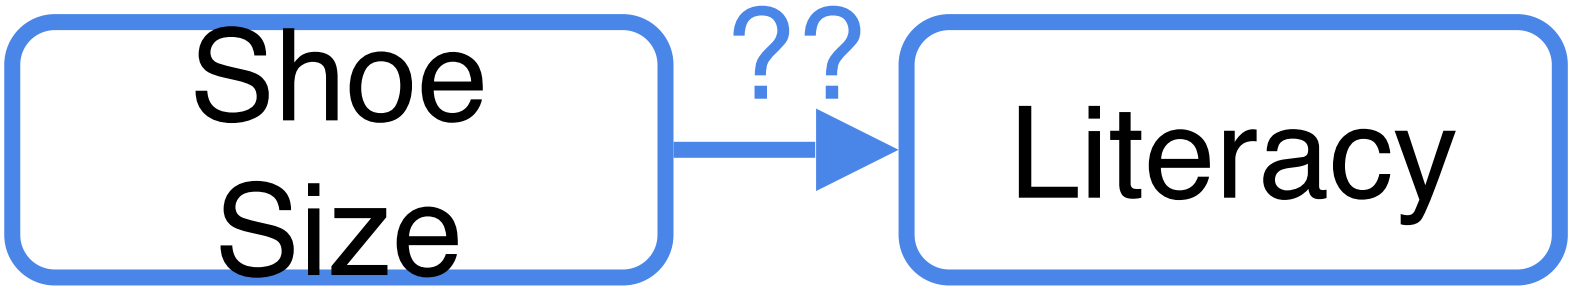
Shoe Size
??
Literacy


Small shoes
Not literate
Child
Big shoes
Literate
Adult

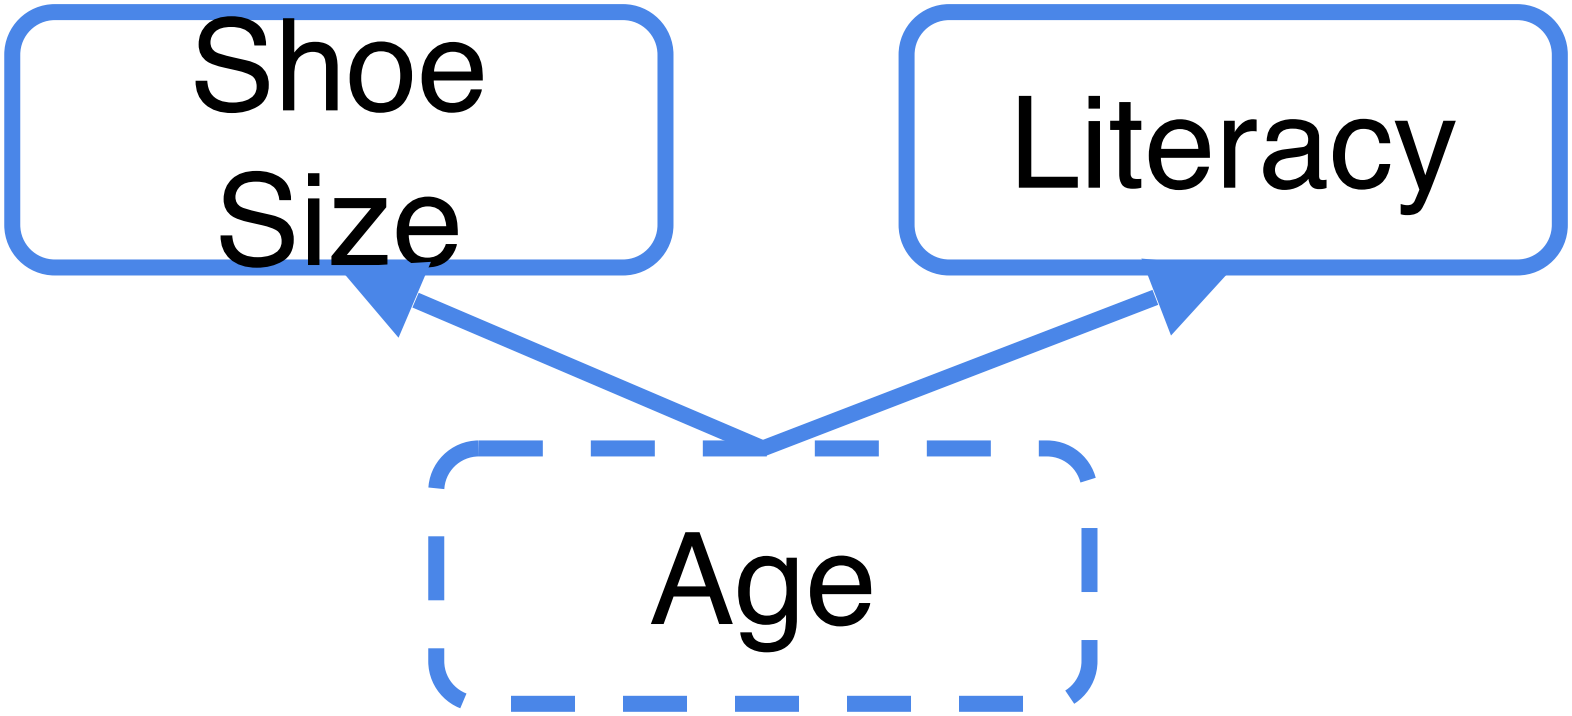
Shoe Size
Literacy
Age

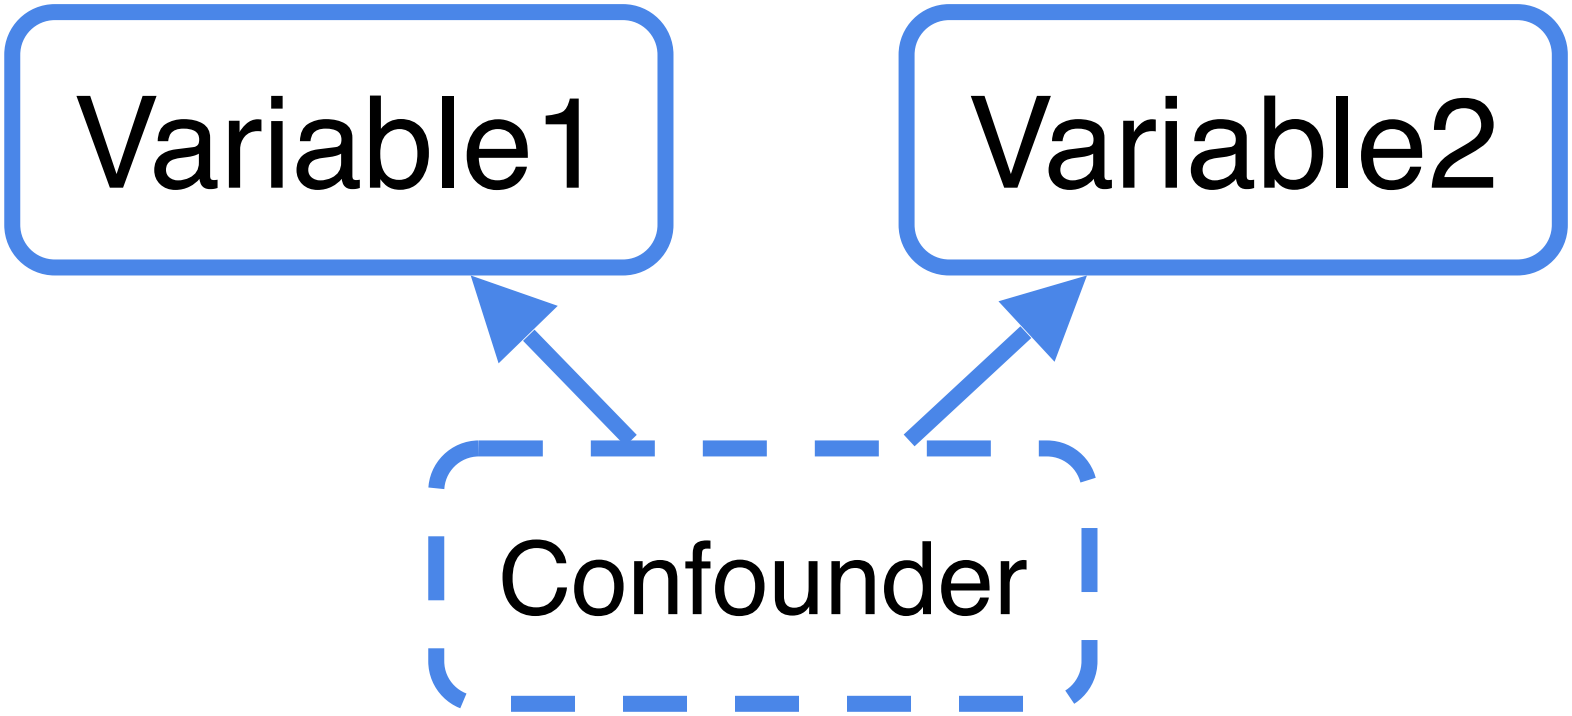
Variable1
Variable2
Confounder

Confounding

popsicles → crime rate

Your analysis sees an increase in crime rate whenever popsicle sales increase. What could confound this analysis?

A
popsicle preference

B
new gun laws

C
temperature

D
changes in popsicle prices

E
new law enforcement officers

You can plan ahead to avoid confounding and/or include confounders in your models to account for their role on the outcome variable.

Ignoring confounders will lead you to draw incorrect conclusions

Stratification changes results

Sample: 400 patients with index vertebral fractures

…looks like vertebroplasty was *way* worse for patients!

Vertebroplasty	Conservative care	Relative risk (95% confidence interval)
30/200 (15%)	15/200 (7.5%)	2.0 (1.1–3.6)

subsequent fractures

https://www.ncbi.nlm.nih.gov/pmc/articles/PMC3503514/

But wait...at time of initial fracture...

	Vertebroplasty N = 200	Conservative care N = 200
Age, y, mean ± SD	78.2 ± 4.1	79.0 ± 5.2
Weight, kg, mean ± SD	54.4 ± 2.3	53.9 ± 2.1
Smoking status, No. (%)	110 (55)	16 (8)

Age and weight are similar between groups. **Smoking Status** differs vastly.

https://www.ncbi.nlm.nih.gov/pmc/articles/PMC3503514/

So...let's stratify those results real quick

Smoke			**No smoke**		
Vertebroplasty	Conservative	RR (95% confidence interval)	Vertebroplasty	Conservative	RR (95% confidence interval)
23/110 (21%)	3/16 (19%)	1.1 (0.4, 3.3)	7/90 (8%)	12/184(7%)	1.2 (0.5, 2.9)

Risk of re-fracture is now similar within group

https://www.ncbi.nlm.nih.gov/pmc/articles/PMC3503514/